

# MPLS Layer 3 VPNs: From Fundamentals to Expert Implementation

## Module 1: VPNs and MPLS - The Foundation

### Part 1: The Conceptual Lecture (The Why)

#### The Business Problem


Imagine you're a large corporation with 50 offices worldwide. Each office needs to communicate securely with every other office as if they were all on the same private network. You have two fundamental challenges:

1. **Security:** Internet traffic is public - anyone can potentially intercept it
2. **Isolation:** Your company's traffic must remain separate from other companies' traffic

This is where Virtual Private Networks (VPNs) come in. Think of a VPN as a private tunnel through public infrastructure - like having your own private lane on a public highway.

### IPsec VPNs vs MPLS VPNs: Two Different Philosophies

#### IPsec VPNs: The Encryption Approach




IPsec VPNs work by:

- **Encrypting** every packet at the source
- Sending it over the **public Internet**
- **Decrypting** it at the destination

Think of it like sending a locked box through regular mail - anyone can handle the box, but only you have the key.

#### MPLS VPNs: The Isolation Approach



MPLS VPNs work by:

- Using **labels** to create isolated paths through the provider network
- Traffic stays within the **provider's controlled network**
- No encryption needed (though can be added)


Think of it like having reserved train cars on a railway - your cargo travels in designated cars that only you can access.

#### Key Differences:

Aspect	IPsec VPN	MPLS VPN
Infrastructure	Public Internet	Provider's private network
Security Method	Encryption	Isolation via labels
Complexity	Complex ( $N \times (N-1)/2$ tunnels for full mesh)	Simple (provider handles complexity)
Scalability	Poor (each site needs tunnel to every other site)	Excellent (provider manages connections)
QoS	Best-effort Internet	Guaranteed by provider
Cost	Low (Internet connection)	Higher (premium service)

### Layer 3 VPNs vs Layer 2 VPNs: Different Network Layers

#### Layer 2 VPNs (L2VPNs)




|  
Devices think they're on the same LAN segment (broadcast domain)

L2VPNs:

- Transport **Ethernet frames** (Layer 2)
- Customer sites appear to be on the **same LAN**
- Customer manages all **routing**
- Provider is transparent - just extends the LAN

### Layer 3 VPNs (L3VPNs)



L3VPNs:

- Transport **IP packets** (Layer 3)
- Provider **participates in routing**
- Each site is a separate network
- Provider PE routers learn and advertise customer routes

## Part 2: The Junos CLI Masterclass (The How)

For Module 1, we'll configure a basic comparison setup showing both VPN types.

### Basic MPLS Configuration (Foundation for MPLS VPNs)

```
[edit]
# First, enable MPLS on interfaces
set interfaces ge-0/0/0 unit 0 family mpls
set interfaces ge-0/0/1 unit 0 family mpls

# Enable MPLS protocol
set protocols mpls interface ge-0/0/0.0
set protocols mpls interface ge-0/0/1.0

# Configure RSVP for label distribution
set protocols rsvp interface ge-0/0/0.0
set protocols rsvp interface ge-0/0/1.0

# Create Label-Switched Paths (LSPs)
set protocols mpls label-switched-path T0-PE2 {
    to 192.168.1.2;      # Destination PE router
    primary STRICT-PATH; # Use explicit path
}

set protocols mpls path STRICT-PATH {
    10.0.0.1 strict;     # P router 1
    10.0.0.2 strict;     # P router 2
}
```

### IPsec VPN Configuration Example

```
[edit security]
# IKE (Phase 1) Configuration
set ike proposal IKE-PROPOSAL {
    authentication-method pre-shared-keys;
    dh-group group14;
    authentication-algorithm sha-256;
    encryption-algorithm aes-256-cbc;
    lifetime-seconds 86400;
}

set ike policy IKE-POLICY {
    mode main;
    proposals IKE-PROPOSAL;
    pre-shared-key ascii-text "SecretKey123!";
}
```

```

set ike gateway IKE-GATEWAY {
    ike-policy IKE-POLICY;
    address 203.0.113.1; # Remote site public IP
    external-interface ge-0/0/0.0;
}

# IPsec (Phase 2) Configuration
set ipsec proposal IPSEC-PROPOSAL {
    protocol esp;
    authentication-algorithm hmac-sha-256-128;
    encryption-algorithm aes-256-cbc;
    lifetime-seconds 3600;
}

set ipsec policy IPSEC-POLICY {
    proposals IPSEC-PROPOSAL;
}

set ipsec vpn SITE-TO-SITE {
    bind-interface st0.0; # Tunnel interface
    ike {
        gateway IKE-GATEWAY;
        ipsec-policy IPSEC-POLICY;
    }
}

# Tunnel Interface Configuration
set interfaces st0 unit 0 {
    family inet {
        address 10.255.255.1/30; # Tunnel IP
    }
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### For MPLS:

```

# Verify MPLS interfaces are up
user@router> show mpls interface
Interface      State      Administrative groups (x: extended)
ge-0/0/0.0     Up        <none>
ge-0/0/1.0     Up        <none>

# Check MPLS LSPs
user@router> show mpls lsp
Ingress LSP: 1 sessions
To          From          State Rt P      ActivePath      LSPname
192.168.1.2 192.168.1.1 Up      0 *      STRICT-PATH    T0-PE2

# Verify RSVP neighbors
user@router> show rsvp neighbor
RSVP neighbor: 2 learned
Address          Idle Up/Dn LastChange HelloInt HelloTx/Rx MsgRcvd
10.0.0.1         0   1/0   00:45:23   9       301/301   125

```

#### For IPsec VPN:

```

# Check IKE Phase 1
user@router> show security ike security-associations
Index  State  Initiator cookie  Responder cookie  Mode           Remote Address
1234567 UP    a1b2c3d4e5f6g7h8 i9j0k1l2m3n4o5p6 Main          203.0.113.1

# Check IPsec Phase 2
user@router> show security ipsec security-associations
Total active tunnels: 1
ID    Algorithm      SPI      Life:sec/kb  Mon lsys Port  Gateway
<2    ESP:aes-cbc-256/sha256 0x1234abcd 3543/ unlim - root 500 203.0.113.1
>2    ESP:aes-cbc-256/sha256 0xabcd1234 3543/ unlim - root 500 203.0.113.1

```

### Common Troubleshooting Scenarios

#### Scenario 1: MPLS LSP Not Establishing

- **Symptom:** show mpls lsp shows state as "Down"

- **Diagnostic Command:**

```
user@router> show mpls lsp extensive
[Output shows: "CSPF failed: no route to 192.168.1.2"]
```

- **Cause:** RSVP not enabled on intermediate routers

- **Solution:**

```
# On each P router:
set protocols rsvp interface all
set protocols mpls interface all
```

#### Scenario 2: IPsec Tunnel Up but No Traffic

- **Symptom:** Tunnel shows UP but ping fails

- **Diagnostic Command:**

```
user@router> show route table inet.0 10.2.0.0/24
[No route found]
```

- **Cause:** Missing static route through tunnel

- **Solution:**

```
set routing-options static route 10.2.0.0/24 next-hop st0.0
```

## Module 2: Layer 3 VPNs Overview - Understanding the Architecture

### Part 1: The Conceptual Lecture (The Why)


#### The Multi-Tenancy Challenge

Service providers face a unique challenge: they need to provide VPN services to thousands of customers using the same physical network infrastructure. Each customer must:

- Have complete **privacy** (Customer A can't see Customer B's traffic)
- Use **any IP addresses** they want (including overlapping addresses)
- Feel like they have their own **private network**

This is like an apartment building where every tenant can paint their walls any color, play any music, and arrange furniture however they want - without affecting their neighbors.

#### L3VPN Terminology - The Cast of Characters



#### The Players:

##### 1. CE (Customer Edge) Router

- The customer's router at each site
- Doesn't know about MPLS or VPNs
- Just does normal IP routing

##### 2. PE (Provider Edge) Router

- The provider's router connecting to customers
- The "smart" router that handles VPN logic
- Maintains separate routing tables per customer

### 3. P (Provider Core) Router

- Core routers in provider network
- Only switches MPLS labels
- Doesn't know about customer routes

### 4. VRF (Virtual Routing and Forwarding)

- A separate routing table for each customer
- Like having multiple virtual routers in one physical router
- Provides the isolation between customers

## The VPN-IPv4 Address: Solving the Overlap Problem

**The Problem:** What if Customer A uses 10.1.0.0/24 and Customer B also uses 10.1.0.0/24? How does the provider network distinguish between them?

### The Solution: VPN-IPv4 Addresses

A regular IPv4 address is 32 bits:

```
10.1.0.1 = [32 bits]
```

A VPN-IPv4 address is 96 bits (32 + 64):

```
RD:IPv4-address = [64-bit Route Distinguisher][32-bit IPv4 address]
65000:100:10.1.0.1 = Unique globally!
65000:200:10.1.0.1 = Different, also unique!
```

### Route Distinguisher (RD) Structure:

The RD is 64 bits with format: Type:Administrator:Assigned

Common formats:

1. **Type 0:** ASN:Number (e.g., 65000:100)
  - ASN is 16 bits ( $2^{16} = 65,536$  values)
  - Number is 32 bits ( $2^{32} = 4.3$  billion values)
2. **Type 1:** IP:Number (e.g., 192.168.1.1:100)
  - IP is 32 bits (IPv4 address)
  - Number is 16 bits ( $2^{16} = 65,536$  values)


Think of it like apartment addressing:

- Regular address: "Apartment 10"
- With RD: "Building-65000-Floor-100-Apartment-10"

Now each apartment is globally unique!

## Route Targets: The Distribution Mechanism

While RDs make routes unique, **Route Targets (RTs)** control which sites receive which routes.



Think of Route Targets like mailing lists:

- **Export RT:** "Publish my routes to this mailing list"
- **Import RT:** "Subscribe to routes from this mailing list"

Common RT Designs:

1. **Full Mesh:** All sites import/export same RT
2. **Hub-and-Spoke:** Spokes export to one RT, import from another
3. **Extranet:** Selective sharing between VPNs

## Part 2: The Junos CLI Masterclass (The How)

### Basic L3VPN Configuration Structure

```

[edit]
# Step 1: Configure the VRF (Virtual Routing and Forwarding Instance)
set routing-instances CUSTOMER-A {
    instance-type vrf;                      # This is a VRF-type instance
    interface ge-0/0/2.0;                     # CE-facing interface
    route-distinguisher 65000:100;           # Makes routes unique
    vrf-target target:65000:100;             # Import and export RT
    vrf-table-label;                        # Allocate label for VRF
}

# Detailed VRF configuration with separate import/export
set routing-instances CUSTOMER-A {
    instance-type vrf;
    interface ge-0/0/2.0;
    route-distinguisher 65000:100;
    vrf-import CUSTOMER-A-IMPORT;          # Import policy
    vrf-export CUSTOMER-A-EXPORT;          # Export policy
    vrf-table-label;
}

# Step 2: Define Import/Export Policies
set policy-options policy-statement CUSTOMER-A-IMPORT {
    term ACCEPT-ROUTES {
        from {
            protocol bgp;
            community RT-CUSTOMER-A;      # Match on RT community
        }
        then accept;
    }
    term REJECT-ALL {
        then reject;
    }
}

set policy-options policy-statement CUSTOMER-A-EXPORT {
    term ADD-RT {
        from {
            protocol [ direct ospf bgp static ]; # Routes to export
        }
        then {
            community add RT-CUSTOMER-A;      # Add RT community
            accept;
        }
    }
}

# Step 3: Define the Route Target Community
set policy-options community RT-CUSTOMER-A {
    members target:65000:100;              # Route Target value
}

# Step 4: Configure BGP between PE and CE (if using BGP)
set routing-instances CUSTOMER-A {
    protocols {
        bgp {
            group CE-PEERS {
                type external;
                peer-as 65001;            # Customer's AS number
                neighbor 10.1.1.1 {       # CE router IP
                    description "Customer A - Site 1 CE";
                }
            }
        }
    }
}

# Step 5: Configure MP-BGP between PE routers
set protocols bgp {
    group IBGP-PEERS {
        type internal;
        local-address 192.168.1.1;      # Loopback address
        family inet-vpn unicast;        # VPN-IPv4 family
        neighbor 192.168.1.2 {         # Remote PE loopback
            description "PE-2";
        }
        neighbor 192.168.1.3 {
    }
}

```

```

        description "PE-3";
    }
}
}

# Step 6: Configure MPLS and LDP on core-facing interfaces
set protocols mpls {
    interface ge-0/0/0.0;                      # Core-facing interface
    interface ge-0/0/1.0;
}

set protocols ldp {
    interface ge-0/0/0.0;
    interface ge-0/0/1.0;
}

```

## Complete PE Router Configuration Example

```

# Full working configuration for PE1 router
[edit]
# System basics
set system host-name PE1

# Interfaces
set interfaces ge-0/0/0 {
    description "To P-Core";
    unit 0 {
        family inet address 10.0.0.1/30;
        family mpls;
    }
}

set interfaces ge-0/0/2 {
    description "To Customer-A CE";
    unit 0 {
        family inet address 10.1.1.254/24;
    }
}

set interfaces lo0 {
    unit 0 {
        family inet address 192.168.1.1/32;
    }
}

# OSPF for infrastructure
set protocols ospf {
    area 0.0.0.0 {
        interface ge-0/0/0.0;
        interface lo0.0 passive;
    }
}

# MPLS and LDP
set protocols mpls interface ge-0/0/0.0;
set protocols ldp interface ge-0/0/0.0;

# MP-BGP configuration
set protocols bgp {
    local-as 65000;
    group IBGP-VPN {
        type internal;
        local-address 192.168.1.1;
        family inet-vpn unicast;
        cluster 192.168.1.1;          # If route reflector
        neighbor 192.168.1.2;
        neighbor 192.168.1.3;
    }
}

# Customer A VRF
set routing-instances CUSTOMER-A {
    instance-type vrf;
    interface ge-0/0/2.0;
    route-distinguisher 192.168.1.1:100; # Using IP:Number format
}
```

```

vrf-target target:65000:100;
vrf-table-label;
protocols {
    bgp {
        group CE {
            type external;
            peer-as 65001;
            neighbor 10.1.1.1 {
                description "Customer A - Site 1";
            }
        }
    }
}
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential L3VPN Verification Commands

```

# 1. Verify VRF configuration
user@PE1> show route instance CUSTOMER-A detail
CUSTOMER-A:
  Router ID: 10.1.1.254
  Type: vrf           State: Active
  Interfaces:
    ge-0/0/2.0
  Route-distinguisher: 192.168.1.1:100
  Vrf-import: [ __vrf-import-CUSTOMER-A-internal__ ]
  Vrf-export: [ __vrf-export-CUSTOMER-A-internal__ ]
  Vrf-import-target: [ target:65000:100 ]
  Vrf-export-target: [ target:65000:100 ]
  Fast-reroute-priority: low
  Tables:
    CUSTOMER-A.inet.0      : 5 routes (5 active, 0 holddown, 0 hidden)

# 2. Check VRF routing table
user@PE1> show route table CUSTOMER-A.inet.0
CUSTOMER-A.inet.0: 5 destinations, 5 routes (5 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

10.1.1.0/24      *[Direct/0] 00:10:23
                  > via ge-0/0/2.0
10.1.1.254/32   *#[Local/0] 00:10:23
                  Local via ge-0/0/2.0
10.2.0.0/24      *#[BGP/170] 00:05:15, MED 0, localpref 100
                  AS path: 65001 I
                  > to 10.1.1.1 via ge-0/0/2.0

# 3. Verify BGP VPN routes
user@PE1> show route table bgp.l3vpn.0
bgp.l3vpn.0: 10 destinations, 10 routes (10 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

192.168.1.1:100:10.1.1.0/24
  *#[BGP/170] 00:15:23, localpref 100, from 192.168.1.1
  AS path: I
  > to 10.0.0.2 via ge-0/0/0.0, Push 299792

# 4. Check MP-BGP neighbors
user@PE1> show bgp summary
Threading mode: BGP I/O
Groups: 2 Peers: 3 Down peers: 0
Table          Tot Paths  Act Paths Suppressed      History Damp State     Pending
bgp.l3vpn.0      10        10        0            0        0        0        0
inet.0          0         0         0            0        0        0        0
Peer           AS       InPkt     OutPkt    OutQ   Flaps Last Up/Dwn State
192.168.1.2    65000     234      245        0      0      1:45:23 Establ
bgp.l3vpn.0: 5/5/5/0
192.168.1.3    65000     228      241        0      0      1:44:15 Establ
bgp.l3vpn.0: 5/5/5/0

```

### Common Troubleshooting Scenarios

#### Scenario 1: Routes Not Appearing in Remote VRF

- **Symptom:** Customer can't reach remote sites

- **Diagnostic Commands:**

```
# Check if routes are being exported
user@PE1> show route advertising-protocol bgp 192.168.1.2 table bgp.l3vpn
[No output - routes not being advertised]
```

```
# Check export policy
user@PE1> show configuration routing-instances CUSTOMER-A vrf-export
[Missing or incorrect]
```

- **Cause:** Missing or incorrect Route Target configuration

- **Solution:**

```
set routing-instances CUSTOMER-A vrf-target target:65000:100
commit
```

### Scenario 2: VPN Routes Present but No Connectivity

- **Symptom:** Routes visible but ping fails

- **Diagnostic Commands:**

```
# Check MPLS forwarding
user@PE1> show route forwarding-table family mpls
Label          Type      Next hop           Interface
299792        VPN       10.1.1.1          ge-0/0/2.0

# Trace MPLS path
user@PE1> traceroute mpls ldp 192.168.1.2
Probe options: ttl 64, retries 3, wait 10, paths 16, exp 7, fanout 16

  ttl   Label  Protocol    Address      Previous Hop     Probe Status
    1    299856  LDP        10.0.0.2    (null)          Success
    2    299840  LDP        10.0.0.6    10.0.0.2        Success
    3            Unknown     192.168.1.2  10.0.0.6        Egress
```

- **Cause:** MPLS not enabled on intermediate interfaces

- **Solution:**

```
# On each P router
set interfaces ge-0/0/0 unit 0 family mpls
set protocols mpls interface ge-0/0/0.0
set protocols ldp interface ge-0/0/0.0
```

### Scenario 3: Overlapping Customer Addresses Not Working

- **Symptom:** Two customers using same IP range, routes getting mixed

- **Diagnostic Commands:**

```
# Check Route Distinguisher
user@PE1> show route table bgp.l3vpn.0 10.1.0.0/24 detail
[Shows same RD for different customers]
```

- **Cause:** Same Route Distinguisher used for different VRFs

- **Solution:**

```
# Use unique RD per VRF
set routing-instances CUSTOMER-A route-distinguisher 192.168.1.1:100
set routing-instances CUSTOMER-B route-distinguisher 192.168.1.1:200
```

## Module 3: L3VPN Operational Characteristics - The Inner Workings

### Part 1: The Conceptual Lecture (The Why)

#### The Control Plane: How Routes Are Distributed

The L3VPN control plane is like a sophisticated postal system where each post office (PE router) needs to know how to reach every address in the network, but only delivers mail to its local customers.

#### The Journey of a Route:

```

Step 1: CE advertises route to PE1
Customer A           PE1
[10.2.0.0/24] ----BGP/OSPF--> [Learns in VRF]

Step 2: PE1 converts to VPN-IPv4
PE1 VRF Table        PE1 BGP VPN Table
[10.2.0.0/24] ---Convert--> [192.168.1.1:100:10.2.0.0/24]
                           + Route Target: 65000:100
                           + MPLS Label: 299792

Step 3: PE1 advertises to other PEs via MP-BGP
PE1 -----MP-BGP-----> PE2, PE3
                           [VPN-IPv4 + RT + Label]

Step 4: Remote PEs import based on Route Target
PE2 BGP VPN Table      PE2 VRF Table
[Matches RT 65000:100] --> [Installs in Customer A VRF]

```

#### Key Control Plane Components:

1. **Multi-Protocol BGP (MP-BGP)**
  - Extension of BGP to carry VPN routes
  - Uses Address Family Identifier (AFI) = 1, Subsequent AFI (SAFI) = 128
  - Carries: VPN-IPv4 prefix + Route Target + MPLS label
2. **BGP Extended Communities**
  - Route Target (RT): Controls route distribution
  - Route Origin: Identifies the originating VRF
  - Additional attributes for QoS, bandwidth, etc.
3. **Label Allocation**
  - Each PE allocates a label for each VRF
  - Label identifies which VRF to use for incoming packets
  - Can be per-VRF (one label) or per-prefix (many labels)

#### The Data Plane: How Traffic Flows

The data plane uses a two-label stack to transport packets across the MPLS network:

```

Packet Flow from CE1 (Site A) to CE2 (Site B):

1. CE1 sends plain IP packet
   [IP Header][Data] --> PE1

2. PE1 adds VPN label (identifies VRF at PE2) and transport label
   [Transport Label][VPN Label][IP Header][Data] --> P1

3. P routers swap transport label (ignore VPN label)
   [New Transport Label][VPN Label][IP Header][Data] --> P2

4. Penultimate hop pops transport label
   [VPN Label][IP Header][Data] --> PE2

5. PE2 uses VPN label to identify VRF, forwards to CE2
   [IP Header][Data] --> CE2

```

#### Label Stack Visualization:

Position in Stack	Purpose	Who Processes
Top Label	Transport to PE	P routers (LDP/RSVP)
Bottom Label	Identify VRF/prefix	Egress PE only
Original IP	Customer data	CE routers only

#### Why Two Labels?

- **Scalability:** P routers only need to know about PE routers, not millions of customer routes
- **Privacy:** P routers never see customer IP addresses

- **Flexibility:** Can change transport without affecting VPN service

## PHP (Penultimate Hop Popping)

The second-to-last router removes the transport label to reduce processing at the PE:

Without PHP: P2 --> PE2 [Label1] [Label2] [IP] PE2: Pop Label1, lookup Label2	With PHP: P2 --> PE2 [Label2] [IP] PE2: Lookup Label2 (faster!)
--	--

## Part 2: The Junos CLI Masterclass (The How)

### Advanced L3VPN Control Plane Configuration

```
[edit]
# Configure MP-BGP with route reflection
set protocols bgp {
    group RR-CLIENTS {
        type internal;
        local-address 192.168.1.1;
        family inet-vpn {
            unicast {
                # Add-path for multiple paths
                add-path {
                    receive;
                    send {
                        path-count 2;
                    }
                }
            }
        }
        cluster 192.168.1.1;          # Route reflector cluster ID
        neighbor 192.168.1.2;
        neighbor 192.168.1.3;
    }
}

# Configure per-prefix label allocation (more granular)
set routing-instances CUSTOMER-A {
    instance-type vrf;
    interface ge-0/0/2.0;
    route-distinguisher 192.168.1.1:100;
    vrf-target target:65000:100;
    # Don't use vrf-table-label for per-prefix
    protocols {
        bgp {
            group CE {
                type external;
                peer-as 65001;
                neighbor 10.1.1.1 {
                    family inet {
                        unicast {
                            # Allocate label per prefix
                            prefix-limit {
                                maximum 1000;
                                teardown 80;
                            }
                        }
                    }
                }
            }
        }
    }
}

# Configure BGP signaled LSPs (BGP-LU)
set protocols bgp {
    group IBGP-LU {
        type internal;
        local-address 192.168.1.1;
        family inet {
            labeled-unicast {
                rib {
                    inet.3;           # Install in inet.3 for resolution
                }
            }
        }
    }
}
```

```

        }
    }
    neighbor 192.168.1.2;
}
}

# Advanced VRF import/export with communities
set policy-options {
    policy-statement CUSTOMER-A-EXPORT {
        term LOCAL-ROUTES {
            from {
                protocol [ direct local static ];
                route-filter 10.1.0.0/16 orlonger;
            }
            then {
                community add RT-CUSTOMER-A;
                community add LOCATION-SITE1;      # Additional marking
                accept;
            }
        }
        term BGP-ROUTES {
            from {
                protocol bgp;
                as-path CUSTOMER-AS;
            }
            then {
                community add RT-CUSTOMER-A;
                local-preference 200;             # Prefer this site
                accept;
            }
        }
    }
}

community RT-CUSTOMER-A members target:65000:100;
community LOCATION-SITE1 members origin:65000:1001;
as-path CUSTOMER-AS "^65001$";
}

# Configure constrained route distribution (route target filtering)
set protocols bgp {
    group IBGP-VPN {
        family route-target;          # Enable RT constraint
        neighbor 192.168.1.2;
    }
}

```

## Data Plane Configuration and Optimization

```

[edit]
# Configure MPLS with traffic engineering
set protocols mpls {
    # Administrative groups for path selection
    admin-groups {
        GOLD 0;
        SILVER 1;
        BRONZE 2;
    }

    label-switched-path T0-PE2-PREMIUM {
        to 192.168.1.2;
        bandwidth 100m;           # Reserve bandwidth
        priority 1 1;             # Setup/hold priority
        admin-group include-any GOLD;
        primary GOLD-PATH;
    }

    path GOLD-PATH {
        10.0.0.1 strict;         # High-capacity links
        10.0.0.5 strict;
    }

    # Enable CSPF (Constrained Shortest Path First)
    traffic-engineering bgp;
}

```

```

# Configure LDP with targeted sessions
set protocols ldp {
    interface ge-0/0/0.0;
    interface ge-0/0/1.0;

    # Targeted LDP for L3VPN
    targeted-hello {
        hello-interval 5;
        hold-time 15;
    }

    session 192.168.1.2 {           # Remote PE
        authentication-key "$9$aes256$encrypted-key$";
    }

    # FEC deaggregation for better convergence
    deaggregate;

    # Session protection for faster recovery
    session-protection;
}

# Configure hash-based ECMP for load balancing
set forwarding-options {
    hash-key {
        family mpls {
            label-1;          # Include first label
            label-2;          # Include second label
            payload {
                ip {
                    layer-3;      # Include IP header
                    layer-4;      # Include TCP/UDP ports
                }
            }
        }
    }
}

# Quality of Service for L3VPN traffic
set class-of-service {
    interfaces ge-0/0/0 {
        scheduler-map MPLS-QOS;
        unit 0 {
            classifiers {
                exp MPLS-CLASSIFIER;
            }
        }
    }

    schedulers {
        PREMIUM-DATA {
            transmit-rate percent 40;
            buffer-size percent 40;
            priority high;
        }
        STANDARD-DATA {
            transmit-rate percent 30;
            buffer-size percent 30;
            priority medium-high;
        }
    }

    scheduler-maps MPLS-QOS {
        forwarding-class premium scheduler PREMIUM-DATA;
        forwarding-class standard scheduler STANDARD-DATA;
    }
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Advanced Control Plane Verification

```

# 1. Verify VPN route details with label information
user@PE1> show route table bgp.l3vpn.0 extensive

```

```

bgp.l3vpn.0: 10 destinations, 10 routes (10 active, 0 holddown, 0 hidden)
192.168.1.1:100:10.2.0.0/24 (1 entry, 1 announced)
TSI:
Page 0 idx 0, (group IBGP-VPN type Internal) Type 1 val 0x9c4f5f0 (adv_entry)
    Advertised metrics:
        Nexthop: 192.168.1.1
        MED: 0
        Localpref: 100
        AS path: [65000] 65001 I
        Communities: target:65000:100 origin:65000:1001
        VPN Label: 299792

# 2. Check label allocation and usage
user@PE1> show route forwarding-table family mpls
Routing table: default.mpls
MPLS:
Destination      Type RtRef Next hop          Type Index NhRef Netif
299792           user   0                 indr  1048574   2
                           10.1.1.1       Push  606     2 ge-0/0/2.0

# 3. Verify MP-BGP session details
user@PE1> show bgp neighbor 192.168.1.2 detail
Peer: 192.168.1.2+179 AS 65000 Local: 192.168.1.1+65123 AS 65000
    Group: IBGP-VPN          Routing-Instance: master
    Forwarding routing-instance: master
    Type: Internal      State: Established   Flags: <Sync>
    Last State: OpenConfirm  Last Event: RecvKeepAlive
    Last Error: None
    Options: <Preference LocalAddress AddressFamily Multipath Refresh>
    Address families configured: inet-vpn-unicast route-target
    Local Address: 192.168.1.1 Holdtime: 90 Preference: 170
    NLRI for restart configured on peer: inet-vpn-unicast route-target
    NLRI advertised by peer: inet-vpn-unicast route-target
    NLRI for this session: inet-vpn-unicast route-target
    Peer supports Refresh capability (2)
    Stale routes from peer are kept for: 300
Table bgp.l3vpn.0 Bit: 30000
    RIB State: BGP restart is complete
    Send state: in sync
    Active prefixes: 5
    Received prefixes: 5
    Accepted prefixes: 5
    Suppressed due to damping: 0
    Advertised prefixes: 5

# 4. Trace route through MPLS network
user@PE1> traceroute routing-instance CUSTOMER-A 10.2.0.1 source 10.1.0.1
traceroute to 10.2.0.1 (10.2.0.1) from 10.1.0.1, 30 hops max, 52 byte packets
1  10.1.1.1 (10.1.1.1)  1.123 ms  0.891 ms  0.765 ms
    MPLS Label=299840 CoS=0 TTL=1 S=0
    MPLS Label=299792 CoS=0 TTL=1 S=1
2  10.0.0.2 (10.0.0.2)  2.234 ms  2.011 ms  1.987 ms
    MPLS Label=299824 CoS=0 TTL=1 S=0
    MPLS Label=299792 CoS=0 TTL=2 S=1
3  10.0.0.6 (10.0.0.6)  3.456 ms  3.234 ms  3.111 ms
    MPLS Label=299792 CoS=0 TTL=1 S=1
4  10.2.0.1 (10.2.0.1)  4.567 ms  4.321 ms  4.234 ms

```

## Data Plane Verification

```

# 1. Check MPLS label operations
user@PE1> show mpls lsp ingress extensive
Ingress LSP: 1 sessions

192.168.1.2
    From: 192.168.1.1, State: Up, ActiveRoute: 0, LSPname: T0-PE2-PREMIUM
    ActivePath: GOLD-PATH (primary)
    LSPtype: Static Configured, Penultimate hop popping
    LoadBalance: Random
    Encoding type: Packet, Switching type: Packet, GPID: IPv4
    *Primary GOLD-PATH      State: Up
        Priorities: 1 1
        Bandwidth: 100Mbps
        SmartOptimizeTimer: 180
    Received RRO (ProtectionFlag 1=Available 2=InUse 4=B/W 8=Node 10=SoftPreempt 20=Node-ID):

```

```

10.0.0.1 10.0.0.2 10.0.0.5 10.0.0.6
Recorded routes:
10.0.0.1 10.0.0.2 10.0.0.5 10.0.0.6
Total 1 displayed, Up 1, Down 0

# 2. Monitor label distribution
user@PE1> show ldp database
Input label database, 192.168.1.1:0--192.168.1.2:0
Labels received: 15
  Label      Prefix
299840    192.168.1.2/32
299856    10.0.0.0/30

Output label database, 192.168.1.1:0--192.168.1.2:0
Labels advertised: 14
  Label      Prefix
299792    192.168.1.1/32
299808    10.0.0.4/30

# 3. Check CoS markings
user@PE1> show class-of-service interface ge-0/0/0
Physical interface: ge-0/0/0, Enabled
  Scheduler map: MPLS-QOS, Index: 45123
  Congestion-notification: Disabled
Logical interface ge-0/0/0.0
  Object          Name          Type          Index
  Classifier      MPLS-CLASSIFIER exp          12345

```

## Common Troubleshooting Scenarios

### Scenario 1: MP-BGP Session Up but No VPN Routes

- **Symptom:** BGP established but no VPN routes exchanged
- **Diagnostic Commands:**

```

user@PE1> show bgp neighbor 192.168.1.2 | match "NLRI"
  NLRI for restart configured on peer: inet-unicast
  NLRI advertised by peer: inet-unicast
  # Missing inet-vpn-unicast!

user@PE1> show configuration protocols bgp group IBGP-VPN
# Missing family inet-vpn unicast

```

- **Cause:** Address family not negotiated
- **Solution:**

```

set protocols bgp group IBGP-VPN family inet-vpn unicast
commit
# Session will reset and renegotiate capabilities

```

### Scenario 2: Suboptimal Routing (Traffic Through Wrong Path)

- **Symptom:** Traffic takes longer path than expected
- **Diagnostic Commands:**

```

user@PE1> show route forwarding-table destination 10.2.0.0/24 table CUSTOMER-A
# Shows unexpected next-hop

user@PE1> show route protocol bgp 10.2.0.0/24 detail table CUSTOMER-A.inet.0
# Multiple paths, wrong one selected
  *BGP   Preference: 170/-201
        Next hop: 10.0.0.10 via ge-0/0/1.0  # Suboptimal
        Local preference: 100                 # Default
  BGP   Preference: 170/-101
        Next hop: 10.0.0.2 via ge-0/0/0.0  # Better path
        Local preference: 100
        AS path: 65001 I (shorter)

```

- **Cause:** Local preference or MED affecting selection
- **Solution:**

```

# Adjust import policy to prefer certain paths
set policy-options policy-statement CUSTOMER-A-IMPORT {
    term PREFER-DIRECT {
        from {
            protocol bgp;
            as-path SHORTER;
        }
        then {
            local-preference 150;
            accept;
        }
    }
}
set policy-options as-path SHORTER ".{0,2}"; # Paths with ≤2 AS hops

```

### Scenario 3: MPLS Label Stack Issues

- **Symptom:** Packets dropped at P router
- **Diagnostic Commands:**

```

# On P router
user@P1> show mpls lsp transit
No transit LSPs # Problem!

user@P1> show ldp neighbor
No LDP neighbors # LDP not running

user@P1> show configuration protocols
# Missing MPLS/LDP configuration

```

- **Cause:** MPLS/LDP not configured on transit router
- **Solution:**

```

# On P router
set interfaces ge-0/0/0 unit 0 family mpls
set interfaces ge-0/0/1 unit 0 family mpls
set protocols mpls interface all
set protocols ldp interface all
commit

```

### Scenario 4: Route Target Mismatch

- **Symptom:** Sites can't communicate despite correct configuration
- **Diagnostic Commands:**

```

user@PE1> show route advertising-protocol bgp 192.168.1.2 detail table bgp.l3vpn
192.168.1.1:100:10.1.0.0/24 (1 entry, 1 announced)
    Communities: target:65000:100

user@PE2> show configuration routing-instances CUSTOMER-A
vrf-target target:65000:200; # Different RT!

```

- **Cause:** Mismatched Route Targets
- **Solution:**

```

# Ensure matching RTs
set routing-instances CUSTOMER-A vrf-target target:65000:100
# Or use import/export for complex scenarios
set routing-instances CUSTOMER-A vrf-import [ IMPORT-100 IMPORT-200 ]
set routing-instances CUSTOMER-A vrf-export EXPORT-100

```

## Summary and Key Takeaways

### Architectural Principles

#### 1. Separation of Concerns

- CE routers: Simple IP routing
- PE routers: VPN intelligence

- P routers: Fast label switching
- Scalability Through Hierarchy**
    - Two-label stack separates transport from VPN
    - P routers don't need customer routes
    - Route Reflectors reduce BGP sessions
  - Flexibility Through Abstraction**
    - Route Distinguishers enable overlapping addresses
    - Route Targets enable any-to-any connectivity
    - Multiple transport options (LDP, RSVP, SR)

## Configuration Best Practices

1. **Always use unique Route Distinguishers** per VRF per PE
2. **Implement Route Target constraints** to reduce unnecessary routes
3. **Use per-VRF labels** for simplicity, per-prefix for granularity
4. **Enable authentication** on BGP and LDP sessions
5. **Configure BFD** for faster failure detection
6. **Use communities** for policy control and traffic engineering

## Troubleshooting Methodology

1. **Verify control plane first** (BGP, route distribution)
2. **Check data plane second** (MPLS labels, forwarding)
3. **Use traceroute** to verify label operations
4. **Compare working vs non-working sites**
5. **Check all routers in path** (CE, PE, P)

# JNCIE-SP Learning Module: Layer 3 VPN Configuration, Verification, and OSPF PE-CE Protocol

## Module 4: Layer 3 VPN Configuration

### Part 1: The Conceptual Lecture (The Why)


#### The Fundamental Problem

Imagine you're a large corporation with offices in New York, London, and Tokyo. Each office has its own network with computers, servers, and printers. You want these offices to communicate as if they were all in the same building, but they're separated by thousands of miles. The traditional solution would be to lease expensive dedicated physical lines between each office, but this is costly and inflexible.

**Enter MPLS Layer 3 VPNs:** A technology that allows a service provider to create isolated, private networks for multiple customers over a shared infrastructure. Think of it like having private, secure tunnels through a public highway system.

#### Core Concepts of L3VPN Configuration

An L3VPN consists of several key components that must be configured:



#### Key Elements to Configure:

1. **VRF (Virtual Routing and Forwarding) Instance:** A virtual router that maintains separate routing tables for each customer
2. **Route Distinguisher (RD):** Makes customer routes globally unique (64 added to IPv4 routes)
3. **Route Target (RT):** Controls which routes are imported/exported between VRFs
4. **PE-CE Interface:** Physical connection between provider edge and customer edge

5. **PE-CE Routing Protocol:** How routes are exchanged (BGP, OSPF, static, etc.)

## The VPN-IPv4 Address Structure

When a PE router receives a regular IPv4 route from a CE, it transforms it:

```
Original Customer Route: 10.1.1.0/24
                         ↓
Route Distinguisher:   65000:100 (ASN:Value format)
                         +
IPv4 Prefix:          10.1.1.0/24
                         =
VPN-IPv4 Route:       65000:100:10.1.1.0/24
```

This 96 address ensures that even if two customers use the same private IP space, their routes remain unique in the provider's network.

## Part 2: The Junos CLI Masterclass (The How)

### Configuration Hierarchy Overview

L3VPN configuration in Junos touches multiple hierarchy levels:

```
[edit]
└── interfaces      # PE-CE interface configuration
└── routing-instances # VRF configuration
└── protocols
    └── bgp        # MP-BGP for VPN route distribution
    └── mpls       # MPLS transport
└── policy-options # Route import/export policies
```

### Step-by-Step L3VPN PE Configuration

#### Step 1: Configure the PE-CE Interface

```
[edit interfaces]
ge-0/0/0 {
    description "To Customer A - CE1";
    unit 0 {
        family inet {
            address 192.168.1.1/30; # PE side of PE-CE link
        }
    }
}
```

#### Step 2: Create the VRF (Routing Instance)

```
[edit routing-instances]
CUSTOMER-A {
    instance-type vrf;           # Defines this as a VRF
    interface ge-0/0/0.0;         # Assigns PE-CE interface to VRF
    route-distinguisher 65000:100; # Makes routes globally unique
    vrf-target target:65000:100;  # RT for import/export (simplified)
    vrf-table-label;             # Allocates label for VRF table
}
```

**Understanding vrf-target:** This is a shorthand that creates both import and export policies:

- **vrf-import:** Which routes to accept into this VRF
- **vrf-export:** Which routes to advertise from this VRF

#### Step 3: Configure PE-CE Routing (Static Example)

```
[edit routing-instances CUSTOMER-A]
routing-options {
    static {
        route 10.1.0.0/24 next-hop 192.168.1.2; # Customer LAN via CE
    }
}
```

#### Step 4: Configure MP-BGP for VPN Route Distribution

```
[edit protocols bgp]
group IBGP-RR {
    type internal;
    local-address 1.1.1.1;          # PE loopback
    family inet-vpn {              # Enables VPN-IPv4 AFI/SAFI
        unicast;
    }
    neighbor 2.2.2.2;             # Route Reflector or remote PE
}
```

#### Step 5: Enable MPLS on Core Interfaces

```
[edit protocols mpls]
interface ge-0/0/1.0; # Core-facing interface
interface lo0.0;      # Loopback

[edit protocols ldp]
interface ge-0/0/1.0; # LDP for label distribution
interface lo0.0;
```

#### Complete Annotated Configuration Block

```
## Complete L3VPN PE Configuration for Customer A ##

interfaces {
    ge-0/0/0 {
        description "PE-CE link to Customer A Site 1";
        unit 0 {
            family inet {
                address 192.168.1.1/30;
            }
        }
    }
    ge-0/0/1 {
        description "PE-P core link";
        unit 0 {
            family inet {
                address 10.0.0.1/30;
            }
            family mpls;           # Enable MPLS encapsulation
        }
    }
    lo0 {
        unit 0 {
            family inet {
                address 1.1.1.1/32;     # PE loopback for BGP
            }
        }
    }
}

routing-instances {
    CUSTOMER-A {
        instance-type vrf;
        interface ge-0/0/0.0;
        route-distinguisher 65000:100;      # Format: ASN:CustomerID
        vrf-target target:65000:100;         # Import and export same RT
        vrf-table-label;                   # Enables single label per VRF
        routing-options {
            static {
                route 10.1.0.0/24 next-hop 192.168.1.2;
            }
        }
    }
}

protocols {
    mpls {
        interface ge-0/0/1.0;
        interface lo0.0;
    }
    ldp {
        interface ge-0/0/1.0;
        interface lo0.0;
    }
}
```

```

}
bgp {
    group IBGP-VPN {
        type internal;
        local-address 1.1.1.1;
        family inet-vpn {
            unicast;
        }
        neighbor 2.2.2.2 {
            description "To PE2";
        }
    }
    ospf {
        area 0.0.0.0 # IGP for PE loopback reachability
        interface ge-0/0/1.0;
        interface lo0.0;
    }
}
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### 1. Verify VRF Configuration

```

user@PE1> show route instance CUSTOMER-A detail
CUSTOMER-A:
  Router ID: 192.168.1.1
  Type: vrf          State: Active
  Interfaces:
    ge-0/0/0.0
  Route-distinguisher: 65000:100
  Vrf-import: [ __vrf-import-CUSTOMER-A-internal__ ]
  Vrf-export: [ __vrf-export-CUSTOMER-A-internal__ ]
  Vrf-import-target: [ target:65000:100 ]
  Vrf-export-target: [ target:65000:100 ]
  Fast-reroute-priority: low
  Tables:
    CUSTOMER-A.inet.0      : 3 routes (3 active, 0 holddown, 0 hidden)

```

#### 2. Verify Routes in VRF

```

user@PE1> show route table CUSTOMER-A.inet.0

CUSTOMER-A.inet.0: 3 destinations, 3 routes (3 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

10.1.0.0/24      *[Static/5] 00:10:32
                  > to 192.168.1.2 via ge-0/0/0.0
192.168.1.0/30   *[Direct/0] 00:10:32
                  > via ge-0/0/0.0
192.168.1.1/32   *[Local/0] 00:10:32
                  Local via ge-0/0/0.0

```

### Common Troubleshooting Scenarios

#### Scenario 1: Routes Not Being Exported from VRF

*Symptom:* Local CE routes visible in VRF but not advertised to remote PEs

*Diagnostic Commands:*

```

user@PE1> show route advertising-protocol bgp 2.2.2.2 table bgp.l3vpn.0
# No output or missing routes

```

*Cause:* Missing or incorrect vrf-target configuration

*Solution:*

```

[edit routing-instances CUSTOMER-A]
set vrf-target target:65000:100

```

## Scenario 2: Remote VPN Routes Not Installing in VRF

*Symptom:* Can see remote routes in bgp.l3vpn.0 but not in VRF table

*Diagnostic Commands:*

```
user@PE1> show route table bgp.l3vpn.0 65000:100:10.2.0.0/24
# Route exists

user@PE1> show route table CUSTOMER-A.inet.0 10.2.0.0/24
# No route found
```

*Cause:* Route target mismatch - import target doesn't match export target of remote PE

*Solution:*

```
[edit routing-instances CUSTOMER-A]
set vrf-target target:65000:100    # Ensure this matches remote PE's export
```

## Scenario 3: CE Cannot Reach Remote CE

*Symptom:* Ping from CE1 to CE2 fails

*Diagnostic Commands:*

```
user@PE1> show route table CUSTOMER-A.inet.0 10.2.0.0/24
# Route exists with correct next-hop

user@PE1> show route forwarding-table vpn CUSTOMER-A
# Check MPLS label is assigned

user@PE1> show ldp neighbor
# Verify LDP sessions are up
```

*Cause:* MPLS transport path broken

*Solution:*

```
# Ensure MPLS/LDP enabled on all core interfaces
[edit protocols mpls]
set interface ge-0/0/1.0
[edit protocols ldp]
set interface ge-0/0/1.0
```

---

## Module 5: Layer 3 VPN Verification

### Part 1: The Conceptual Lecture (The Why)

#### Why Verification Matters

Building an L3VPN is like constructing a complex highway system with multiple levels of tunnels and interchanges. Each component must work perfectly for traffic to flow. Verification ensures:

1. **Control Plane Health:** Routes are being learned and distributed correctly
2. **Data Plane Integrity:** Packets are being forwarded with correct labels
3. **End-to-End Connectivity:** Customer sites can communicate

#### The L3VPN Verification Framework

Verification Layers:

Application Layer (Customer Traffic)	← Can CE1 ping CE2?
VRF Routing (PE VRF Tables)	← Are routes in VRF?
BGP VPNv4 (Route Distribution)	← Are VPN routes exchanged?
MPLS Transport (Labels)	← Is label path established?

## Part 2: The Junos CLI Masterclass (The How)

### Systematic Verification Approach

#### Level 1: IGP and MPLS Transport Verification

```
# Verify PE loopback reachability via IGP
user@PE1> show route 2.2.2.2/32

inet.0: 10 destinations, 10 routes (10 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

2.2.2.2/32      *[OSPF/10] 00:05:23, metric 2
                 > to 10.0.0.2 via ge-0/0/1.0

# Verify MPLS LSP exists
user@PE1> show route 2.2.2.2/32 table inet.3

inet.3: 2 destinations, 2 routes (2 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

2.2.2.2/32      *[LDP/9] 00:05:23, metric 1
                 > to 10.0.0.2 via ge-0/0/1.0, Push 299776

# Verify LDP neighbors
user@PE1> show ldp neighbor
Address          Interface          Label space ID        Hold time
10.0.0.2         ge-0/0/1.0       2.2.2.2:0           13
```

#### Level 2: BGP VPNv4 Verification

```
# Verify BGP session with family inet-vpn
user@PE1> show bgp summary
Threading mode: BGP I/O
Groups: 1 Peers: 1 Down peers: 0
Table          Tot Paths  Act Paths Suppressed   History Damp State    Pending
bgp.l3vpn.0
                  2          2          0          0          0          0
Peer          AS        InPkt     OutPkt   OutQ   Flaps Last Up/Dwn State
2.2.2.2       65000     453      451       0       0      3:23:45 Establ
bgp.l3vpn.0: 2/2/2/0

# View VPNv4 routes received
user@PE1> show route receive-protocol bgp 2.2.2.2 table bgp.l3vpn.0

bgp.l3vpn.0: 4 destinations, 4 routes (4 active, 0 holddown, 0 hidden)
Prefix          Nexthop          MED      Lclpref   AS path
65000:100:10.2.0.0/24
*                2.2.2.2          100        I
```

#### Level 3: VRF Table Verification

```
# Check VRF routing table
user@PE1> show route table CUSTOMER-A.inet.0 detail

CUSTOMER-A.inet.0: 5 destinations, 5 routes (5 active, 0 holddown, 0 hidden)
10.2.0.0/24 (1 entry, 1 announced)
  *BGP  Preference: 170/-101
        Route Distinguisher: 65000:100
        Next hop type: Indirect, Index: 1048574
        Address: 0x9d4e8f4
        Next-hop reference count: 2
        Source: 2.2.2.2
        Next hop type: Router, Next hop index: 612
        Next hop: 10.0.0.2 via ge-0/0/1.0, selected
        Label operation: Push 16, Push 299776(top)
        Label TTL action: prop-ttl, prop-ttl(top)
        Load balance label: Label 16: None;
        Label element ptr: 0x9d4e7c8
        Label parent element ptr: 0x9d4e4a8
        Label element references: 1
```

```

Label element child references: 0
Label element lsp id: 0
Session Id: 0x141
Protocol next hop: 2.2.2.2
Composite next hop: 0x8f3c2f4 565 INH Session ID: 0x149
Indirect next hop: 0x9bd9404 1048574 INH Session ID: 0x149
State: <Secondary Active Int Ext>
Local AS: 65000 Peer AS: 65000
Age: 10:23 Metric2: 2
Validation State: unverified
Task: BGP_65000.2.2.2.2
Announcement bits (1): 0-KRT
AS path: I
Communities: target:65000:100
Import Accepted
VPN Label: 16
Localpref: 100
Router ID: 2.2.2.2
Primary Routing Table bgp.l3vpn.0

```

#### Level 4: Data Plane Verification

```

# Verify MPLS forwarding table
user@PE1> show route forwarding-table family mpls
Routing table: default.mpls
MPLS:
Destination      Type RtRef Next hop          Type Index NhRef Netif
default          perm  0
0                user   0
1                user   0
2                user   0
13               user   0
16               user   0
16(S=0)          user   0
299776          user   0 10.0.0.2           Swap 299792 ge-0/0/1.0

# Verify VPN forwarding entry
user@PE1> show route forwarding-table vpn CUSTOMER-A
Routing table: CUSTOMER-A.inet
Internet:
Destination      Type RtRef Next hop          Type Index NhRef Netif
10.1.0.0/24       user   0 192.168.1.2        ucst    587     2 ge-0/0/0.0
10.2.0.0/24       user   0

```

### Part 3: Verification & Troubleshooting (The What-If)

#### Comprehensive Verification Checklist

```

# Quick Health Check Script
show bgp summary | match inet-vpn          # BGP session up?
show route summary table CUSTOMER-A.inet.0 # Routes in VRF?
show ldp neighbor | match "^[0-9]"         # LDP neighbors up?
show mpls lsp summary                      # If using RSVP
ping 10.2.0.1 routing-instance CUSTOMER-A # End-to-end test

```

#### Advanced Troubleshooting Scenarios

##### Scenario 1: Asymmetric Routing in L3VPN

*Symptom:* Traffic from CE1→CE2 works, but CE2→CE1 fails

*Diagnostic Commands:*

```

# On PE1
user@PE1> show route table CUSTOMER-A.inet.0 10.2.0.0/24
# Route exists

# On PE2
user@PE2> show route table CUSTOMER-A.inet.0 10.1.0.0/24
# Route missing or wrong next-hop

# Check route export from PE1
user@PE1> show route advertising-protocol bgp 2.2.2.2 65000:100:10.1.0.0/24 detail

```

**Cause:** Route not being exported from PE1 or wrong RT

**Solution:**

```
# Ensure bidirectional route exchange
[edit routing-instances CUSTOMER-A]
set vrf-target target:65000:100 # Same on both PEs
```

## Scenario 2: MTU Issues in MPLS Network

**Symptom:** Small pings work, large pings fail

**Diagnostic Commands:**

```
user@CE1> ping 10.2.0.1 size 1500 do-not-fragment
# Fails

user@CE1> ping 10.2.0.1 size 1472 do-not-fragment
# Works
```

**Cause:** MPLS adds overhead (typically 8-12 bytes), causing fragmentation

**Solution:**

```
# Increase MTU on all MPLS interfaces
[edit interfaces ge-0/0/1]
set mtu 1514 # Or higher to accommodate MPLS labels
```

---

# Module 6: OSPF as the PE-to-CE Protocol

## Part 1: The Conceptual Lecture (The Why)

### The PE-CE Protocol Challenge

When connecting customer sites through an L3VPN, we need a method to exchange routing information between the customer's equipment (CE) and the provider's equipment (PE). While static routes work for simple scenarios, dynamic routing protocols like OSPF provide:

1. **Automatic route discovery:** No manual configuration for every subnet
2. **Dynamic failover:** Automatic rerouting when links fail
3. **Scalability:** Easier to manage hundreds of routes

### The OSPF-in-VPN Complexity

Running OSPF between PE and CE in an L3VPN environment creates unique challenges:

Traditional OSPF Domain:	OSPF in L3VPN:
R1---R2---R3 \      / R4---R5	CE1--PE1~~[MPLS]~~PE2--CE2                    OSPF     BGP  VPNv4    OSPF
Single OSPF Database	OSPF Domain Fragmented by BGP in the Middle!

**The Fundamental Problem:** OSPF expects a contiguous domain where all routers exchange Link State Advertisements (LSAs). In L3VPN, OSPF domains at different sites are separated by BGP, breaking this continuity.

### OSPF Loop Prevention Mechanisms

Because customer sites might have backdoor connections, we need loop prevention:

1. **DN (Down) Bit:** Prevents Type-3 LSA loops
2. **VPN Route Tag:** Prevents Type-5 LSA loops
3. **Domain ID:** Identifies routes from the same OSPF domain

Loop Scenario Without Protection:  
CE1 ---> PE1 ---> PE2 ---> CE2  
  ^                  |  
  |                  |  
+----Backdoor Connection----+

```
Route advertised: CE1→PE1→PE2→CE2→CE1→PE1 (Loop!)
```

## Part 2: The Junos CLI Masterclass (The How)

### OSPF PE-CE Configuration Structure

The configuration involves three main components:

1. **OSPF instance within VRF**: PE runs separate OSPF process per customer
2. **BGP extended communities**: Preserve OSPF attributes across MPLS
3. **Route redistribution**: OSPF↔BGP conversion at each PE

### Step-by-Step OSPF PE-CE Configuration

#### Step 1: Configure OSPF in the VRF

```
[edit routing-instances CUSTOMER-A]
protocols {
    ospf {
        # Domain ID ensures routes from same OSPF domain are recognized
        domain-id 0.0.0.100;                      # Format: 0.0.0.X or IP address

        # Optional: Set OSPF router-id for this VRF
        router-id 192.168.1.1;

        # Define OSPF area and interface
        area 0.0.0.0 {
            interface ge-0/0/0.0 {                  # PE-CE interface
                interface-type p2p;                 # Point-to-point (no DR election)
                metric 10;                         # Optional: Set OSPF cost
            }
        }

        # Critical: Export BGP routes into OSPF
        export BGP-T0-OSPF;                     # Policy to redistribute BGP→OSPF
    }
}
```

#### Step 2: Configure Route Export Policy (BGP→OSPF)

```
[edit policy-options]
policy-statement BGP-T0-OSPF {
    term BGP-ROUTES {
        from protocol bgp;                   # Routes learned via BGP
        then {
            metric 100;                      # Set OSPF metric
            external {                       # Advertise as Type-5 LSA
                type 2;                      # E2 (metric doesn't increment)
            }
            accept;
        }
    }
    term REJECT {
        then reject;                      # Don't leak other routes
    }
}
```

#### Step 3: Configure BGP Extended Communities

```
[edit routing-instances CUSTOMER-A]
vrf-target target:65000:100;                      # Standard RT for VPN membership

protocols {
    ospf {
        domain-id 0.0.0.100;                  # Critical for loop prevention

        # VPN route-type community automatically added for OSPF routes
        # Format: area:route-type:option
        # This happens automatically when OSPF runs in VRF
    }
}
```

#### Step 4: Understanding Automatic BGP Communities

When PE1 redistributes OSPF routes into BGP, it automatically adds:

```
Route-Type Community: 0x8000:area:route-type:option  
Domain-ID Community: 0x8005:0:domain-id
```

Example for Area 0 Type-3 LSA:

- Route-Type: 0x8000:0:3:0
- Domain-ID: 0x8005:0:0.0.100

#### Step 5: Complete PE Configuration

```
## Complete OSPF PE-CE Configuration ##

interfaces {
    ge-0/0/0 {
        description "PE-CE OSPF Link to Customer A";
        unit 0 {
            family inet {
                address 192.168.1.1/30;
            }
        }
    }
}

routing-instances {
    CUSTOMER-A {
        instance-type vrf;
        interface ge-0/0/0.0;
        route-distinguisher 65000:100;
        vrf-target target:65000:100;
        vrf-table-label;

        protocols {
            ospf {
                domain-id 0.0.0.100;           # Same across all PEs for this customer
                router-id 192.168.1.1;         # Unique per PE

                area 0.0.0.0 {
                    interface ge-0/0/0.0 {
                        interface-type p2p;
                        authentication {          # Optional: OSPF authentication
                            md5 1 key "$9$H.5FCAu"; # Encrypted key
                        }
                    }
                }
                export BGP-T0-OSPF;          # Redistribute remote sites
            }
        }
    }
}

policy-options {
    policy-statement BGP-T0-OSPF {
        term BGP-VPN-ROUTES {
            from {
                protocol bgp;
                community TARGET-65000-100; # Only VPN routes
            }
            then {
                metric 100;
                external {
                    type 2;                  # E2 external routes
                }
                accept;
            }
        }
        term DIRECT-ROUTES {           # Optional: Connected routes
            from protocol direct;
            then {
                metric 10;
                external {
                    type 1;                  # E1 external routes
                }
            }
        }
    }
}
```

```

        }
        accept;
    }
}

community TARGET-65000-100 members target:65000:100;
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential OSPF-VPN Verification Commands

#### 1. Verify OSPF Neighborship

```

user@PE1> show ospf neighbor instance CUSTOMER-A
Address      Interface      State      ID      Pri  Dead
192.168.1.2  ge-0/0/0.0  Full      10.1.1.1  128   38

```

#### 2. Verify OSPF Database in VRF

```

user@PE1> show ospf database instance CUSTOMER-A

OSPF database, Area 0.0.0.0
Type      ID      Adv Rtr      Seq      Age  Opt  Cksum Len
Router    192.168.1.1  192.168.1.1  0x80000003  234  0x22 0x9d5f  48
Router *  10.1.1.1    10.1.1.1    0x80000005  233  0x22 0x5a9c  60
Summary   10.2.0.0    192.168.1.1  0x80000001  234  0x22 0x4d5e  28

OSPF AS SCOPE link state database
Type      ID      Adv Rtr      Seq      Age  Opt  Cksum Len
Extern   10.2.0.0    192.168.1.1  0x80000001  234  0x22 0x3f6b  36

```

#### 3. Check BGP Extended Communities

```

user@PE1> show route table bgp.l3vpn.0 extensive 65000:100:10.1.10.0/24

bgp.l3vpn.0: 4 destinations, 4 routes (4 active, 0 holddown, 0 hidden)
65000:100:10.1.10.0/24 (1 entry, 0 announced)
  *BGP    Preference: 170/-101
  Route Distinguisher: 65000:100
  Next hop type: Indirect
  Source: 1.1.1.1
  Protocol next hop: 1.1.1.1
  Communities: target:65000:100
                domain-id:0.0.0.100      # Domain ID preserved
                rt-type:0.0.0.0:1:0       # Area 0, Type 1 LSA
  Import Accepted
  VPN Label: 16
  Localpref: 100
  Router ID: 1.1.1.1

```

## Common OSPF-VPN Troubleshooting Scenarios

### Scenario 1: OSPF Routes Not Redistributed into BGP

*Symptom:* CE routes visible in PE's OSPF database but not in bgp.l3vpn.0

*Diagnostic Commands:*

```

user@PE1> show ospf route instance CUSTOMER-A
Topology default Route Table:
Prefix      Path  Route      NH      Metric  NextHop      Nexthop
          Type  Type      Type           Interface      Address/LSP
10.1.10.0/24  Intra Network  IP      11  ge-0/0/0.0  192.168.1.2

user@PE1> show route advertising-protocol bgp 2.2.2.2 65000:100:10.1.10.0/24
# No output

```

*Cause:* VRF missing proper export configuration

*Solution:*

```
[edit routing-instances CUSTOMER-A]
set vrf-table-label          # Required for OSPF routes
set protocols ospf domain-id 0.0.0.100  # Required for proper community
```

### Scenario 2: Routes Received as External Instead of Inter-Area

*Symptom:* Routes from remote site appear as O E2 instead of O IA

*Diagnostic Commands:*

```
user@CE1> show route protocol ospf
10.2.0.0/24      *[OSPF/150] 00:05:12, metric 100, tag 3489660929
                  > to 192.168.1.1 via ge-0/0/0.0

# Tag 3489660929 = 0xD0000001 (VPN route tag)
```

*Cause:* Domain-ID mismatch between PEs

*Solution:*

```
# Ensure same domain-id on all PEs for this VPN
[edit routing-instances CUSTOMER-A protocols ospf]
set domain-id 0.0.0.100  # Must match on PE1 and PE2
```

### Scenario 3: DN Bit Preventing Route Installation

*Symptom:* Type-3 LSA received but not installed in routing table

*Diagnostic Commands:*

```
user@PE1> show ospf database lsa-id 10.2.0.0 detail instance CUSTOMER-A

Summary  10.2.0.0      192.168.1.1      0x80000001  234  0x22 0x4d5e  28
mask 255.255.255.0
Topology default (ID 0)
Type: Transit, Node ID: 2.2.2.2
Metric: 100, Bidirectional
Options: 0x22 (*|*|DC|*|*|*|E|*)
          ^
          DN bit set
```

*Cause:* DN bit prevents PE from using LSA (loop prevention working correctly)

*Solution:*

```
# This is expected behavior - no action needed
# DN bit prevents loops when customer has backdoor links
# PE ignores Type-3 LSAs with DN bit set
```

### Scenario 4: Suboptimal Routing via Backdoor Link

*Symptom:* Traffic prefers slower backdoor link over MPLS VPN

*Diagnostic Commands:*

```
user@CE1> show route 10.2.0.0/24
10.2.0.0/24  *[OSPF/10] via backdoor-link    # Cost 50
              [OSPF/150] via PE-link        # Cost 100 (External)
```

*Cause:* OSPF prefers intra-area/inter-area routes over external

*Solution - Use Sham Links:*

```
# Create loopback in VRF for sham-link endpoint
[edit routing-instances CUSTOMER-A]
set interface lo0.100 family inet address 100.100.100.1/32

# Configure sham-link to remote PE
set protocols ospf area 0.0.0.0 sham-link local 100.100.100.1
set protocols ospf area 0.0.0.0 sham-link remote 100.100.100.2 metric 10
```

```
# Export sham-link endpoint to BGP
set routing-options static route 100.100.100.1/32 discard
```

## Complete Verification Flow for OSPF PE-CE

```
# 1. Check OSPF neighbor
show ospf neighbor instance CUSTOMER-A

# 2. Verify OSPF routes learned from CE
show ospf route instance CUSTOMER-A

# 3. Confirm routes in VRF table
show route table CUSTOMER-A.inet.0 protocol ospf

# 4. Verify routes advertised to BGP
show route advertising-protocol bgp 2.2.2.2 table bgp.l3vpn.0

# 5. Check remote routes received via BGP
show route receive-protocol bgp 2.2.2.2 table bgp.l3vpn.0

# 6. Verify routes advertised to CE via OSPF
show ospf database advertising-router 192.168.1.1 instance CUSTOMER-A

# 7. End-to-end connectivity test
ping 10.2.0.1 routing-instance CUSTOMER-A source 192.168.1.1
```

## Summary

These three modules have covered the complete lifecycle of L3VPN implementation:

1. **Configuration:** Building VRFs, RD/RT, PE-CE interfaces, and MP-BGP
2. **Verification:** Systematic checking from IGP/MPLS up through VPN routes
3. **OSPF Integration:** Dynamic routing with loop prevention mechanisms

The key to mastering L3VPNs is understanding the interaction between:

- **Local routing** (PE-CE via static/OSPF/BGP)
- **VPN route distribution** (MP-BGP with RD/RT)
- **MPLS forwarding** (Label stack for transport)

Each layer must function correctly for the complete solution to work. Practice these configurations in lab environments, and always verify each layer systematically when troubleshooting.

## Module 7: OSPF—Optimal Routing in L3VPNs

### Part 1: The Conceptual Lecture (The Why)

#### The Fundamental Problem

Imagine you're managing a large company with multiple branch offices. Each office has its own local network (think of it as a neighborhood), and you've hired a service provider to connect all these offices together using MPLS L3VPN technology (like a private highway system). Now, some of your offices decide to create their own direct connections between each other for backup purposes - these are called "backdoor links" (like building your own private roads between offices).

When you use OSPF (Open Shortest Path First) as the routing protocol between your offices and the service provider, three critical problems emerge:

1. **Routing Loops:** Information about networks can circulate endlessly
2. **Suboptimal Routing:** Traffic might take the slower backdoor link instead of the fast MPLS highway
3. **Control Plane Confusion:** The same route information appears from multiple sources

#### Understanding the DN Bit (Down Bit)

The DN bit is a special flag in OSPF LSA Type 3 (Summary LSAs) that acts like a "do not re-advertise" stamp. Think of it as a postal stamp that says "delivered by MPLS - do not forward back."



Without DN Bit:

1. CE1 advertises 10.1.1.0/24 to PE1
2. PE1 redistributes into MP-BGP
3. PE2 receives via MP-BGP and redistributes to OSPF
4. PE2 sends to CE2 as Type 3 LSA
5. CE2 forwards this LSA to CE1 via backdoor
6. CE1 sees its own route coming back! (LOOP)

With DN Bit:

- 1-4. Same as above
5. PE2 sets DN bit when sending Type 3 LSA to CE2
6. CE2 forwards LSA to CE1 via backdoor
7. CE1 sees DN bit set and IGNORES the LSA

## Understanding the VPN Route Tag

The VPN Route Tag is a 32-bit value attached to OSPF External routes (Type 5 LSAs) that serves as a unique identifier, like a barcode on a package. It prevents the same route from being re-imported into BGP.

The tag structure encodes the BGP Autonomous System Number (ASN):

- For 2-byte ASN: Tag = 0xD000 + ASN (example: AS 65001 → Tag 0xD0FDE9)
- For 4-byte ASN: Tag = 0xD800 + lower 16 bits of ASN

## Understanding Sham Links

A sham link is a virtual OSPF adjacency across the MPLS backbone that makes the MPLS path appear as an intra-area link. Think of it as creating a "virtual tunnel" that tricks OSPF into believing the MPLS path is a direct connection.

Without Sham Link:

```
CE1 ---Area 0--- CE2 (Backdoor = Intra-Area route, Cost 10)
|           |
PE1 ===MPLS==== PE2 (MPLS = Inter-Area route, Cost 5)
```

OSPF Preference: Intra-Area > Inter-Area

Result: Traffic uses backdoor despite higher cost!

With Sham Link:

```
CE1 ---Area 0--- CE2 (Backdoor = Intra-Area, Cost 10)
|           |
PE1 ===SHAM===== PE2 (MPLS = Now Intra-Area, Cost 5)
```

Result: Traffic uses MPLS path (lower cost wins)

## Part 2: The Junos CLI Masterclass (The How)

### OSPF PE-CE Configuration with Loop Prevention

First, let's establish the baseline OSPF configuration hierarchy:

```
[edit routing-instances CUSTOMER-A]
instance-type vrf;
interface ge-0/0/0.100;    # Interface to CE
route-distinguisher 65000:100;
vrf-target target:65000:100;
vrf-table-label;

protocols {
    ospf {
        # DN bit is automatically handled by Junos for Type 3 LSAs
        # VPN route tag is automatically generated for Type 5 LSAs

        export bgp-to-ospf;    # Policy to redistribute BGP into OSPF
        area 0.0.0.0 {
            interface ge-0/0/0.100;
        }
    }
}
```

### Configuring Domain ID for Optimal Route Type Selection

```
[edit routing-instances CUSTOMER-A protocols ospf]
domain-id 0.0.0.1;      # Same domain-id across all PEs for this VPN

# Why domain-id matters:
# - Same domain-id = Routes appear as Type 3 (Inter-area)
# - Different domain-id = Routes appear as Type 5 (External)
# - Type 3 preferred over Type 5 in OSPF
```

## Configuring Sham Links

Step 1: Create loopback interfaces for sham link endpoints:

```
# On PE1
[edit routing-instances CUSTOMER-A]
interface lo0.100;

[edit interfaces lo0]
unit 100 {
    family inet {
        address 192.168.255.1/32;
    }
}

# On PE2
[edit routing-instances CUSTOMER-A]
interface lo0.100;

[edit interfaces lo0]
unit 100 {
    family inet {
        address 192.168.255.2/32;
    }
}
```

Step 2: Configure the sham link on both PEs:

```
# On PE1
[edit routing-instances CUSTOMER-A protocols ospf area 0.0.0.0]
sham-link-remote 192.168.255.2 {
    local 192.168.255.1;
    metric 1;          # Cost of the sham link
    ipsec-sa SHAM-IPSEC; # Optional IPsec protection
}

# On PE2
[edit routing-instances CUSTOMER-A protocols ospf area 0.0.0.0]
sham-link-remote 192.168.255.1 {
    local 192.168.255.2;
    metric 1;
}
```

Step 3: Configure redistribution policy:

```
[edit policy-options]
policy-statement bgp-to-ospf {
    term 1 {
        from protocol bgp;
        then {
            metric 5;
            external {
                type 2;    # Use Type-2 external metric
            }
            accept;
        }
    }
}
```

## Complete Reference Configuration

```
# Complete PE1 Configuration
routing-instances {
    CUSTOMER-A {
```

```

instance-type vrf;
interface ge-0/0/0.100;
interface lo0.100;
route-distinguisher 65000:100;
vrf-target target:65000:100;
vrf-table-label;
protocols {
    ospf {
        export bgp-to-ospf;
        domain-id 0.0.0.1;
        area 0.0.0.0 {
            interface ge-0/0/0.100;
            interface lo0.100 {
                passive;
            }
            sham-link-remote 192.168.255.2 {
                local 192.168.255.1;
                metric 1;
            }
        }
    }
}
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

```

# Verify OSPF neighbor status including sham links
user@PE1> show ospf neighbor instance CUSTOMER-A
Address      Interface      State      ID          Pri  Dead
10.1.1.1     ge-0/0/0.100  Full       1.1.1.1    128   38
192.168.255.2 shamlink      Full       2.2.2.2    0     35

# Check OSPF database for DN bit
user@PE1> show ospf database instance CUSTOMER-A detail | match "DN-bit|Type|Advertising"
Type: Summary
Advertising Router: 2.2.2.2
Options: 0x2 (DN-bit set)

# Verify VPN route tag
user@PE1> show ospf database external instance CUSTOMER-A detail | match tag
Tag: 0xd0fde9

# Check sham link status
user@PE1> show ospf sham-link instance CUSTOMER-A
Sham link to 192.168.255.2, State: Up
  Local address: 192.168.255.1
  Metric: 1, Type: Point-to-Point

```

### Troubleshooting Scenario 1: Routing Loop Detection

**Symptom:** High CPU usage, OSPF database constantly changing

```

user@CE1> show ospf database summary | count
Count: 500 # Unusually high number

user@CE1> show ospf database detail | match "DN|Age"
Age: 2      # Very young LSAs
Options: 0x0 # DN bit NOT set - Problem!

```

**Cause:** PE is not setting DN bit properly **Solution:**

```

# On PE, verify VRF configuration
[edit routing-instances CUSTOMER-A]
set vrf-table-label # This enables proper DN bit handling

```

### Troubleshooting Scenario 2: Suboptimal Routing via Backdoor

**Symptom:** Traffic preferring backdoor link over MPLS

```
user@CE1> show route 10.2.2.0/24
10.2.2.0/24  *[OSPF/10] via ge-0/0/1.0 # Backdoor link
              [OSPF/150] via ge-0/0/0.0 # MPLS path (not preferred)
```

```
user@CE1> show ospf route 10.2.2.0/24 detail
Route Type: Intra-area (via backdoor)
Route Type: Inter-area (via MPLS)
```

**Cause:** MPLS path seen as inter-area while backdoor is intra-area **Solution:** Configure sham link or adjust OSPF costs

### Troubleshooting Scenario 3: Sham Link Not Establishing

**Symptom:** Sham link remains down

```
user@PE1> show ospf sham-link instance CUSTOMER-A
Sham link to 192.168.255.2, State: Down
  Error: Remote endpoint unreachable

user@PE1> show route table CUSTOMER-A.inet.0 192.168.255.2
# No route found
```

**Cause:** Sham link endpoints not redistributed into BGP **Solution:**

```
[edit policy-options policy-statement vrf-export]
term sham-endpoints {
    from {
        protocol direct;
        interface lo0.100;
    }
    then {
        community add vpn-community;
        accept;
    }
}
```

## Module 8: Route Leaking in L3VPN Environment

### Part 1: The Conceptual Lecture (The Why)

#### The Fundamental Problem

Picture a large corporation with three divisions: Sales, Engineering, and Corporate Services. Each division has its own private network (VRF - Virtual Routing and Forwarding instance), completely isolated from the others for security. But here's the challenge: Corporate Services runs shared resources (like email servers and printers) that both Sales and Engineering need to access.

Route leaking is like creating controlled doorways between these isolated rooms. Without it, you'd need to duplicate all shared resources in each VRF, which is expensive and inefficient.

#### Understanding Route Leaking Concepts

Route leaking allows selective sharing of routes between:

1. **VRF-to-VRF:** Between different customer VPNs
2. **VRF-to-Global:** Between a VPN and the Internet routing table
3. **Global-to-VRF:** From Internet to a specific VPN

Think of it as a sophisticated postal system where certain packages can cross between otherwise separate delivery zones:

#### Route Leaking Scenarios

**VRF-to-VRF:**  
[Sales VRF] <--leak--> [Corporate VRF] <--leak--> [Engineering VRF]  
| | |  
Sales Network Shared Servers Engineering Network

**VRF-to-Global (Internet Access):**  
[Customer VRF] --leak--> [inet.0] ---> Internet  
^

## Route Leaking Methods

There are three primary methods, each with different use cases:

1. **RIB Groups**: Direct copying of routes between routing tables
  - Like photocopying documents and placing them in different filing cabinets
2. **Logical Tunnel Interfaces (lt-)**: Creating virtual connections
  - Like building a hallway between two separate buildings
3. **Instance Import/Export**: Using routing policies
  - Like having a customs office that decides what can cross borders

## Part 2: The Junos CLI Masterclass (The How)

### Method 1: RIB Groups Configuration

RIB groups allow routes learned in one table to be copied to other tables:

```
# Step 1: Define the RIB group
[edit routing-options]
rib-groups {
    LEAK-CORPORATE-TO-SALES {
        import-rib [ CORPORATE.inet.0 SALES.inet.0 ];
        # First table is primary, others receive copies
    }

    LEAK-OSPF-TO-VRFS {
        import-rib [ inet.0 SALES.inet.0 ENGINEERING.inet.0 ];
        import-policy SHARED-SERVICES-ONLY; # Filter what gets leaked
    }
}

# Step 2: Apply RIB group to the protocol
[edit routing-instances CORPORATE protocols ospf]
rib-group LEAK-CORPORATE-TO-SALES;

# For static routes
[edit routing-instances CORPORATE routing-options static]
rib-group LEAK-CORPORATE-TO-SALES;
```

### Method 2: Logical Tunnel Interfaces

Logical tunnels create a virtual point-to-point connection:

```
# Step 1: Create logical tunnel interfaces
[edit interfaces]
lt-0/0/0 {
    unit 0 {
        encapsulation ethernet;
        peer-unit 1;
        family inet {
            address 10.255.255.1/30;
        }
    }
    unit 1 {
        encapsulation ethernet;
        peer-unit 0;
        family inet {
            address 10.255.255.2/30;
        }
    }
}

# Step 2: Assign interfaces to routing instances
[edit routing-instances]
SALES {
    interface lt-0/0/0.0;
    protocols {
        bgp {
            group to-corporate {
                peer-as 65002;
            }
        }
    }
}
```

```

        neighbor 10.255.255.2;
        export SALES-TO-CORPORATE;
    }
}
}

CORPORATE {
    interface lt-0/0/0.1;
    protocols {
        bgp {
            group to-sales {
                peer-as 65001;
                neighbor 10.255.255.1;
                export CORPORATE-TO-SALES;
            }
        }
    }
}

```

## Method 3: Instance Import/Export

Most elegant for complex scenarios:

```

# Step 1: Tag routes for export
[edit routing-instances CORPORATE]
routing-options {
    instance-export MARK-CORPORATE-ROUTES;
}

[edit policy-options policy-statement MARK-CORPORATE-ROUTES]
term shared-services {
    from {
        protocol ospf;
        route-filter 10.100.0.0/16 orlonger;
    }
    then {
        community add SHARED-SERVICES;
        accept;
    }
}
term reject {
    then reject;
}

# Step 2: Import tagged routes
[edit routing-instances SALES]
routing-options {
    instance-import IMPORT-SHARED-SERVICES;
}

[edit policy-options policy-statement IMPORT-SHARED-SERVICES]
term shared {
    from {
        instance CORPORATE;
        community SHARED-SERVICES;
    }
    then accept;
}
term reject {
    then reject;
}

[edit policy-options community SHARED-SERVICES]
members 65000:9999;

```

## Internet Access via Route Leaking

Common scenario - providing Internet access to VRF:

```

# Option 1: Default route injection
[edit routing-instances CUSTOMER-A]
routing-options {
    static {
        route 0.0.0.0/0 next-table inet.0;
    }
}

```

```

        }
    }

# Option 2: Selective Internet routes
[edit routing-options]
rib-groups {
    INTERNET-TO-VRF {
        import-rib [ inet.0 CUSTOMER-A.inet.0 ];
        import-policy INTERNET-ROUTES-ONLY;
    }
}

[edit protocols bgp group INTERNET]
family inet {
    unicast {
        rib-group INTERNET-TO-VRF;
    }
}

```

## Complete Reference Configuration

```

# Complete route leaking setup for shared services scenario
routing-options {
    rib-groups {
        SHARED-TO-ALL {
            import-rib [ SHARED.inet.0 SALES.inet.0 ENGINEERING.inet.0 ];
        }
    }
}

routing-instances {
    SHARED {
        instance-type vrf;
        interface ge-0/0/0.100;
        route-distinguisher 65000:999;
        vrf-target target:65000:999;
        protocols {
            ospf {
                rib-group SHARED-TO-ALL;
                area 0.0.0.0 {
                    interface ge-0/0/0.100;
                }
            }
        }
        routing-options {
            instance-export TAG-SHARED;
        }
    }
}

SALES {
    instance-type vrf;
    interface ge-0/0/1.100;
    route-distinguisher 65000:100;
    vrf-target target:65000:100;
    routing-options {
        instance-import ACCEPT-SHARED;
    }
}

policy-options {
    policy-statement TAG-SHARED {
        term 1 {
            from protocol ospf;
            then {
                community add SHARED-COMM;
                accept;
            }
        }
    }
}

policy-statement ACCEPT-SHARED {
    term 1 {
        from {
            instance SHARED;

```

```

        community SHARED-COMM;
    }
    then accept;
}
term default {
    then reject;
}
}

community SHARED-COMM members 65000:9999;
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

```

# Verify routes are leaked between instances
user@PE1> show route table SALES.inet.0 10.100.0.0/24
10.100.0.0/24 *[OSPF/10] via ge-0/0/0.100
[Secondary] Protocol: OSPF, Source: SHARED.inet.0

# Check RIB group operation
user@PE1> show rib-groups
RIB Group: SHARED-TO-ALL
Import RIB: SHARED.inet.0 SALES.inet.0 ENGINEERING.inet.0
Active routes: 25

# Verify instance import/export
user@PE1> show route instance SALES detail | match import
Import: [ ACCEPT-SHARED ]

# Check route source
user@PE1> show route 10.100.0.0/24 detail table SALES.inet.0
*OSPF Preference: 10
Source: 10.100.0.1
Secondary Tables: ENGINEERING.inet.0

```

### Troubleshooting Scenario 1: Routes Not Leaking

**Symptom:** Routes visible in source VRF but not destination VRF

```

user@PE1> show route table SHARED.inet.0 10.100.0.0/24
10.100.0.0/24 *[OSPF/10] via ge-0/0/0.100

user@PE1> show route table SALES.inet.0 10.100.0.0/24
# No output - route not present

```

**Cause:** RIB group not applied to protocol **Solution:**

```

[edit routing-instances SHARED protocols ospf]
set rib-group SHARED-TO-ALL

```

### Troubleshooting Scenario 2: Routing Loop After Leaking

**Symptom:** Route count increasing, TTL exceeded messages

```

user@PE1> show route summary table SALES.inet.0
Total routes: 50000 # Abnormally high

user@PE1> traceroute routing-instance SALES 10.100.0.1
1 10.1.1.1
2 10.2.2.2
3 10.1.1.1 # Loop detected!

```

**Cause:** Leaked routes being re-advertised back **Solution:**

```

[edit policy-options policy-statement ACCEPT-SHARED]
term prevent-loop {
    from {
        protocol ospf;
        interface ge-0/0/1.100; # Don't accept routes learned from SALES
    }
}

```

```

    then reject;
}

```

## Troubleshooting Scenario 3: Asymmetric Routing

**Symptom:** Traffic works one direction but not return

```

user@CE1> ping 10.100.0.1 source 10.1.1.1
# No response

user@PE1> show route table SHARED.inet.0 10.1.1.1
# No route back to source

```

**Cause:** Leaked route only in one direction **Solution:** Ensure bidirectional leaking or use static routes for return path:

```

[edit routing-instances SHARED routing-options static]
route 10.1.1.0/24 next-table SALES.inet.0;

```

## Module 9: Hub-and-Spoke L3VPN Topologies

### Part 1: The Conceptual Lecture (The Why)

#### The Fundamental Problem

Imagine a retail company with one headquarters (HQ) and 100 branch stores. Each store needs to:

1. Access central servers at HQ (inventory, email, HR systems)
2. Send all Internet traffic through HQ's security firewall
3. NOT communicate directly with other stores (security requirement)

Building full connectivity (every site to every site) would require  $\frac{n(n-1)}{2}$  connections. For 101 sites, that's 5,050 connections! Hub-and-spoke reduces this to just 100 connections (one per spoke to the hub).

#### Understanding Hub-and-Spoke Architecture

Hub-and-spoke is like an airline's route map. All flights (data) go through a main hub airport (headquarters) to reach any destination:

Traditional Full Mesh:	Hub-and-Spoke:
Site-A ↔ Site-B	Spoke-A
↖ ↘ ↗	↓
HQ	HUB
↙ ↘ ↗	↑
Site-C ↔ Site-D	Spoke-B
6 connections	2 connections
Complex routing	Simple routing
Expensive	Cost-effective

#### The Route-Target Challenge

In standard L3VPN, all sites in a VPN share the same route-target, creating full-mesh connectivity. Hub-and-spoke requires asymmetric route distribution:

1. **Hub must learn all spoke routes** (to reach all branches)
2. **Spokes must learn only hub routes** (not other spoke routes)
3. **Hub must re-advertise spoke routes** (for spoke-to-spoke via hub)

This is achieved using different import/export route-targets:

Route-Target Flow:	
Spoke Sites:	Hub Site:
Export: 65000:100	Export: 65000:1
Import: 65000:1	Import: 65000:100

Spoke-Hub: Routes tagged with 65000:100, Hub imports them  
 Hub-Spoke: Routes tagged with 65000:1, Spokes import them  
 Spoke-Spoke: Must transit through Hub (no direct RT match)

## Understanding the Data Flow

When Spoke-A sends traffic to Spoke-B:

1. Spoke-A has no route to Spoke-B (different RTs)
2. Spoke-A sends traffic to Hub (default route or aggregate)
3. Hub receives packet, has route to Spoke-B
4. Hub forwards packet to Spoke-B
5. Return traffic follows same pattern

This creates a "hairpin" effect at the hub - traffic comes in one interface and goes out another interface on the same PE router.

## Part 2: The Junos CLI Masterclass (The How)

### Basic Hub-and-Spoke Configuration

#### Hub Site PE Configuration:

```
[edit routing-instances HUB-VRF]
instance-type vrf;
interface ge-0/0/0.100; # To Hub CE
route-distinguisher 65000:1;

# Hub imports spoke routes, exports hub routes
vrf-import HUB-IMPORT;
vrf-export HUB-EXPORT;
vrf-table-label; # Required for same-PE spoke-to-spoke

protocols {
    bgp {
        group hub-ce {
            type external;
            peer-as 65100;
            neighbor 10.0.0.1;
        }
    }
}

[edit policy-options]
policy-statement HUB-IMPORT {
    term from-spokes {
        from community RT-SPOKES;
        then accept;
    }
    term deny {
        then reject;
    }
}

policy-statement HUB-EXPORT {
    term to-spokes {
        from protocol bgp;
        then {
            community add RT-HUB;
            accept;
        }
    }
    # Critical: Re-advertise spoke routes back to spokes
    term spoke-to-spoke {
        from community RT-SPOKES;
        then {
            community add RT-HUB;
            community delete RT-SPOKES;
            accept;
        }
    }
}

community RT-HUB members target:65000:1;
community RT-SPOKES members target:65000:100;
```

#### Spoke Site PE Configuration:

```

[edit routing-instances SPOKE-A-VRF]
instance-type vrf;
interface ge-0/0/1.100; # To Spoke CE
route-distinguisher 65000:101; # Unique per spoke

# Spoke imports hub routes, exports spoke routes
vrf-import SPOKE-IMPORT;
vrf-export SPOKE-EXPORT;

protocols {
    bgp {
        group spoke-ce {
            type external;
            peer-as 65101;
            neighbor 10.1.1.1;
        }
    }
}

[edit policy-options]
policy-statement SPOKE-IMPORT {
    term from-hub {
        from community RT-HUB;
        then accept;
    }
    term deny {
        then reject;
    }
}

policy-statement SPOKE-EXPORT {
    term to-hub {
        from protocol bgp;
        then {
            community add RT-SPOKES;
            accept;
        }
    }
}

```

## Advanced: Multiple Hubs for Redundancy

For high availability, deploy primary and backup hubs:

```

# Primary Hub - Higher Local Preference
[edit policy-options policy-statement HUB-EXPORT]
term to-spokes {
    from protocol bgp;
    then {
        community add RT-HUB;
        local-preference 200; # Primary hub preferred
        accept;
    }
}

# Backup Hub - Lower Local Preference
[edit policy-options policy-statement BACKUP-HUB-EXPORT]
term to-spokes {
    from protocol bgp;
    then {
        community add RT-HUB;
        local-preference 100; # Backup hub
        accept;
    }
}

```

## Spoke-to-Spoke Communication on Same PE

When multiple spokes connect to the same PE router:

```

# Enable local VRF route leaking
[edit routing-instances]
SPOKE-A-VRF {
    routing-options {

```

```

        auto-export {
            family inet {
                unicast;
            }
        }
    }

# Or use vrf-table-label for MPLS switching
[edit routing-instances HUB-VRF]
vrf-table-label; # Creates VPN label for local switching

```

## Default Route Advertisement from Hub

Common requirement - Hub advertises default route to all spokes:

```

[edit routing-instances HUB-VRF]
routing-options {
    static {
        route 0.0.0.0/0 {
            next-hop 10.0.0.254; # To firewall/Internet
            preference 250;
        }
    }
}

[edit policy-options policy-statement HUB-EXPORT]
term default-to-spokes {
    from {
        protocol static;
        route-filter 0.0.0.0/0 exact;
    }
    then {
        community add RT-HUB;
        accept;
    }
}

```

## Complete Reference Configuration

```

# Complete Hub PE Configuration
routing-instances {
    HUB-VRF {
        instance-type vrf;
        interface ge-0/0/0.100;
        interface lo0.100; # Hub loopback
        route-distinguisher 65000:1;
        vrf-import HUB-IMPORT;
        vrf-export HUB-EXPORT;
        vrf-table-label;

        routing-options {
            static {
                route 0.0.0.0/0 next-hop 10.0.0.254;
            }
            autonomous-system 65000;
        }

        protocols {
            bgp {
                group hub-ce {
                    type external;
                    family inet {
                        unicast;
                    }
                    peer-as 65100;
                    neighbor 10.0.0.1 {
                        export TO-HUB-CE;
                    }
                }
            }
        }
    }
}

```

```

policy-options {
    policy-statement HUB-IMPORT {
        term from-spokes {
            from community RT-SPOKES;
            then {
                local-preference 100;
                accept;
            }
        }
        term reject {
            then reject;
        }
    }

    policy-statement HUB-EXPORT {
        term hub-and-default {
            from {
                protocol [ bgp static direct ];
            }
            then {
                community add RT-HUB;
                accept;
            }
        }
        term re-advertise-spokes {
            from community RT-SPOKES;
            then {
                community delete RT-SPOKES;
                community add RT-HUB;
                next-hop self;
                accept;
            }
        }
    }

    policy-statement T0-HUB-CE {
        term send-all {
            then accept;
        }
    }
}

community RT-HUB members target:65000:1;
community RT-SPOKES members target:65000:100;
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

```

# Verify hub has all spoke routes
user@HUB-PE> show route table HUB-VRF.inet.0 summary
Total routes: 250 # Should see all spoke prefixes

# Check spoke only has hub routes
user@SPOKE-PE> show route table SPOKE-A-VRF.inet.0 summary
Total routes: 10 # Much fewer - only hub routes

# Verify route targets
user@HUB-PE> show route advertising-protocol bgp 192.168.1.2 detail
* 10.1.0.0/24 (1 entry)
  Communities: target:65000:1 # Hub RT applied

# Check VPN label allocation for same-PE communication
user@PE1> show route table mpls.0 label 299776
299776 *[VPN/0]
      to table HUB-VRF.inet.0

# Trace spoke-to-spoke path
user@SPOKE-A-CE> traceroute 10.2.2.1 source 10.1.1.1
1 10.1.1.254 (Spoke-A to PE)
2 10.0.0.1 (PE to Hub)
3 10.0.0.254 (Hub to PE)
4 10.2.2.1 (PE to Spoke-B)

```

## Troubleshooting Scenario 1: Spoke-to-Spoke Communication Fails

**Symptom:** Spoke A cannot reach Spoke B

```
user@SPOKE-A-CE> ping 10.2.2.1
No route to host

user@SPOKE-PE> show route table SPOKE-A-VRF.inet.0 10.2.2.0/24
# No route found
```

**Cause:** Hub not re-advertising spoke routes **Solution:**

```
[edit policy-options policy-statement HUB-EXPORT]
term re-advertise {
    from community RT-SPOKES;
    then {
        community delete RT-SPOKES;
        community add RT-HUB;
        accept; # Must accept to re-advertise
    }
}
```

## Troubleshooting Scenario 2: Routing Loop in Hub-and-Spoke

**Symptom:** TTL exceeded, packets looping

```
user@HUB-CE> show interfaces ge-0/0/0.100 statistics
Input packets: 5000000 # Very high
Output packets: 5000000 # Matches input - loop!

user@HUB-CE> show route 10.2.2.0/24
10.2.2.0/24 *[BGP/170] via 10.0.0.254 # Points back to PE
```

**Cause:** Hub CE sending spoke routes back to PE **Solution:** Filter spoke routes from CE to PE

```
# On Hub CE
[edit policy-options policy-statement T0-PE]
term no-spoke-routes {
    from {
        route-filter 10.1.0.0/16 orlonger; # Spoke prefixes
    }
    then reject;
}
term accept {
    then accept;
}
```

## Troubleshooting Scenario 3: Same-PE Spokes Cannot Communicate

**Symptom:** Two spokes on same PE cannot reach each other via hub

```
user@PE1> show route table SPOKE-A-VRF.inet.0 10.2.2.0/24
10.2.2.0/24 *[BGP/170] via HUB-VRF

user@PE1> ping routing-instance SPOKE-A 10.2.2.1
# No response

user@PE1> show route forwarding-table destination 10.2.2.0/24
Destination Type NextHop Interface
10.2.2.0/24 user 0:dead:beef # Bad next-hop
```

**Cause:** Missing vrf-table-label configuration **Solution:**

```
[edit routing-instances HUB-VRF]
set vrf-table-label
```

## Troubleshooting Scenario 4: Suboptimal Routing in Dual-Hub Setup

**Symptom:** Traffic uses backup hub despite primary being available

```

user@SPOKE-PE> show route table SPOKE-A-VRF.inet.0 0.0.0.0/0 detail
0.0.0.0/0 (2 entries)
*BGP Preference: 170, LocalPref: 100 # Backup hub
BGP Preference: 170, LocalPref: 200 # Primary hub (not active)

user@SPOKE-PE> show route hidden table SPOKE-A-VRF.inet.0
0.0.0.0/0 [BGP] Hidden by routing policy # Primary hub route hidden

```

**Cause:** Import policy filtering primary hub routes **Solution:** Check and correct import policy

```

[edit policy-options policy-statement SPOKE-IMPORT]
term from-all-hubs {
    from community [ RT-HUB RT-BACKUP-HUB ];
    then accept;
}

```

These comprehensive modules provide you with the theoretical foundation, practical configuration skills, and troubleshooting expertise needed for OSPF optimal routing, route leaking, and hub-and-spoke L3VPN topologies in preparation for the JNCIE-SP exam.

## Module 10: Layer 3 VPN Class of Service (CoS)

### Part 1: The Conceptual Lecture (The Why)

#### Understanding the Need for QoS in L3VPNs

Imagine you're managing a highway system where multiple companies share the same roads (your MPLS network). Some trucks carry medical supplies (voice traffic), others carry regular packages (data), and some carry bulk materials (file transfers). Without traffic management, all vehicles travel at the same priority, potentially causing critical shipments to be delayed.

**Class of Service (CoS)** in Layer 3 VPNs solves this exact problem in networks. When service providers offer L3VPN services to multiple customers, they need to ensure that each customer's traffic receives appropriate treatment based on its importance.


#### The Fundamental Challenge

In an MPLS L3VPN environment, we face several quality challenges:

1. **Customer Differentiation:** Different customers pay for different service levels
2. **Application Requirements:** Voice needs low latency, video needs consistent bandwidth, bulk data can tolerate delays
3. **Resource Contention:** Limited bandwidth must be shared fairly and efficiently
4. **SLA Compliance:** Service Level Agreements promise specific performance metrics

#### How CoS Works in L3VPNs

The CoS architecture in L3VPNs operates at multiple layers:




Flow: IP Packet → MPLS Encapsulation → Core Transport → MPLS Decapsulation → IP Packet

#### Key Components:

1. **Classification:** Identifying traffic types at PE ingress
2. **Marking:** Setting QoS bits (DSCP in IP, EXP in MPLS)
3. **Queuing:** Placing packets in appropriate queues
4. **Scheduling:** Determining transmission order
5. **Shaping/Policing:** Controlling traffic rates

#### The DSCP to EXP Mapping Challenge

When an IP packet enters the MPLS network, its QoS markings must be preserved:



## **QoS Models in L3VPNs**

### **1. Uniform Mode:**

- PE routers modify customer DSCP values
- Provides consistent QoS across provider network
- Provider controls end-to-end QoS

### **2. Pipe Mode:**

- Customer DSCP values are preserved end-to-end
- Provider uses EXP bits internally
- Customer QoS is tunneled through provider network

### **3. Short-Pipe Mode:**

- Hybrid approach
- Provider QoS in core, customer QoS at egress
- Most flexible for complex requirements

## **Part 2: The Junos CLI Masterclass (The How)**

### **Configuration Hierarchy**

The CoS configuration in Junos for L3VPNs involves multiple configuration stanzas:

```
[edit class-of-service] # Main CoS configuration
[edit interfaces]       # Apply CoS to interfaces
[edit routing-instances] # VRF-specific considerations
[edit firewall]         # Classification filters
```

### **Step-by-Step L3VPN CoS Configuration**

#### **Step 1: Define DSCP to Forwarding Class Mappings**

```
[edit class-of-service]
set dscp-map L3VPN-DSCP-MAP import default
set dscp-map L3VPN-DSCP-MAP forwarding-class expedited-forwarding {
    loss-priority low code-points ef;
}
set dscp-map L3VPN-DSCP-MAP forwarding-class assured-forwarding {
    loss-priority low code-points [ af31 af32 af33 ];
    loss-priority high code-points [ af11 af12 af13 ];
}
set dscp-map L3VPN-DSCP-MAP forwarding-class best-effort {
    loss-priority low code-points be;
}
```

**Purpose:** This creates a custom DSCP classifier that maps customer DSCP values to internal forwarding classes.

#### **Step 2: Configure EXP Classifiers for MPLS**

```
[edit class-of-service]
set exp-map MPLS-EXP-MAP import default
set exp-map MPLS-EXP-MAP forwarding-class expedited-forwarding {
    loss-priority low code-points 5;
}
set exp-map MPLS-EXP-MAP forwarding-class assured-forwarding {
    loss-priority low code-points 3;
    loss-priority high code-points 2;
}
set exp-map MPLS-EXP-MAP forwarding-class best-effort {
    loss-priority low code-points 0;
}
```

**Purpose:** Maps MPLS EXP bits to forwarding classes within the core network.

#### **Step 3: Configure Rewrite Rules**

```
[edit class-of-service]
# DSCP rewrite for customer-facing interfaces
set rewrite-rule dscp L3VPN-DSCP-REWRITE import default
```

```

set rewrite-rule dscp L3VPN-DSCP-REWRITE forwarding-class expedited-forwarding {
    loss-priority low code-point ef;
}

# EXP rewrite for core-facing interfaces
set rewrite-rule exp MPLS-EXP-REWRITE import default
set rewrite-rule exp MPLS-EXP-REWRITE forwarding-class expedited-forwarding {
    loss-priority low code-point 5;
}

```

**Purpose:** Ensures proper marking of packets as they traverse network boundaries.

#### Step 4: Configure Schedulers and Scheduler Maps

```

[edit class-of-service]
# Define schedulers for each forwarding class
set schedulers VOICE-SCHEDULER {
    transmit-rate percent 30;
    buffer-size percent 5;
    priority strict-high;
}

set schedulers DATA-SCHEDULER {
    transmit-rate percent 40;
    buffer-size percent 25;
    priority low;
}

set schedulers BE-SCHEDULER {
    transmit-rate remainder;
    buffer-size remainder;
    priority low;
}

# Create scheduler map
set scheduler-maps L3VPN-SCHEDULER-MAP {
    forwarding-class expedited-forwarding scheduler VOICE-SCHEDULER;
    forwarding-class assured-forwarding scheduler DATA-SCHEDULER;
    forwarding-class best-effort scheduler BE-SCHEDULER;
}

```

**Purpose:** Defines how bandwidth is allocated and packets are scheduled for transmission.

#### Step 5: Apply CoS to PE Interfaces

```

[edit class-of-service]
# Customer-facing interface (CE-facing)
set interfaces ge-0/0/0 {
    unit 100 { # L3VPN sub-interface
        classifiers {
            dscp L3VPN-DSCP-MAP;
        }
        rewrite-rules {
            dscp L3VPN-DSCP-REWRITE;
        }
    }
}

# Core-facing interface (P-facing)
set interfaces ge-0/0/1 {
    scheduler-map L3VPN-SCHEDULER-MAP;
    unit 0 {
        classifiers {
            exp MPLS-EXP-MAP;
        }
        rewrite-rules {
            exp MPLS-EXP-REWRITE;
        }
    }
}

```

**Purpose:** Applies all CoS components to the appropriate interfaces.

#### Step 6: Configure VRF-Specific CoS Policies (Optional)

```
[edit routing-instances CUSTOMER-A]
set routing-options forwarding-table export COS-POLICY

[edit policy-options]
set policy-statement COS-POLICY {
    term VOICE-TRAFFIC {
        from {
            protocol bgp;
            community VOICE-COMMUNITY;
        }
        then {
            class expedited-forwarding;
            accept;
        }
    }
}
```

**Purpose:** Allows per-VRF or per-route CoS classification based on BGP attributes.

## Complete Reference Configuration

```
## Complete L3VPN CoS Configuration ##

class-of-service {
    # DSCP Classifier
    dscp-map L3VPN-DSCP-MAP {
        import default;
        forwarding-class expedited-forwarding {
            loss-priority low code-points ef;
        }
        forwarding-class assured-forwarding {
            loss-priority low code-points [ af31 af32 af33 ];
        }
        forwarding-class best-effort {
            loss-priority low code-points be;
        }
    }

    # EXP Classifier
    exp-map MPLS-EXP-MAP {
        import default;
        forwarding-class expedited-forwarding {
            loss-priority low code-points 5;
        }
        forwarding-class assured-forwarding {
            loss-priority low code-points 3;
        }
        forwarding-class best-effort {
            loss-priority low code-points 0;
        }
    }

    # Rewrite Rules
    rewrite-rule dscp L3VPN-DSCP-REWRITE {
        import default;
        forwarding-class expedited-forwarding {
            loss-priority low code-point ef;
        }
    }

    rewrite-rule exp MPLS-EXP-REWRITE {
        import default;
        forwarding-class expedited-forwarding {
            loss-priority low code-point 5;
        }
    }

    # Schedulers
    schedulers {
        VOICE-SCHEDULER {
            transmit-rate percent 30;
            buffer-size percent 5;
            priority strict-high;
        }
        DATA-SCHEDULER {
    }
```

```

        transmit-rate percent 40;
        buffer-size percent 25;
        priority low;
    }
}

# Scheduler Map
scheduler-maps {
    L3VPN-SCHEDULER-MAP {
        forwarding-class expedited-forwarding scheduler VOICE-SCHEDULER;
        forwarding-class assured-forwarding scheduler DATA-SCHEDULER;
    }
}

# Interface Application
interfaces {
    ge-0/0/0 {
        unit 100 {
            classifiers {
                dscp L3VPN-DSCP-MAP;
            }
        }
    }
    ge-0/0/1 {
        scheduler-map L3VPN-SCHEDULER-MAP;
        unit 0 {
            classifiers {
                exp MPLS-EXP-MAP;
            }
        }
    }
}
}

```

## **Part 3: Verification & Troubleshooting (The What-If)**

## Essential Verification Commands

## 1. Verify CoS Configuration

```
user@PE1> show class-of-service configuration
```

## Good Output:

```
Physical interface: ge-0/0/0
  Unit: 100
    Classifier: L3VPN-DSCP-MAP (dscp)
    Rewrite: L3VPN-DSCP-REWRITE (dscp)
```

```
Physical interface: ge-0/0/1
Scheduler map: L3VPN-SCHEDULER-MAP
Unit: 0
Classifier: MPLS-EXP-MAP (exp)
Rewrite: MPLS-EXP-REWRITE (exp)
```

## 2. Monitor Queue Statistics

```
user@PE1> show interfaces queue ge-0/0/1
```

## Good Output:

```
Queue: 0, Forwarding class: best-effort
  Packets: 1234567, Bytes: 987654321
  Tail-drop packets: 0, RED packets: 0

Queue: 1, Forwarding class: expedited-forwarding
  Packets: 234567, Bytes: 34567890
  Tail-drop packets: 0, RED packets: 0
```

### **3. Verify DSCP/EXP Markings**

```
user@PE1> show class-of-service classifier name L3VPN-DSCP-MAP
```

## Common Troubleshooting Scenarios

### Scenario 1: Voice Quality Issues Despite CoS Configuration

**Symptom:** Customer reports poor voice quality, jitter, and delays

**Diagnostic Commands:**

```
user@PE1> show interfaces queue ge-0/0/1 | match "drop|expedited"
Queue: 1, Forwarding class: expedited-forwarding
Tail-drop packets: 15234, RED packets: 0 # Problem: Drops in voice queue!

user@PE1> show interfaces ge-0/0/1 extensive | match "output rate"
Output rate: 95000000 bps # Interface near saturation
```

**Cause:** Voice queue is too small or bandwidth allocation insufficient

**Solution:**

```
[edit class-of-service schedulers VOICE-SCHEDULER]
set transmit-rate percent 40; # Increase from 30%
set buffer-size percent 10; # Increase from 5%

[edit class-of-service]
set interfaces ge-0/0/1 shaping-rate 900m; # Add shaping to prevent overrun
```

### Scenario 2: DSCP Values Not Preserved Through MPLS Network

**Symptom:** CE2 receives all traffic with DSCP 0 (best-effort)

**Diagnostic Commands:**

```
user@PE2> show class-of-service interface ge-0/0/0.100
# No rewrite rules shown!

user@PE2> monitor traffic interface ge-0/0/0.100 no-resolve
15:23:45.123456 In IP 10.1.1.1 > 10.2.2.2: Flags [P.], TOS 0x0 # Wrong!
```

**Cause:** Missing DSCP rewrite rules on egress PE

**Solution:**

```
[edit class-of-service interfaces ge-0/0/0]
set unit 100 rewrite-rules dscp L3VPN-DSCP-REWRITE protocol mpls;
```

### Scenario 3: CoS Not Working for Specific VRF Traffic

**Symptom:** Traffic from one VRF doesn't receive proper QoS treatment

**Diagnostic Commands:**

```
user@PE1> show route forwarding-table vpn CUSTOMER-A
# Check for class-of-service information

user@PE1> show policy COS-POLICY
# Verify policy is correctly configured
```

**Cause:** VRF-specific CoS policy not applied or incorrectly configured

**Solution:**

```
[edit routing-instances CUSTOMER-A]
set vrf-table-label; # Enable VRF table label for CoS processing

[edit routing-instances CUSTOMER-A routing-options]
set forwarding-table export COS-POLICY;
```

### Scenario 4: Asymmetric CoS Behavior

**Symptom:** CoS works PE1→PE2 but not PE2→PE1

**Diagnostic Commands:**

```
user@PE2> show configuration class-of-service | display set  
# Compare with PE1 configuration  
  
user@PE2> show route table mpls.0 label 300123 detail  
# Check MPLS path and EXP bit handling
```

**Cause:** Inconsistent CoS configuration across PEs

**Solution:**

```
# Ensure symmetric configuration on all PEs  
[edit groups L3VPN-COS]  
set class-of-service <complete-cos-config>  
  
[edit]  
set apply-groups L3VPN-COS;
```

## Module 11: Layer 3 VPN Protection Mechanisms

### Part 1: The Conceptual Lecture (The Why)

#### Understanding the Need for Protection

Imagine you're running a critical business application across your L3VPN. Suddenly, a fiber cut occurs between PE and CE routers. In traditional networks, BGP would take 30-180 seconds to detect the failure and reconverge. For voice calls, financial transactions, or real-time applications, even a few seconds of downtime is catastrophic.

**Protection mechanisms** solve this by pre-computing backup paths and enabling millisecond-level failover.

#### The Convergence Challenge

Traditional BGP convergence involves multiple steps:

Timeline of Traditional Failure:

T+0s: Link fails x  
T+3s: BGP Hold Timer starts  
T+30–180s: BGP session times out  
T+31s: Routes withdrawn  
T+32s: RIB recalculation  
T+33s: FIB update  
T+34s: Traffic restored

Total Downtime: 30+ seconds!

#### BGP PIC Edge (Prefix Independent Convergence)

BGP PIC Edge revolutionizes failover by pre-installing backup paths:

Traditional BGP FIB:

Prefix: 10.1/24
NH: PE2
(Single Path)

BGP PIC Edge FIB:

Prefix: 10.1/24
Primary NH: PE2
Backup NH: PE3
(Pre-installed)

When PE2 fails:

- Delete old NH
- Calculate new NH
- Install new NH
- Time: Seconds

When PE2 fails:

- Switch pointer to backup
- No calculation needed
- Instant failover
- Time: Milliseconds

#### How BGP PIC Edge Works

The mechanism operates in three phases:

## 1. Pre-computation Phase:

PE1 receives multiple paths to same prefix:  
Path 1: 10.1.1.0/24 via PE2 (Local Pref: 100, Active)  
Path 2: 10.1.1.0/24 via PE3 (Local Pref: 90, Backup)

Instead of discarding Path 2, PE1:  
– Installs Path 1 as primary in FIB  
– Pre-installs Path 2 as backup in FIB  
– Creates protection group linking them

## 2. Failure Detection Phase:

Detection Methods:  
– BFD: 50ms detection  
– Link Down: Immediate  
– IGP Convergence: <500ms

## 3. Switchover Phase:

Upon detection:  
– Hardware switches traffic to backup path  
– No BGP recalculation required  
– No RIB-to-FIB download needed  
– Failover time: <50ms with BFD

# Provider Edge Link Protection

PE Link Protection addresses the last-mile problem:



Problem: Single point of failure at PE-CE link  
Solution: Redundant connections with coordinated failover

## Protection Mechanisms:

### 1. CE Multihoming:

- CE connects to multiple PEs
- Uses BGP multipath or active/standby

### 2. PE Node Protection:

- Protects against entire PE failure
- Backup PE pre-provisions VRFs


### 3. Access Link Protection:

- Protects PE-CE link
- Can use link aggregation or diverse paths

# Protection Coordination

All protection mechanisms must coordinate:

Protection Hierarchy:



(Core protection - 50ms)

## Part 2: The Junos CLI Masterclass (The How)

### BGP PIC Edge Configuration

#### Step 1: Enable BGP PIC Edge Globally

```
[edit routing-options]
set protect core;

[edit protocols bgp]
set group IBGP family inet-vpn unicast protection;
```

**Purpose:** Enables PIC Edge computation for BGP routes and protects core-facing interfaces.

#### Step 2: Configure Per-VRF PIC Edge

```
[edit routing-instances CUSTOMER-A]
set routing-options protect core;

[edit routing-instances CUSTOMER-A protocols bgp]
set group CE family inet unicast {
    protection;
    add-path receive;
}
```

**Purpose:** Enables protection within specific VRFs and allows multiple path reception.

#### Step 3: Configure BGP Add-Path for Backup Path Advertisement

```
[edit protocols bgp group IBGP]
set family inet-vpn unicast add-path {
    send {
        path-count 2;
        path-selection-mode all-paths;
    }
    receive;
}
```

**Purpose:** Ensures PEs advertise both primary and backup paths to each other.

#### Step 4: Enable Edge Protection with BFD

```
[edit protocols bgp group CE]
set neighbor 192.168.1.1 {
    bfd-liveness-detection {
        minimum-interval 50;
        multiplier 3;
        transmit-interval {
            minimum-interval 50;
        }
    }
}
```

**Purpose:** Enables fast failure detection (150ms in this case).

### Provider Edge Link Protection Configuration

#### Step 1: Configure CE Multihoming with Different PEs

On Primary PE (PE1):

```
[edit routing-instances CUSTOMER-A]
set instance-type vrf;
set interface ge-0/0/0.100;
set route-distinguisher 65000:100;
set vrf-target target:65000:100;

[edit routing-instances CUSTOMER-A protocols bgp]
set group CE {
```

```

type external;
peer-as 65001;
neighbor 192.168.1.1 {
    local-address 192.168.1.2;
    local-preference 200; # Higher preference
}
}

```

On Backup PE (PE2):

```

[edit routing-instances CUSTOMER-A]
set instance-type vrf;
set interface ge-0/0/0.200;
set route-distinguisher 65000:101; # Different RD
set vrf-target target:65000:100; # Same RT

[edit routing-instances CUSTOMER-A protocols bgp]
set group CE {
    type external;
    peer-as 65001;
    neighbor 192.168.2.1 {
        local-address 192.168.2.2;
        local-preference 100; # Lower preference
    }
}

```

**Purpose:** Creates active/standby PE protection with preference-based selection.

## Step 2: Configure Virtual Router Redundancy Protocol (VRRP) for CE Protection

```

[edit interfaces ge-0/0/0]
set unit 100 {
    family inet {
        address 192.168.1.2/24 {
            vrrp-group 10 {
                virtual-address 192.168.1.254;
                priority 200; # Primary
                preempt;
                accept-data;
                track {
                    interface ge-0/0/1.0 priority-cost 50;
                }
            }
        }
    }
}

```

**Purpose:** Provides IP-level redundancy for CE with single connection.

## Step 3: Configure Protected Core Links

```

[edit protocols mpls]
set interface ge-0/0/1.0 {
    link-protection;
}

[edit protocols rsvp]
set interface ge-0/0/1.0 {
    link-protection;
}

```

**Purpose:** Enables MPLS fast reroute for core-facing interfaces.

## Complete Reference Configuration for Protection

```

## PE1 Configuration – Primary PE ##

# Global Protection Settings
routing-options {
    protect core;

    # Resolution for backup next-hops
    resolution {

```

```

        rib inet.0 {
            resolution-ribs [ inet.3 inet.0 ];
        }
        rib bgp.l3vpn.0 {
            resolution-ribs [ inet.3 inet.0 ];
        }
    }
}

# BGP Configuration with PIC Edge
protocols {
    bgp {
        group IBGP {
            type internal;
            local-address 1.1.1.1;
            family inet-vpn {
                unicast {
                    protection;
                    add-path {
                        send {
                            path-count 2;
                        }
                        receive;
                    }
                }
            }
            neighbor 2.2.2.2;
            neighbor 3.3.3.3;
        }
    }
}

# Enable MPLS Protection
mpls {
    interface ge-0/0/1.0 {
        link-protection;
    }
}

# Enable BFD for Fast Detection
bfd {
    traceoptions {
        file bfd-log;
        flag all;
    }
}
}

# VRF Configuration with Protection
routing-instances {
    CUSTOMER-A {
        instance-type vrf;
        interface ge-0/0/0.100;
        route-distinguisher 65000:100;
        vrf-target target:65000:100;
        routing-options {
            protect core;
        }
        protocols {
            bgp {
                group CE {
                    type external;
                    peer-as 65001;
                    family inet {
                        unicast {
                            protection;
                        }
                    }
                    neighbor 192.168.1.1 {
                        local-address 192.168.1.2;
                        local-preference 200;
                        bfd-liveness-detection {
                            minimum-interval 50;
                            multiplier 3;
                        }
                    }
                }
            }
        }
    }
}
}

```

```

        }
    }
}

# Interface Configuration with VRRP
interfaces {
    ge-0/0/0 {
        unit 100 {
            vlan-id 100;
            family inet {
                address 192.168.1.2/24 {
                    vrrp-group 10 {
                        virtual-address 192.168.1.254;
                        priority 200;
                        preempt;
                        accept-data;
                        authentication-type simple;
                        authentication-key "$9$secret";
                    }
                }
            }
        }
    }
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### 1. Verify BGP PIC Edge Status

```
user@PE1> show route protection 10.1.1.0/24 table CUSTOMER-A.inet.0
```

**Good Output:**

```
10.1.1.0/24
  Protection: Enabled
  Primary: via ge-0/0/0.100
  Backup: via ge-0/0/1.200
  Protection Group: 1024
  State: Both paths installed
```

#### 2. Check Protection Groups

```
user@PE1> show route forwarding-table vpn CUSTOMER-A
```

**Good Output:**

Destination	Type	Next hop	Index	NhRef	Protection
10.1.1.0/24	user	192.168.1.1	1048	2	Group: 1024
		192.168.2.1(B)	1049	1	Backup

#### 3. Monitor BFD Sessions

```
user@PE1> show bfd session
```

**Good Output:**

		Detect	Transmit	
Address	State	Time	Interval	Multiplier
192.168.1.1	Up	0.150	0.050	3

### Common Troubleshooting Scenarios

#### Scenario 1: BGP PIC Edge Not Installing Backup Paths

**Symptom:** Only primary path visible, no backup despite multiple paths available

**Diagnostic Commands:**

```
user@PE1> show route receive-protocol bgp 2.2.2.2 10.1.1.0/24 detail
# Shows: Only one path received

user@PE1> show bgp neighbor 2.2.2.2 | match add-path
# Shows: Add-path not negotiated
```

**Cause:** Add-path capability not negotiated between PEs

**Solution:**

```
[edit protocols bgp group IBGP]
delete neighbor 2.2.2.2;
set neighbor 2.2.2.2 family inet-vpn unicast add-path receive;

# Force session reset
restart bgp neighbor 2.2.2.2
```

## Scenario 2: Slow Failover Despite PIC Edge Configuration

**Symptom:** Failover takes several seconds instead of milliseconds

**Diagnostic Commands:**

```
user@PE1> show route 10.1.1.0/24 extensive | match "protection|detect"
# No BFD information shown

user@PE1> show interfaces ge-0/0/0.100 | match bfd
# BFD not enabled on interface
```

**Cause:** BFD not configured or not operational

**Solution:**

```
[edit protocols bgp group CE neighbor 192.168.1.1]
set bfd-liveness-detection {
    minimum-interval 50;
    multiplier 3;
    transmit-interval minimum-interval 50;
}

[edit interfaces ge-0/0/0 unit 100]
set family inet address 192.168.1.2/24;
```

## Scenario 3: VRRP Flapping Causing Instability

**Symptom:** VRRP mastership keeps changing, causing intermittent connectivity

**Diagnostic Commands:**

```
user@PE1> show vrrp detail
Interface: ge-0/0/0.100, Group: 10, State: backup # Wrong state!
Priority: 150 # Reduced from 200

user@PE1> show log messages | match VRRP
VRRP_INTF_DOWN: Interface ge-0/0/1.0 down, reducing priority
```

**Cause:** Track interface is flapping, causing priority changes

**Solution:**

```
[edit interfaces ge-0/0/0 unit 100 family inet address 192.168.1.2/24]
set vrrp-group 10 {
    track {
        interface ge-0/0/1.0 {
            priority-cost 30; # Reduce from 50
            bandwidth-threshold 100m priority-cost 20;
        }
    }
    preempt-delay 60; # Add delay to prevent flapping
}
```

## Scenario 4: Protection Not Working After PE Reload

**Symptom:** After PE reboot, protection paths not installed

**Diagnostic Commands:**

```
user@PE1> show route forwarding-table vpn CUSTOMER-A summary
Routes: 1000 Protection-enabled: 0 # No protected routes!

user@PE1> show route summary | match protect
Protection: Core protection not enabled
```

**Cause:** Protection core not enabled after configuration restore

**Solution:**

```
[edit routing-options]
set protect core;

# Verify protection is in startup config
show configuration routing-options | match protect | display set

# Request route recalculation
request routing-table build
```

## Module 12: Layer 3 VPN Scaling

### Part 1: The Conceptual Lecture (The Why)

#### Understanding the Scaling Challenge

Imagine a service provider starting with 10 L3VPN customers, each with 5 sites. That's manageable: 50 sites, perhaps 500 routes. Fast forward two years: 1,000 customers, average 20 sites each, 50 routes per site. Suddenly you're dealing with 1 million routes! Without proper scaling techniques, your PE routers would collapse under the computational and memory load.

#### The Fundamental Scaling Bottlenecks

L3VPNs face multiple scaling challenges:

Scaling Dimensions:

Control Plane	
• BGP Sessions	→ Thousands of peers
• Route Count	→ Millions of prefixes
• VRF Count	→ Thousands of VRFs
• Update Rate	→ Thousands/second

Data Plane	
• FIB Entries	→ Hardware limits
• Label Space	→ $2^{20}$ labels max
• Tunnel Scale	→ LSP/tunnel limits
• Packet Rate	→ Millions pps

#### Scaling Technique 1: Route Reflection

Without route reflection, every PE needs full mesh connectivity:

```
Full Mesh (n PEs = n*(n-1)/2 sessions):
4 PEs = 6 sessions
10 PEs = 45 sessions
100 PEs = 4,950 sessions!
1000 PEs = 499,500 sessions!!!
```

```
With Route Reflector:
1000 PEs = 1000 sessions (to RR only)
```

#### Scaling Technique 2: BGP Route Target Constraint (RTC)

Traditional BGP sends all VPN routes to all PEs:

Without RTC:	With RTC:
PE1 (Hosts VRF-A, VRF-B)	PE1 (Hosts VRF-A, VRF-B)
Receives:	Receives:
- VRF-A routes ✓	- VRF-A routes ✓
- VRF-B routes ✓	- VRF-B routes ✓
- VRF-C routes ✗ (Wasted)	- (VRF-C filtered at RR)
- VRF-D routes ✗ (Wasted)	- (VRF-D filtered at RR)
- VRF-E routes ✗ (Wasted)	- (VRF-E filtered at RR)

Memory saved: 60–90% in large deployments!

## Scaling Technique 3: Hierarchical VPNs (HoVPN)

Hierarchical VPNs create a hub-and-spoke architecture:

Traditional Flat Design:	Hierarchical Design:
Every PE connects to every PE	PE connects only to local hub
PE1 ↔ PE2	Level 2: [SuperPE1]---[SuperPE2]
↑ ✗ ↓	
PE3 ↔ PE4	Level 1: PE1,PE2     PE3,PE4

Reduction:  $O(n^2)$  to  $O(n)$

## Scaling Technique 4: BGP ORF (Outbound Route Filtering)

ORF allows PEs to tell route reflectors what they don't want:

PE1 → RR: "I only want routes with RT 65000:100 and 65000:200"  
RR: "Acknowledged, filtering all others before sending"

Result: Reduced processing and bandwidth usage

## Scaling Technique 5: VRF Localization

Instead of all PEs hosting all VRFs:

Inefficient:	Efficient:
PE1: VRF-A,B,C,D,E,F	PE1: VRF-A,B (Regional)
PE2: VRF-A,B,C,D,E,F	PE2: VRF-C,D (Regional)
PE3: VRF-A,B,C,D,E,F	PE3: VRF-E,F (Regional)

Inter-region via Super-PE

## Memory and CPU Optimization

Understanding resource consumption:

Per-Route Memory:
- IPv4 Route: ~200 bytes
- VPnv4 Route: ~300 bytes
- With all attributes: ~500 bytes

1 Million routes = ~500MB just for routes!  
Add FIB programming, adjacencies, etc = 2–3GB

CPU Impact:
- BGP Update processing: $O(n*m)$ where $n=routes$ , $m=peers$
- FIB Programming: $O(n)$
- Route Resolution: $O(n*\log n)$

## Part 2: The Junos CLI Masterclass (The How)

### Configuration for Route Reflection

#### Step 1: Configure Route Reflector

```
[edit protocols bgp group RR-CLIENTS]
set type internal;
```

```

set local-address 10.0.0.1;
set family inet-vpn unicast;
set cluster 10.0.0.1;
set neighbor 10.0.0.0/24 {
    allow; # Dynamic neighbors from range
}

[edit protocols bgp group RR-CLIENTS]
set multipath;
set add-path {
    send {
        path-count 2;
    }
}

```

**Purpose:** Establishes router as route reflector for VPNV4 routes with redundancy support.

## Step 2: Configure Route Reflector Clients

```

[edit protocols bgp group T0-RR]
set type internal;
set local-address 10.0.0.10;
set family inet-vpn unicast;
set neighbor 10.0.0.1 {
    description "Primary Route Reflector";
}
set neighbor 10.0.0.2 {
    description "Backup Route Reflector";
}

```

**Purpose:** PE establishes sessions only with RRs instead of all PEs.

## Configuration for BGP Route Target Constraint

### Step 1: Enable Route Target Constraint

```

[edit protocols bgp]
set family route-target advertise-default;

[edit protocols bgp group RR-CLIENTS]
set family route-target;

```

**Purpose:** Enables RTC capability advertisement.

### Step 2: Configure Automatic Route Target Filtering

```

[edit routing-options]
set route-target-filter automatic;

[edit protocols bgp group T0-RR]
set family route-target {
    advertise-default;
    external-paths 2;
}

```

**Purpose:** Automatically generates RT filters based on configured VRFs.

## Configuration for BGP ORF

### Step 1: Configure ORF Capability

```

[edit protocols bgp group T0-RR]
set family inet-vpn {
    unicast {
        orf-prefix-list;
    }
}

[edit policy-options]
set prefix-list VPN-PREFIXES {
    apply-path "routing-instances <*> vrf-target <*>";
}

```

```
[edit protocols bgp group T0-RR]
set family inet-vpn unicast {
    prefix-limit {
        maximum 100000;
        teardown 80 idle-timeout 30;
    }
}
```

**Purpose:** Enables Outbound Route Filtering and sets safety limits.

## Configuration for Hierarchical VPNs

### Step 1: Configure Area Border PE (Super-PE)

```
[edit routing-instances HUB-VRF]
set instance-type vrf;
set route-distinguisher 65000:9999;
set vrf-target {
    import target:65000:1000; # Import from all spokes
    export target:65000:2000; # Export to spokes
}

[edit routing-instances HUB-VRF routing-options]
set aggregate {
    route 10.0.0.0/8 {
        policy AGGREGATE-POLICY;
    }
}

[edit policy-options policy-statement AGGREGATE-POLICY]
set term AGGREGATE {
    from {
        route-filter 10.0.0.0/8 longer;
    }
    then accept;
}
```

**Purpose:** Creates hierarchical hub that aggregates routes.

### Step 2: Configure Spoke PEs

```
[edit routing-instances SPOKE-VRF]
set instance-type vrf;
set route-distinguisher 65000:1001;
set vrf-target {
    import target:65000:2000; # From hub
    export target:65000:1000; # To hub
}
set routing-options {
    static {
        route 0.0.0.0/0 next-table HUB-VRF.inet.0;
    }
}
```

**Purpose:** Spoke PEs use default route to hub instead of full routing.

## Advanced Scaling Optimizations

### Step 1: Configure BGP Session Parameters for Scale

```
[edit protocols bgp]
set precision-timers;
set hold-time 180;
set keep all;
set log-updown;

[edit protocols bgp group T0-RR]
set out-delay 5;
set damping;
set multipath;
set tcp-mss 4096;
set mtu-discovery;
```

**Purpose:** Optimizes BGP for high-scale environments.

## Step 2: Configure FIB Optimization

```
[edit routing-options]
set forwarding-table {
    export FIB-FILTER;
    ecmp-fast-reroute;
    indirect-next-hop;
    unicast-reverse-path feasible-paths;
}

[edit policy-options policy-statement FIB-FILTER]
set term CRITICAL-ROUTES {
    from {
        route-filter 0.0.0.0/0 upto /24;
    }
    then accept;
}
set term SUPPRESS-HOST-ROUTES {
    from {
        route-filter 0.0.0.0/0 prefix-length-range /32-/32;
    }
    then reject;
}
```

**Purpose:** Limits FIB entries to aggregated routes only.

## Step 3: Configure VRF Limits and Monitoring

```
[edit routing-instances CUSTOMER-A]
set routing-options {
    maximum-prefixes 10000 {
        threshold 80;
        log-interval 30;
    }
    maximum-paths 16;
}

[edit routing-instances CUSTOMER-A protocols bgp]
set group CE {
    family inet {
        unicast {
            prefix-limit {
                maximum 5000;
                teardown 80 idle-timeout 10;
            }
        }
    }
}
```

**Purpose:** Prevents single VRF from consuming excessive resources.

## Complete Scaling Configuration Reference

```
## Comprehensive L3VPN Scaling Configuration ##

# System-level Scaling Parameters
system {
    processes {
        routing {
            bgp {
                rib-sharding;
                update-threading;
            }
        }
    }
}

# Routing Options for Scale
routing-options {
    protect core;
    route-target-filter automatic;
```

```

forwarding-table {
    export FIB-OPTIMIZE;
    ecmp-fast-reroute;
    indirect-next-hop;
}

rib {
    inet.0 {
        aggregate-nexthop-install;
    }
    bgp.l3vpn.0 {
        aggregate-label {
            community target:65000:999;
        }
    }
}
}

# BGP Configuration for Scale
protocols {
    bgp {
        precision-timers;
        log-updown;
        damping;

        # Route Reflector Configuration
        group RR-CLIENTS {
            type internal;
            local-address 10.0.0.1;
            cluster 10.0.0.1;

            family inet-vpn {
                unicast {
                    add-path {
                        send {
                            path-count 2;
                            path-selection-mode all-paths;
                        }
                    }
                }
            }
        }

        family route-target;

        neighbor 10.0.0.0/24 {
            allow;
            tcp-mss 4096;
            mtu-discovery;
        }
    }
}

# ORF Configuration
group PE-GROUP {
    family inet-vpn {
        unicast {
            orf-prefix-list;
            prefix-limit {
                maximum 1000000;
                teardown 90 idle-timeout 60;
            }
        }
    }
}
}

# Hierarchical VPN Hub Configuration
routing-instances {
    REGIONAL-HUB {
        instance-type vrf;
        route-distinguisher 65000:9999;

        vrf-target {
            import target:65000:1000;
            export target:65000:2000;
        }
    }
}

```

```

routing-options {
    aggregate {
        route 10.0.0.0/8;
        route 172.16.0.0/12;
        route 192.168.0.0/16;
    }

    maximum-prefixes 100000 {
        threshold 80;
    }
}

protocols {
    bgp {
        group SPOKE-PEs {
            family inet {
                unicast {
                    aggregate-label {
                        community 65000:9999;
                    }
                }
            }
        }
    }
}

# Policy for FIB Optimization
policy-options {
    policy-statement FIB-OPTIMIZE {
        term ALLOW-AGGREGATES {
            from {
                route-filter 0.0.0.0/0 upto /24;
            }
            then accept;
        }
        term SUPPRESS-HOST {
            from {
                route-filter 0.0.0.0/0 prefix-length-range /30-/32;
            }
            then reject;
        }
    }
}
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### 1. Monitor Scaling Metrics

```
user@PE1> show bgp summary
```

##### Good Output:

Groups: 2 Peers: 2 Down peers: 0						
Table	Tot Paths	Act Paths	Suppressed	History	Pending	
bgp.l3vpn.0	250000	125000	0	0	0	
inet.0	10000	5000	0	0	0	
Peer	AS	InPkt	OutPkt	OutQ	Flaps	State
10.0.0.1	65000	5000000	10000000	0	0	Established

#### 2. Check Route Target Constraint

```
user@PE1> show route table bgp.rttarget.0
```

##### Good Output:

```
bgp.rttarget.0: 20 destinations, 20 routes
65000:100/96
 *[BGP/170] 2d 00:00:00
   Local
65000:200/96
 *[BGP/170] 2d 00:00:00
   Local
```

### 3. Monitor Memory Usage

```
user@PE1> show task memory detail | match "BGP|RIB"
```

#### Good Output:

```
BGP:          500 MB (Within limits)
RIB Manager:  800 MB (Within limits)
FIB:          300 MB (Within limits)
```

## Common Troubleshooting Scenarios

### Scenario 1: PE Running Out of Memory

**Symptom:** High memory usage, BGP sessions flapping, system sluggish

#### Diagnostic Commands:

```
user@PE1> show system processes extensive | match rpd
 PID  USERNAME  VSZ    RSS   %MEM  COMMAND
 1234  root      4096M  3800M 95%  /usr/sbin/rpd  # Problem!

user@PE1> show route summary
Total routes: 2,500,000  # Too many!
```

**Cause:** No route filtering or limits configured

#### Solution:

```
# Implement Route Target Constraint
[edit protocols bgp]
set family route-target;

# Set prefix limits per VRF
[edit routing-instances <*>]
set routing-options maximum-prefixes 10000 threshold 80;

# Enable FIB filtering
[edit routing-options forwarding-table]
set export FILTER-FIB;

# Restart RPD to clear memory
restart routing immediate
```

### Scenario 2: Route Reflector Overloaded

**Symptom:** RR CPU at 100%, slow update propagation

#### Diagnostic Commands:

```
user@RR> show bgp neighbor | match "updates|queue"
  Output queue: 50000  # Huge queue!
  Updates sent: 10000000

user@RR> show system processes | match rpd
  CPU: 99%  # Overloaded
```

**Cause:** RR handling too many clients without optimization

#### Solution:

```
# Enable update threading
[edit system processes]
```

```

set routing bgp update-threading;

# Configure multiple RR clusters
[edit protocols bgp]
set group RR-CLIENTS-1 cluster 10.0.0.1;
set group RR-CLIENTS-2 cluster 10.0.0.2;

# Enable rib-sharding
[edit system processes routing]
set rib-sharding;

# Add more RRs for load distribution

```

### Scenario 3: FIB Programming Delays

**Symptom:** Routes in RIB but not in FIB, forwarding failures

**Diagnostic Commands:**

```

user@PE1> show route forwarding-table summary
Routes: 100000 Pending: 50000 # Half pending!

user@PE1> show pfe statistics traffic
Discard route: 1000000 pps # Drops due to missing FIB

```

**Cause:** FIB programming cannot keep up with updates

**Solution:**

```

# Enable indirect next-hop
[edit routing-options forwarding-table]
set indirect-next-hop;

# Reduce FIB entries via aggregation
[edit routing-instances <*> routing-options]
set aggregate {
    route 10.0.0.0/8;
    route 172.16.0.0/12;
}

# Increase FIB update priority
[edit routing-options]
set forwarding-table {
    export FIB-PRIORITY;
    priority-update;
}

```

### Scenario 4: BGP Convergence Too Slow

**Symptom:** Takes minutes for routes to propagate across network

**Diagnostic Commands:**

```

user@PE1> show bgp neighbor 10.0.0.1 | match "Last traffic"
Last traffic (seconds): Received 300 Sent 300 # 5 min gaps!

user@PE1> monitor traffic interface lo0.0
# Shows bursty BGP updates every 30 seconds

```

**Cause:** Default BGP timers and no optimization

**Solution:**

```

# Enable precision timers
[edit protocols bgp]
set precision-timers;

# Reduce out-delay
[edit protocols bgp group T0-RR]
set out-delay 2;

# Enable TCP optimizations
[edit protocols bgp group T0-RR]
set tcp-mss 4096;

```

```

set mtu-discovery;

# Configure update packing
[edit protocols bgp]
set update-packing {
    timeout 100;
}

```

## Scaling Guidelines Summary

Small Deployment (< 100 PEs, < 10K routes):

- Basic configuration
- No special optimization needed

Medium Deployment (100–500 PEs, 10K–100K routes):

- Implement Route Reflectors
- Enable Route Target Constraint
- Configure prefix limits

Large Deployment (500–1000 PEs, 100K–1M routes):

- Multiple Route Reflector clusters
- Hierarchical VPNs
- FIB optimization
- ORF implementation

Extra Large Deployment (> 1000 PEs, > 1M routes):

- All above optimizations
- Dedicated RR hierarchy
- Regional aggregation
- Hardware offload
- Consider SDN controller

This completes the comprehensive training modules for L3VPN CoS, Protection Mechanisms, and Scaling. Each module has been structured to build understanding from fundamental concepts through practical implementation to real-world troubleshooting scenarios.


## Module 13: BGP Route Target Filtering

### Part 1: The Conceptual Lecture (The Why)

#### The Fundamental Problem

Imagine you're running a massive service provider network with thousands of Layer 3 VPN customers. Each Provider Edge (PE) router in your network receives VPN routing information from every other PE router, regardless of whether it actually needs that information. This is like receiving mail for every single house in your city when you only need mail for the three houses on your street.

Let's visualize this problem:



Problem: PE-1 receives routes for VPN-C,D,E,F,G even though it only needs routes for VPN-A and VPN-B!

### What is BGP Route Target Filtering?

BGP Route Target Filtering (RT Filtering or RTF) is a mechanism that allows PE routers to tell Route Reflectors (RRs) exactly which VPN routes they want to receive. It's like having a subscription service where you only get updates about the specific topics you're interested in.

The technology uses a new BGP address family called **RT-Constraint** (Route Target Constraint) to communicate filtering requirements upstream to Route Reflectors.

## How RT Filtering Works

The mechanism operates in three phases:

### Phase 1: RT Membership Advertisement

```
PE Router          Route Reflector
|                 |
| "I want routes with:   |
|   RT: 65000:100      |
|   RT: 65000:200"     |
|----->|
| (RT Constraint NLRI) |
```

### Phase 2: Route Reflector Builds Filter

Route Reflector's RT Filter Table:

PE	Interested RTs
PE-1	65000:100, 65000:200
PE-2	65000:300, 65000:400
PE-3	65000:500, 65000:600
PE-4	65000:100, 65000:700

### Phase 3: Selective Route Distribution

When RR receives VPN route with RT 65000:100:

- Check filter table
- Send only to PE-1 and PE-4
- Do NOT send to PE-2 and PE-3

## The RT-Constraint Address Family

RT-Constraint uses a special NLRI (Network Layer Reachability Information) format:

```
+-----+
| Origin AS (4B) | <- Who originated this RT membership
+-----+
| RT Value (8B) | <- The actual Route Target
+-----+
```

Example NLRI: 65000:0:65000:100  
  ^AS   ^Type ^RT Value

## Benefits of RT Filtering

1. **Memory Savings:** PEs only store relevant VPN routes
2. **CPU Savings:** Fewer routes to process and evaluate
3. **Bandwidth Savings:** Less BGP update traffic
4. **Faster Convergence:** Smaller routing tables converge faster

## Part 2: The Junos CLI Masterclass (The How)

### Configuration Hierarchy

RT Filtering configuration happens at two levels:

1. **BGP Protocol Level:** Enable the RT-Constraint family
2. **Routing Options:** Configure automatic RT membership

### Step-by-Step Configuration

#### Step 1: Enable RT-Constraint Family on PE Routers

```
[edit protocols bgp]
set group RR-CLIENTS family route-target

[edit protocols bgp group RR-CLIENTS]
set type internal
```

```

set local-address 10.0.0.1
set neighbor 10.0.0.100 description "Route-Reflector-1"
set neighbor 10.0.0.101 description "Route-Reflector-2"

```

**Why here?** The RT-Constraint family must be negotiated as a BGP capability during session establishment.

## Step 2: Configure Automatic RT Advertisement on PE

```

[edit routing-options]
set route-target-filter advertise-default

[edit routing-instances VPN-A]
set vrf-target target:65000:100

```

**Pattern Recognition:** When you configure a `vrf-target` in a routing instance, Junos automatically:

1. Imports routes with that RT
2. Exports routes with that RT
3. With `advertise-default`, advertises RT membership to RRs

## Step 3: Configure Route Reflector for RT Filtering

```

[edit protocols bgp group PE-CLIENTS]
set type internal
set local-address 10.0.0.100
set cluster 10.0.0.100
set family inet-vpn unicast
set family route-target

[edit routing-options]
set route-target-filter advertise-default

```

## Step 4: Advanced RT Filter Configuration

For more granular control:

```

[edit routing-options route-target-filter]
set 65000:100 local      # Accept locally configured VRFs
set 65000:200 neighbor    # Accept from BGP neighbors
set 65000:300 both        # Accept from both local and neighbors

```

## Complete Reference Configuration

### PE Router Configuration:

```

## BGP Configuration
protocols {
    bgp {
        group RR-CLIENTS {
            type internal;
            local-address 10.0.0.1;
            family inet-vpn {
                unicast;
            }
            family route-target; ## Critical for RT filtering
            neighbor 10.0.0.100 {
                description "Primary-RR";
            }
            neighbor 10.0.0.101 {
                description "Backup-RR";
            }
        }
    }
}

## Routing Options
routing-options {
    autonomous-system 65000;
    route-target-filter {
        advertise-default; ## Automatically advertise RT membership
    }
}

```

```

## VRF Configuration
routing-instances {
    CUSTOMER-A {
        instance-type vrf;
        vrf-target target:65000:100; ## This RT will be advertised
        vrf-table-label;
    }
}

```

### Route Reflector Configuration:

```

protocols {
    bgp {
        group PE-CLIENTS {
            type internal;
            local-address 10.0.0.100;
            cluster 10.0.0.100;
            family inet-vpn {
                unicast;
            }
            family route-target; ## Process RT constraints
            neighbor 10.0.0.1 {
                description "PE-1";
            }
            neighbor 10.0.0.2 {
                description "PE-2";
            }
        }
    }
}

routing-options {
    autonomous-system 65000;
    route-target-filter {
        advertise-default;
    }
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### 1. Verify RT-Constraint Family Negotiation

```

user@PE1> show bgp neighbor 10.0.0.100
Peer: 10.0.0.100+179 AS 65000 Local: 10.0.0.1+54321 AS 65000
    Type: Internal      State: Established
    NLRI advertised by peer: inet-vpn-unicast route-target
    NLRI for this session: inet-vpn-unicast route-target

```

**What to look for:** Both "advertised by peer" and "for this session" should show route-target

#### 2. Check RT Membership Advertisements

```

user@PE1> show route advertising-protocol bgp 10.0.0.100 table bgp.rttarget.0

bgp.rttarget.0: 2 destinations, 2 routes (2 active)
  Prefix          Nexthop          MED     Lclpref   AS path
  65000:0:65000:100/96
*           Self                  100       I
  65000:0:65000:200/96
*           Self                  100       I

```

#### 3. Verify Received RT Constraints on RR

```

user@RR> show route table bgp.rttarget.0

bgp.rttarget.0: 8 destinations, 8 routes (8 active)
+ = Active Route

+ 65000:0:65000:100/96

```

```
*[BGP/170] 00:05:23, localpref 100
  AS path: I
  > to 10.0.0.1 via ge-0/0/0.0
+ 65000:0:65000:200/96
*[BGP/170] 00:05:23, localpref 100
  AS path: I
  > to 10.0.0.2 via ge-0/0/0.0
```

## Common Troubleshooting Scenarios

### Scenario 1: PE Not Receiving Expected VPN Routes

#### Symptom:

```
user@PE1> show route table VPN-A.inet.0
# No routes or missing expected routes
```

#### Diagnostic Commands:

```
user@PE1> show route advertising-protocol bgp 10.0.0.100 table bgp.rttarget.0
# Empty or missing RT advertisements

user@PE1> show bgp neighbor 10.0.0.100 | match route-target
# NLRI for this session: inet-vpn-unicast
# Missing route-target family!
```

**Cause:** RT-Constraint family not negotiated with RR

#### Solution:

```
[edit protocols bgp group RR-CLIENTS]
set family route-target
commit
clear bgp neighbor 10.0.0.100
```

### Scenario 2: RT Advertisements Not Being Sent

#### Symptom:

```
user@RR> show route receive-protocol bgp 10.0.0.1 table bgp.rttarget.0
# No RT constraints received from PE
```

#### Diagnostic Commands:

```
user@PE1> show configuration routing-options route-target-filter
# Configuration missing or incomplete

user@PE1> show route table bgp.rttarget.0 hidden
# Routes may be hidden due to policy
```

**Cause:** Missing advertise-default configuration

#### Solution:

```
[edit routing-options]
set route-target-filter advertise-default
commit
```

### Scenario 3: RR Still Sending All Routes Despite RT Filtering

#### Symptom:

```
user@PE1> show bgp summary
Groups: 1 Peers: 1 Down peers: 0
Table          Tot Paths   Act Paths Suppressed
bgp.l3vpn.0      5000        5000         0 # Too many routes!
```

#### Diagnostic Commands:

```

user@RR> show route table bgp.rtarget.0 | match "65000:100|from PE1"
# Verify RT membership is received

user@RR> show configuration protocols bgp group PE-CLIENTS
# Check if family route-target is configured

```

**Cause:** Route Reflector not processing RT constraints

**Solution:**

```

[edit protocols bgp group PE-CLIENTS]
set family route-target
commit

```

## Scenario 4: Partial RT Filtering (Some Routes Missing)

**Symptom:**

```

user@PE1> show route table VPN-A.inet.0
# Some expected routes present, others missing

```

**Diagnostic Commands:**

```

user@PE1> show route table bgp.rtarget.0 detail
65000:0:65000:100/96 (1 entry)
  State: <FlashAll>
  Route Preference: 170
  Source: 10.0.0.1

user@PE1> show route table bgp.l3vpn.0 extensive | match "Communities|prefix"
# Check if routes have expected RT communities

```

**Cause:** Multiple RTs configured but not all advertised

**Solution:**

```

[edit routing-instances VPN-A]
set vrf-target target:65000:100
set vrf-target target:65000:101 # Add all required RTs

# Or manually configure RT filter
[edit routing-options route-target-filter]
set 65000:100 local
set 65000:101 local
commit

```

# Module 14: Layer 3 VPNs and Internet Access


## Part 1: The Conceptual Lecture (The Why)

### The Fundamental Challenge

Imagine you're a bank with 100 branches, all connected via MPLS L3VPN. Each branch needs to:


1. Communicate with headquarters (via VPN)
2. Access the public Internet (for web services)

The challenge: VPN traffic travels in a private, isolated routing table (VRF), while Internet traffic needs to reach the global routing table. How do you provide Internet access to VPN sites without compromising security or creating routing loops?



## Three Methods of Internet Access in L3VPNs


### Method 1: Central Internet Gateway (Shared PE)



All Internet traffic flows through PE-Central

**Concept:** One PE router acts as the gateway between VPN and Internet. It has both VRF and global routing table access.


### Method 2: Internet Access at Each PE (Distributed)



Each PE has its own Internet connection

**Concept:** Every PE router has direct Internet access. Local Internet breakout at each site.

### Method 3: Internet VRF (Separate VPN for Internet)



**Concept:** Create a special "Internet VRF" that all sites can access, maintaining VPN isolation.

## Route Leaking Fundamentals

The core mechanism is "route leaking" - selectively copying routes between routing tables:

VRF Routing Table	Global Routing Table
+-----+	+-----+
10.1.0.0/24	0.0.0.0/0 (ISP)
10.2.0.0/24	8.8.8.8/32
192.168.0.0/16	200.1.1.0/24
+-----+	+-----+
Private	Public

Route Leaking allows controlled exchange

## NAT Considerations

Since VPN sites typically use RFC1918 addresses (10.x, 172.16.x, 192.168.x), Network Address Translation (NAT) is essential:

Before NAT:	After NAT:
Src: 10.1.1.100	Src: 200.1.1.50
Dst: 8.8.8.8	Dst: 8.8.8.8
(Private)	(Public)

## Part 2: The Junos CLI Masterclass (The How)

### Method 1: Central Internet Gateway Configuration

This is the most common deployment model. We'll configure a PE router to provide Internet access to VPN sites.

#### Step 1: Configure the VRF with Default Route Generation

```
[edit routing-instances CUSTOMER-A]
set instance-type vrf
```

```

set vrf-target target:65000:100
set vrf-table-label

## Generate default route in VRF pointing to global table
set routing-options static route 0.0.0.0/0 next-table inet.0

```

**Why next-table?** This creates a route that "jumps" from VRF to global table.

## Step 2: Configure Reverse Route from Global to VRF

```

[edit routing-options]
set static route 10.1.0.0/16 next-table CUSTOMER-A.inet.0

## Or use RIB groups for dynamic import
[edit routing-options]
set rib-groups VRF-TO-GLOBAL import-rib [ CUSTOMER-A.inet.0 inet.0 ]

```

## Step 3: Configure NAT for Internet-Bound Traffic

```

[edit security nat source]
set rule-set VRF-TO-INTERNET from routing-instance CUSTOMER-A
set rule-set VRF-TO-INTERNET to interface ge-0/0/0.0 # Internet-facing
set rule-set VRF-TO-INTERNET rule NAT-RULE match source-address 10.0.0.0/8
set rule-set VRF-TO-INTERNET rule NAT-RULE then source-nat interface

```

## Step 4: Configure Firewall Filter (Security)

```

[edit firewall family inet filter VRF-INTERNET-FILTER]
set term ALLOW-ESTABLISHED from protocol tcp
set term ALLOW-ESTABLISHED from tcp-established
set term ALLOW-ESTABLISHED then accept

set term ALLOW-DNS from protocol udp
set term ALLOW-DNS from destination-port 53
set term ALLOW-DNS then accept

set term ALLOW-HTTP from protocol tcp
set term ALLOW-HTTP from destination-port [ 80 443 ]
set term ALLOW-HTTP then accept

set term DEFAULT then reject

[edit interfaces ge-0/0/1 unit 100] # VRF interface
set family inet filter output VRF-INTERNET-FILTER

```

## Method 2: Internet VRF Configuration

### Step 1: Create Internet VRF

```

[edit routing-instances INTERNET-VRF]
set instance-type vrf
set vrf-target target:65000:999
set vrf-table-label

## Import Internet routes from global table
set routing-options static route 0.0.0.0/0 next-hop 200.1.1.1

```

### Step 2: Configure Route Leaking Between VRFs

```

## Create policy to export default route to customer VRFs
[edit policy-options policy-statement EXPORT-DEFAULT]
set term 1 from route-filter 0.0.0.0/0 exact
set term 1 then community add INTERNET-RT
set term 1 then accept

[edit policy-options community INTERNET-RT]
set members target:65000:999

## Apply to Internet VRF
[edit routing-instances INTERNET-VRF]
set vrf-export EXPORT-DEFAULT

```

### Step 3: Import Internet Routes in Customer VRF

```
[edit routing-instances CUSTOMER-A]
set vrf-target target:65000:100
set vrf-target import target:65000:999 # Import Internet routes

## Or use vrf-import policy for more control
[edit policy-options policy-statement IMPORT-VPN-AND-INTERNET]
set term VPN from community VPN-RT
set term VPN then accept
set term INTERNET from community INTERNET-RT
set term INTERNET from route-filter 0.0.0.0/0 exact
set term INTERNET then accept
set term DEFAULT then reject

[edit routing-instances CUSTOMER-A]
set vrf-import IMPORT-VPN-AND-INTERNET
```

## Complete Reference Configuration

### Central Internet Gateway PE:

```
## Interfaces
interfaces {
    ge-0/0/0 {
        description "To-Internet";
        unit 0 {
            family inet {
                address 200.1.1.2/30;
            }
        }
    }
    ge-0/0/1 {
        description "To-CE";
        unit 100 {
            family inet {
                address 172.16.1.1/30;
            }
        }
    }
    lo0 {
        unit 0 {
            family inet {
                address 10.0.0.1/32;
            }
        }
    }
}

## MPLS and LDP
protocols {
    mpls {
        interface all;
    }
    ldp {
        interface all;
    }
}

## BGP Configuration
protocols {
    bgp {
        group IBGP-RR {
            type internal;
            local-address 10.0.0.1;
            family inet-vpn unicast;
            neighbor 10.0.0.100;
        }
        group EBGP-INTERNET {
            type external;
            peer-as 701;
            neighbor 200.1.1.1;
        }
    }
}
```

```

## VRF Configuration
routing-instances {
    CUSTOMER-A {
        instance-type vrf;
        interface ge-0/0/1.100;
        vrf-target target:65000:100;
        vrf-table-label;
        routing-options {
            static {
                route 0.0.0.0/0 next-table inet.0;
            }
        }
        protocols {
            bgp {
                group CE {
                    type external;
                    peer-as 65001;
                    neighbor 172.16.1.2;
                }
            }
        }
    }
}

## Global Routing Options
routing-options {
    autonomous-system 65000;
    static {
        route 10.0.0.0/8 next-table CUSTOMER-A.inet.0;
    }
}

## NAT Configuration
security {
    nat {
        source {
            rule-set VRF-T0-INTERNET {
                from routing-instance CUSTOMER-A;
                to interface ge-0/0/0.0;
                rule NAT-ALL {
                    match {
                        source-address 0.0.0.0/0;
                    }
                    then {
                        source-nat {
                            interface;
                        }
                    }
                }
            }
        }
    }
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### 1. Verify Default Route in VRF

```

user@PE1> show route table CUSTOMER-A.inet.0 0.0.0.0/0

CUSTOMER-A.inet.0: 15 destinations, 15 routes
+ = Active Route

+ 0.0.0.0/0
  *[Static/5] 00:10:15
    > to table inet.0

```

#### 2. Check NAT Translations

```

user@PE1> show security nat source summary
Total pools: 1

```

Pool name	Address range	Port
interface-pool	200.1.1.2-200.1.1.2	1024-63487
Total rules: 1		
Rule name	Rule set	From
NAT-ALL	VRF-TO-INTERNET	CUSTOMER-A
Action		
interface		

### 3. Verify Traffic Flow

```
user@PE1> show security flow session source-prefix 10.1.1.0/24

Session ID: 30001, Policy name: default-policy/2
In: 10.1.1.100/45231 --> 8.8.8.8/443;tcp
Out: 8.8.8.8/443 --> 200.1.1.2/45231;tcp
```

## Common Troubleshooting Scenarios

### Scenario 1: No Internet Connectivity from VPN Site

**Symptom:**

```
user@CE> ping 8.8.8.8 routing-instance CUSTOMER-A
PING 8.8.8.8: 56 data bytes
^C
--- 8.8.8.8 ping statistics ---
5 packets transmitted, 0 packets received, 100% packet loss
```

**Diagnostic Commands:**

```
user@PE1> show route table CUSTOMER-A.inet.0 0.0.0.0/0
# No default route present

user@PE1> show route forwarding-table vpn CUSTOMER-A
# Check if next-table route is in forwarding table
```

**Cause:** Missing default route in VRF

**Solution:**

```
[edit routing-instances CUSTOMER-A routing-options static]
set route 0.0.0.0/0 next-table inet.0
commit
```

### Scenario 2: Internet Traffic Works but Return Traffic Fails

**Symptom:**

```
user@CE> traceroute 8.8.8.8
traceroute to 8.8.8.8
1 172.16.1.1 (172.16.1.1) 1.123 ms
2 200.1.1.1 (200.1.1.1) 2.456 ms
3 * * * (no response)
```

**Diagnostic Commands:**

```
user@PE1> show security nat source rule all
# Check if NAT is applied

user@PE1> show route table inet.0 10.1.0.0/16
# No route back to VPN subnet
```

**Cause:** Missing return route from global table to VRF

**Solution:**

```
[edit routing-options static]
set route 10.1.0.0/16 next-table CUSTOMER-A.inet.0
commit
```

### Scenario 3: NAT Not Working

**Symptom:**

```
user@PE1> show security nat source translations
# No translations shown

user@PE1> monitor traffic interface ge-0/0/0.0
# Shows private source IPs going to Internet
```

**Diagnostic Commands:**

```
user@PE1> show configuration security nat
# Check NAT configuration

user@PE1> show security nat source pool all
# Verify address pools
```

**Cause:** NAT rule not matching VRF traffic

**Solution:**

```
[edit security nat source rule-set VRF-T0-INTERNET]
set from routing-instance CUSTOMER-A # Must specify VRF
set to interface ge-0/0/0.0
commit
```

## Scenario 4: Routing Loop Between VRF and Global Table

**Symptom:**

```
user@PE1> show route table CUSTOMER-A.inet.0 8.8.8.8
# Route pointing to inet.0

user@PE1> show route table inet.0 8.8.8.8
# Route pointing back to CUSTOMER-A.inet.0
# Loop detected!
```

**Diagnostic Commands:**

```
user@PE1> traceroute 8.8.8.8 routing-instance CUSTOMER-A
# TTL expires, shows same hops repeating

user@PE1> show route table inet.0 | match CUSTOMER-A
# Check for incorrect route imports
```

**Cause:** Incorrect route leaking configuration

**Solution:**

```
## Use specific prefixes, not 0/0 in both directions
[edit routing-instances CUSTOMER-A routing-options static]
delete route 0.0.0.0/0 next-table inet.0

[edit routing-instances CUSTOMER-A routing-options static]
set route 0.0.0.0/0 next-hop 200.1.1.1

## Or use policy to control what gets leaked
[edit policy-options policy-statement LEAK-T0-GLOBAL]
set term 1 from route-filter 10.0.0.0/8 orlonger
set term 1 then accept
set term 2 then reject
commit
```

---

## Module 15: Inter-AS Layer 3 VPNs

### Part 1: The Conceptual Lecture (The Why)

#### The Fundamental Business Problem

Imagine two scenarios:

- Merger/Acquisition:** Company A (AS 65001) acquires Company B (AS 65002). Both have existing MPLS networks. How do you connect their L3VPNs without rebuilding everything?
- Multi-Provider Services:** A multinational corporation needs VPN services across three continents, each served by different providers. How do you create one seamless VPN across multiple autonomous systems?



## The Technical Challenge

MPLS VPNs were designed to work within a single AS because:

- BGP VPN routes (VPNv4) are typically not advertised via eBGP
- MPLS labels have only local significance
- Route Targets might overlap between providers

## Three Inter-AS Options

The IETF defined three options (RFC 4364), nicknamed Option A, B, and C:

Option A: Back-to-Back VRFs (VRF-to-VRF)

Simple but doesn't scale

Option B: ASBR-to-ASBR VPNv4 Exchange


Moderate complexity, good scale

Option C: Multihop eBGP Between RRs

Complex but most scalable

## Option A: Back-to-Back VRFs

**Concept:** ASBRs treat inter-AS connections as customer connections.



One sub-interface per VPN per ASBR link

### How it works:


- Each ASBR creates VRFs for each VPN
- ASBRs peer using IPv4 (not VPNv4) per VRF
- Standard routing protocols (BGP/OSPF/Static) between ASBRs

### Characteristics:

- ✓ Simple to understand and configure
- ✓ Complete isolation between VPNs
- ✗ Requires configuration for each VPN on ASBRs
- ✗ ASBRs must maintain VRF state

## Option B: ASBR-to-ASBR VPNv4

**Concept:** ASBRs exchange VPNv4 routes directly but don't maintain VRFs.



### Label Stack Operation:

```

Packet from CE1 to CE2:
At PE1: [VPN Label][Transport Label][IP Packet]
At ASBR1: [New VPN Label][IP Packet] # Swap VPN label
At ASBR2: [VPN Label][Transport Label][IP Packet]
At PE2: [IP Packet] # Pop all labels


```

#### Characteristics:

- ✓ ASBRs don't need VRF configuration
- ✓ Scales better than Option A
- ✓ Single BGP session handles all VPNs
- ✗ ASBRs must process all VPNv4 routes

### Option C: Multihop eBGP (RR-to-RR)

**Concept:** Route Reflectors exchange VPNv4 routes; ASBRs only handle labeled IPv4.



#### Three-Label Stack:

```

At Ingress PE:
[VPN Label][BGP Label][IGP Label][IP Packet]
|           |           |
|           |           +-- To reach ASBR1
|           +-- To reach PE2 (via ASBR2)
+-- For the VPN at PE2

```

#### Characteristics:

- ✓ Most scalable solution
- ✓ ASBRs only handle infrastructure routes
- ✓ Clean separation of control and data planes
- ✗ Most complex to configure
- ✗ Requires BGP labeled unicast

## Part 2: The Junos CLI Masterclass (The How)

### Option A Configuration

#### ASBR1 (AS 65001) Configuration:

```

## Create VRFs for each customer
[edit routing-instances]
set CUSTOMER-A instance-type vrf
set CUSTOMER-A interface ge-0/0/1.100 # To ASBR2
set CUSTOMER-A vrf-target target:65001:100
set CUSTOMER-A vrf-table-label

set CUSTOMER-B instance-type vrf
set CUSTOMER-B interface ge-0/0/1.200 # To ASBR2
set CUSTOMER-B vrf-target target:65001:200
set CUSTOMER-B vrf-table-label

## Configure sub-interfaces to peer ASBR
[edit interfaces ge-0/0/1]
set vlan-tagging
set unit 100 vlan-id 100
set unit 100 family inet address 172.16.1.1/30
set unit 200 vlan-id 200
set unit 200 family inet address 172.16.2.1/30

## BGP sessions per VRF to remote ASBR
[edit routing-instances CUSTOMER-A protocols bgp group INTER-AS]

```

```

set type external
set peer-as 65002
set neighbor 172.16.1.2

[edit routing-instances CUSTOMER-B protocols bgp group INTER-AS]
set type external
set peer-as 65002
set neighbor 172.16.2.2

```

## Option B Configuration

### ASBR1 Configuration:

```

## Enable VPNv4 family on inter-AS link
[edit protocols bgp group EBGP-INTER-AS]
set type external
set local-address 10.1.1.1
set peer-as 65002
set neighbor 10.2.2.2
set family inet-vpn unicast

## Configure next-hop-self for VPNv4 routes
[edit protocols bgp group IBGP-PE]
set type internal
set local-address 10.1.1.1
set family inet-vpn unicast
set neighbor 10.1.0.0 # Route Reflector
set vpn-apply-export # Apply VRF export policies

## Policy to handle RT community
[edit policy-options policy-statement INTER-AS-EXPORT]
set term 1 from family inet-vpn
set term 1 then next-hop self
set term 1 then accept

[edit protocols bgp group EBGP-INTER-AS]
set export INTER-AS-EXPORT

## Ensure labeled-unicast for transport
[edit protocols mpls]
set interface ge-0/0/0.0 # Inter-AS interface
set label-switched-path T0-REMOTE-ASBR to 10.2.2.2

```

### ASBR2 Configuration (AS 65002):

```

[edit protocols bgp group EBGP-INTER-AS]
set type external
set local-address 10.2.2.2
set peer-as 65001
set neighbor 10.1.1.1
set family inet-vpn unicast

## Accept and readvertise VPNv4 routes
[edit policy-options policy-statement ACCEPT-VPN4]
set term 1 from family inet-vpn
set term 1 then next-hop self
set term 1 then accept

[edit protocols bgp group EBGP-INTER-AS]
set import ACCEPT-VPN4

```

## Option C Configuration (Most Complex)

### Step 1: Configure BGP Labeled Unicast Between ASBRs

#### ASBR1 (AS 65001):

```

## Enable labeled-unicast for PE loopbacks
[edit protocols bgp group EBGP-LABELED]
set type external
set local-address 10.1.1.1
set peer-as 65002
set neighbor 10.2.2.2

```

```

set family inet labeled-unicast

## Advertise PE loopbacks with labels
[edit policy-options policy-statement EXPORT-PE-LOOPBACKS]
set term PE-LOOPBACKS from route-filter 10.1.0.0/24 orlonger
set term PE-LOOPBACKS from protocol bgp
set term PE-LOOPBACKS then accept

[edit protocols bgp group EBGP-LABELED]
set export EXPORT-PE-LOOPBACKS

## Configure MPLS on inter-AS interface
[edit protocols mpls]
set interface ge-0/0/0.0 # To ASBR2

```

## Step 2: Configure Multihop eBGP Between Route Reflectors

### RR1 (AS 65001):

```

## Multihop eBGP session to remote RR
[edit protocols bgp group EBGP-RR]
set type external
set multihop ttl 255
set local-address 10.1.0.100 # RR1 loopback
set peer-as 65002
set neighbor 10.2.0.100 # RR2 loopback
set family inet-vpn unicast

## Policy to handle next-hop resolution
[edit policy-options policy-statement NHS-T0-ASBR]
set term 1 from protocol bgp
set term 1 from rib bgp.l3vpn.0
set term 1 then next-hop 10.1.1.1 # Local ASBR
set term 1 then accept

[edit protocols bgp group EBGP-RR]
set export NHS-T0-ASBR

## Static route to reach remote RR
[edit routing-options static]
set route 10.2.0.100/32 next-hop 10.1.1.1

```

### RR2 (AS 65002):

```

[edit protocols bgp group EBGP-RR]
set type external
set multihop ttl 255
set local-address 10.2.0.100
set peer-as 65001
set neighbor 10.1.0.100
set family inet-vpn unicast

[edit routing-options static]
set route 10.1.0.100/32 next-hop 10.2.2.2

```

## Step 3: Configure Resolution of BGP Next-Hops

### PE1 (AS 65001):

```

## Enable resolution using labeled BGP routes
[edit routing-options]
set resolution rib bgp.l3vpn.0 resolution-ribs inet.3
set resolution rib bgp.l3vpn.0 resolution-ribs inet.0

## Import labeled routes into inet.3
[edit protocols bgp]
set family inet labeled-unicast rib inet.3

```

## Complete Option C Reference Configuration

### Complete ASBR1 Configuration (AS 65001):

```

interfaces {
    ge-0/0/0 {
        description "To-ASBR2";
        unit 0 {
            family inet {
                address 192.168.1.1/30;
            }
            family mpls;
        }
    }
    lo0 {
        unit 0 {
            family inet {
                address 10.1.1.1/32;
            }
        }
    }
}

protocols {
    mpls {
        interface ge-0/0/0.0;
        interface all;
    }
    bgp {
        group EBGP-LABELED {
            type external;
            local-address 10.1.1.1;
            peer-as 65002;
            family inet {
                labeled-unicast {
                    rib {
                        inet.3;
                    }
                }
            }
            export EXPORT-PE-LOOPBACKS;
            neighbor 10.2.2.2;
        }
        group IBGP-RR {
            type internal;
            local-address 10.1.1.1;
            family inet {
                labeled-unicast;
            }
            family inet-vpn {
                unicast;
            }
            neighbor 10.1.0.100; # Local RR
        }
    }
    ldp {
        interface all;
        interface ge-0/0/0.0 {
            disable; # Don't run LDP on inter-AS link
        }
    }
}

policy-options {
    policy-statement EXPORT-PE-LOOPBACKS {
        term PE-LOOPS {
            from {
                route-filter 10.1.0.0/16 orlonger;
                protocol [ bgp direct static ];
            }
            then {
                next-hop self;
                accept;
            }
        }
        term REJECT {
            then reject;
        }
    }
}

```

```

routing-options {
    autonomous-system 65001;
    forwarding-table {
        export LOAD-BALANCE;
    }
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### 1. Verify Inter-AS BGP Sessions

```

# Option A - Check per-VRF sessions
user@ASBR1> show bgp summary instance CUSTOMER-A
Groups: 1 Peers: 1 Down peers: 0
Peer          AS      InPkt     OutPkt   State
172.16.1.2    65002    150       148     Established

# Option B - Check VPNv4 session
user@ASBR1> show bgp neighbor 10.2.2.2
Peer: 10.2.2.2 AS 65002 Local: 10.1.1.1 AS 65001
    Type: External    State: Established
    NLRI advertised by peer: inet-vpn-unicast
    NLRI for this session: inet-vpn-unicast

# Option C - Check labeled-unicast session
user@ASBR1> show bgp neighbor 10.2.2.2 | match "NLRI|State"
    State: Established
    NLRI advertised by peer: inet-labeled-unicast

```

#### 2. Verify VPNv4 Route Exchange

```

# Check received VPNv4 routes
user@ASBR1> show route receive-protocol bgp 10.2.2.2 table bgp.l3vpn.0

bgp.l3vpn.0: 50 destinations
  65002:100:192.168.100.0/24
    *[BGP/170] 00:05:30, MED 0, localpref 100
      AS path: 65002 I
      > to 10.2.2.2 via ge-0/0/0.0

# Check advertised VPNv4 routes
user@ASBR1> show route advertising-protocol bgp 10.2.2.2 table bgp.l3vpn.0

```

#### 3. Verify Label Operations

```

# Option B - Check VPN label allocation
user@ASBR1> show route table mpls.0 protocol bgp

mpls.0: 25 destinations
  300016 (1 entry)
    *[BGP/170] 00:10:00
      > to 10.2.2.2 via ge-0/0/0.0, Pop

# Option C - Check three-label stack
user@PE1> show route table VRF-A.inet.0 192.168.100.0/24 detail
VRF-A.inet.0: 192.168.100.0/24
  Protocol next hop: 10.2.0.1 # Remote PE
  Label operation: Push 300100(top) Push 300200 Push 300300
    VPN Label^    BGP Label^    IGP Label^

```

### Common Troubleshooting Scenarios

#### Scenario 1: Option A - No Routes Exchanged Between ASBRs

##### Symptom:

```

user@ASBR1> show route receive-protocol bgp 172.16.1.2 instance CUSTOMER-A
# No routes received

```

```
user@CE1> ping 192.168.100.1 routing-instance CUSTOMER-A
# No route to host
```

#### Diagnostic Commands:

```
user@ASBR1> show bgp neighbor 172.16.1.2 instance CUSTOMER-A
# State: Active (not Established)

user@ASBR1> show configuration routing-instances CUSTOMER-A
# Check if interface is in VRF

user@ASBR1> monitor traffic interface ge-0/0/1.100
# Check if BGP packets are being sent
```

**Cause:** VRF interface misconfiguration or BGP not configured in VRF

#### Solution:

```
[edit routing-instances CUSTOMER-A]
set interface ge-0/0/1.100 # Add interface to VRF
set protocols bgp group INTER-AS type external
set protocols bgp group INTER-AS peer-as 65002
set protocols bgp group INTER-AS neighbor 172.16.1.2
commit
```

## Scenario 2: Option B - VPNv4 Routes Not Installing

#### Symptom:

```
user@PE1> show route table VRF-A.inet.0 192.168.100.0/24
# Route missing despite being in bgp.l3vpn.0

user@ASBR1> show route table bgp.l3vpn.0 192.168.100.0/24
bgp.l3vpn.0: 65002:100:192.168.100.0/24
 *[BGP/170] 00:15:00
   AS path: 65002 I
     > to 10.2.2.2 via ge-0/0/0.0
# Route present but not imported to VRF
```

#### Diagnostic Commands:

```
user@PE1> show route table bgp.l3vpn.0 hidden extensive
# Check for "unusable" next-hop

user@PE1> show route resolution unresolved
# Next-hop 10.2.2.2 unresolved
```

**Cause:** BGP next-hop not reachable or label path broken

#### Solution:

```
# On ASBR1 – Set next-hop-self for iBGP
[edit protocols bgp group IBGP-PE]
set family inet-vpn unicast next-hop self
commit

# Or create static route to remote ASBR
[edit routing-options static]
set route 10.2.2.2/32 next-hop 192.168.1.2
commit
```

## Scenario 3: Option C - Three-Label Stack Not Forming

#### Symptom:

```
user@PE1> show route forwarding-table vpn VRF-A destination 192.168.100.0/24
Destination      Type  Next hop          Index NhRef
192.168.100.0/24    user  0:10.1.1.1       1048   2
                      Push 300100      # Only VPN label!

# Missing BGP and IGP labels
```

#### Diagnostic Commands:

```
user@PE1> show route 10.2.0.1 table inet.3
# No route to remote PE in inet.3

user@ASBR1> show route protocol bgp table inet.3
# No labeled routes

user@ASBR1> show bgp neighbor 10.2.2.2 | match labeled
# labeled-unicast not negotiated
```

**Cause:** BGP labeled-unicast not configured or not working

#### Solution:

```
# On ASBR1
[edit protocols bgp group EBGP-LABELED]
set family inet labeled-unicast rib inet.3
commit

# On PE1 - Enable recursive resolution
[edit routing-options]
set resolution rib bgp.l3vpn.0 resolution-ribs [ inet.3 inet.0 ]
commit
```

### Scenario 4: Option C - RR Multihop Session Won't Establish

#### Symptom:

```
user@RR1> show bgp neighbor 10.2.0.100
Peer: 10.2.0.100 AS 65002
  Type: External      State: Active  # Not Established
  Last State: Connect
  Last Error: Cease
```

#### Diagnostic Commands:

```
user@RR1> ping 10.2.0.100 source 10.1.0.100
# No route to host

user@RR1> show route 10.2.0.100
# No route found

user@RR1> traceroute 10.2.0.100
# Dies at first hop
```

**Cause:** No IP reachability between RRs (multihop requires underlying connectivity)

#### Solution:

```
# Option 1: Static routes
[edit routing-options static]
set route 10.2.0.100/32 next-hop 10.1.1.1  # Via local ASBR
commit

# Option 2: Use recursive next-hop via ASBRs
[edit protocols bgp group EBGP-RR]
set multihop ttl 255
set local-address 10.1.0.100
set neighbor 10.2.0.100 multihop  # Enable multihop explicitly
commit
```

### Advanced Troubleshooting: Label Stack Analysis

```
# Trace label operations hop-by-hop
user@PE1> traceroute 192.168.100.1 routing-instance VRF-A

# Monitor MPLS packets
user@ASBR1> monitor traffic interface ge-0/0/0.0 matching "mpls"

# Check MPLS forwarding table
user@ASBR1> show route forwarding-table family mpls
```

```

Label          Type  Next hop        Index NhRef
300100        Swap   192.168.1.2    1050   2
                           Swap label 400100

```

```

# Verify end-to-end LSP
user@PE1> ping mpls ldp 10.2.0.1/32

```

These three modules provide a comprehensive foundation for MPLS L3VPN internetworking and scaling. Remember that in production networks, always:

1. **Test in lab first** - Inter-AS configurations can be complex
2. **Document everything** - Especially RT values and label allocations
3. **Monitor continuously** - Set up BGP session monitoring and label usage alerts
4. **Plan for growth** - Choose the option that matches your scale requirements

## Module 17: Troubleshooting Layer 3 VPN—Overview

### Part 1: The Conceptual Lecture (The Why)

#### The Fundamental Problem

Imagine you've built a complex L3VPN network connecting hundreds of customer sites. Suddenly, Site A cannot reach Site B. The customer is frustrated, and you need to find the problem quickly. But where do you start? With multiple layers of protocols (MPLS, BGP, IGP, PE-CE routing), dozens of routers, and thousands of routes, finding the issue is like finding a needle in a haystack.

L3VPN troubleshooting requires a **systematic approach** because failures can occur at multiple layers:

```

Application Layer      | "Users can't access the application"
                         ▼
IP Layer               | "Ping fails between sites"
                         ▼
L3VPN Layer            | "VPN routes not propagating"
                         ▼
MPLS Layer             | "Labels not being imposed/swapped"
                         ▼
Transport Layer        | "Physical link down"

```


#### The L3VPN Troubleshooting Methodology

Think of L3VPN troubleshooting like diagnosing a car that won't start. You follow a logical sequence:

1. **Is there fuel?** (Are routes present?)
2. **Is the fuel reaching the engine?** (Are routes being advertised?)
3. **Is the engine getting spark?** (Are labels working?)
4. **Is the transmission engaged?** (Is forwarding working?)

#### The Three Planes of L3VPN

Understanding L3VPN troubleshooting requires knowing the three operational planes:



#### PE-to-CE Verification Methods

PE-to-CE verification is critical because this is where the customer network meets the service provider network. Problems here affect everything downstream.

## Method 1: Bottom-Up Verification

Start from Layer 1 and work up:


```
Physical → Data Link → IP → Routing Protocol → VPN  
↓           ↓           ↓           ↓           ↓  
Link UP? ARP work? Ping CE? Routes learned? In VRF?
```

## Method 2: Top-Down Verification

Start from the VPN service and drill down:

```
VPN Service → Route Tables → Protocol → Connectivity → Physical  
↓           ↓           ↓           ↓           ↓  
Sites talk? Routes in VRF? BGP up? Ping work? Link up?
```

## The Troubleshooting Decision Tree



## Why Systematic Troubleshooting Matters

1. **Efficiency:** Avoids random checking that wastes time
2. **Completeness:** Ensures no layer is overlooked
3. **Documentation:** Creates a trail for future reference
4. **Learning:** Builds pattern recognition for common issues

## Part 2: The Junos CLI Masterclass (The How)

### The L3VPN Troubleshooting Toolkit

In Junos, your primary troubleshooting tools are organized by function:

Verification Commands	Diagnostic Commands	Testing Commands
└── show route	└── show log messages	└── ping
└── show bgp	└── monitor traffic	└── traceroute
└── show ospf	└── show system	└── test
└── show ldp	└── show interfaces	
└── show mpls	└── extensive	
└── show route forwarding-table		

## Step-by-Step Basic L3VPN Troubleshooting

### Step 1: Verify Physical Connectivity

```
user@PE1> show interfaces ge-0/0/0 terse  
Interface       Admin Link Proto   Local                   Remote  
ge-0/0/0        up      up      inet    10.1.1.1/30  
  
## More detailed check  
user@PE1> show interfaces ge-0/0/0.100 extensive | match "link|error|drop"  
Link-level type: Ethernet, MTU: 1518  
Link flags : None
```

```

Input errors:
  Errors: 0, Drops: 0, Framing errors: 0, Policed discards: 0
Output errors:
  Carrier transitions: 0, Errors: 0, Drops: 0, MTU errors: 0

```

#### What to look for:

- Admin: up, Link: up (both must be up)
- No errors or drops
- Correct IP addressing

### Step 2: Verify IP Connectivity to CE

```

user@PE1> ping 10.1.1.2 source 10.1.1.1 routing-instance CUSTOMER-A count 3
PING 10.1.1.2: 56 data bytes
64 bytes from 10.1.1.2: icmp_seq=0 ttl=255 time=1.234 ms
64 bytes from 10.1.1.2: icmp_seq=1 ttl=255 time=1.145 ms
64 bytes from 10.1.1.2: icmp_seq=2 ttl=255 time=1.098 ms

--- 10.1.1.2 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max/stddev = 1.098/1.159/1.234/0.056 ms

```

**Critical:** Always specify the routing-instance for VRF-based pings!

### Step 3: Verify PE-CE Routing Protocol

#### For BGP PE-CE:

```

user@PE1> show bgp summary instance CUSTOMER-A
Groups: 1 Peers: 1 Down peers: 0
Table          Tot Paths  Act Paths Suppressed      History Damp State      Pending
CUSTOMER-A.inet.0
                  10          10          0          0          0          0
Peer             AS       InPkt     OutPkt   OutQ    Flaps Last Up/Dwn State
10.1.1.2        65001    450      445      0       0       3:21:45 Established

```

#### For OSPF PE-CE:

```

user@PE1> show ospf neighbor instance CUSTOMER-A
Address           Interface           State      ID           Pri  Dead
10.1.1.2         ge-0/0/0.100       Full      192.168.1.2  128  36

```

#### For Static Routes:

```

user@PE1> show route protocol static table CUSTOMER-A.inet.0
CUSTOMER-A.inet.0: 15 destinations, 15 routes (15 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

192.168.10.0/24  *[Static/5] 1w2d 03:45:12
                  > to 10.1.1.2 via ge-0/0/0.100

```

### Step 4: Verify Routes in VRF Table

```

user@PE1> show route table CUSTOMER-A.inet.0 protocol bgp

CUSTOMER-A.inet.0: 25 destinations, 30 routes (25 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

172.16.1.0/24    *[BGP/170] 00:45:23, localpref 100
                  AS path: 65001 I, validation-state: unverified
                  > to 10.1.1.2 via ge-0/0/0.100
172.16.2.0/24    *[BGP/170] 00:45:23, localpref 100
                  AS path: 65001 I, validation-state: unverified
                  > to 10.1.1.2 via ge-0/0/0.100

```

### Step 5: Verify VPNv4 Route Advertisement

```

user@PE1> show route advertising-protocol bgp 192.168.100.1 table CUSTOMER-A.inet.0 detail

```

```

CUSTOMER-A.inet.0: 25 destinations, 30 routes (25 active, 0 holddown, 0 hidden)
* 172.16.1.0/24 (1 entry, 1 announced)
BGP group IBGP type Internal
  Route Distinguisher: 65000:1
  VPN Label: 299776
  Nexthop: Self
  Flags: Nexthop Change
  AS path: [65000] 65001 I
  Communities: target:65000:1

```

## PE-to-CE Verification Command Reference

```

## Complete PE-to-CE Verification Sequence
## =====

## 1. Interface Status
show interfaces <interface> terse
show interfaces <interface> extensive | match "error|drop|flap"

## 2. ARP/ND Resolution
show arp interface <interface> vpn <vpn-name>
show ipv6 neighbors interface <interface> vpn <vpn-name>

## 3. Basic Connectivity
ping <ce-ip> source <pe-ip> routing-instance <vrf-name>
ping <ce-ip> source <pe-ip> routing-instance <vrf-name> size 1500 do-not-fragment

## 4. Routing Protocol Status
## For BGP:
show bgp summary instance <vrf-name>
show bgp neighbor <ce-ip> instance <vrf-name>

## For OSPF:
show ospf neighbor instance <vrf-name>
show ospf interface instance <vrf-name>
show ospf database instance <vrf-name>

## For RIP:
show rip neighbor instance <vrf-name>
show rip statistics instance <vrf-name>

## For EIGRP (if supported):
show eigrp neighbor instance <vrf-name>
show eigrp topology instance <vrf-name>

## 5. Route Learning Verification
show route receive-protocol bgp <ce-ip> table <vrf-name>.inet.0
show route protocol ospf table <vrf-name>.inet.0
show route protocol direct table <vrf-name>.inet.0

## 6. Route Installation Verification
show route table <vrf-name>.inet.0 <prefix> detail
show route forwarding-table vpn <vrf-name> destination <prefix>

## 7. VPN Label Verification
show route table <vrf-name>.inet.0 <prefix> detail | match label
show mpls route vpn <vrf-name>

## 8. Export Policy Verification
show policy <export-policy-name>
test policy <export-policy-name> <prefix>

```

## Advanced Diagnostic Commands

```

## Hidden Commands for Deep Troubleshooting
## =====

## View BGP RIB-IN before policy processing
user@PE1> show route receive-protocol bgp 10.1.1.2 table CUSTOMER-A.inet.0 hidden all

## View BGP RIB-OUT after policy processing
user@PE1> show route advertising-protocol bgp 192.168.100.1 table CUSTOMER-A.inet.0 detail hidden

## Check VRF route resolution

```

```

user@PE1> show route resolution table CUSTOMER-A.inet.0 unresolved

## View VRF-specific forwarding table
user@PE1> show route forwarding-table vpn CUSTOMER-A extensive

## Monitor real-time route changes
user@PE1> monitor route table CUSTOMER-A.inet.0

```

## Part 3: Verification & Troubleshooting (The What-If)

### Standard Verification Checklist

#### Healthy L3VPN State Indicators

```

## 1. PE-CE Interface Status
user@PE1> show interfaces ge-0/0/0.100 terse
Interface          Admin Link Proto   Local           Remote
ge-0/0/0.100       up     up    inet   10.1.1.1/30
                           mpls
                                         ## Important for some designs

## 2. PE-CE BGP Session (Good State)
user@PE1> show bgp neighbor 10.1.1.2 instance CUSTOMER-A | match State
Type: External      State: Established    ## Must be Established
Last State: OpenConfirm  Last Event: RecvKeepAlive

## 3. Routes Learned from CE (Good State)
user@PE1> show route summary table CUSTOMER-A.inet.0
CUSTOMER-A.inet.0: 25 destinations, 30 routes (25 active, 0 holddown, 0 hidden)
    BGP:    20 routes,    20 active
    Static:  3 routes,    3 active
    Direct:  2 routes,    2 active

## 4. VPNv4 Routes Advertised (Good State)
user@PE1> show route advertising-protocol bgp 192.168.100.1 | match "CUSTOMER-A|entries"
CUSTOMER-A.inet.0: 25 destinations, 30 routes (25 active, 0 holddown, 0 hidden)

```

### Troubleshooting Scenario 1: CE Routes Not Appearing in PE VRF

**Symptom:** Customer reports their routes aren't being learned by the PE

#### Initial Check:

```

user@PE1> show route table CUSTOMER-A.inet.0 172.16.1.0/24

## No output - route not in VRF table

```

#### Diagnostic Process:

```

## Step 1: Check if CE is advertising the route
user@PE1> show route receive-protocol bgp 10.1.1.2 table CUSTOMER-A.inet.0 172.16.1.0/24 hidden all

CUSTOMER-A.inet.0: 25 destinations, 30 routes (25 active, 0 holddown, 5 hidden)
172.16.1.0/24 (1 entry, 0 announced)
  Import: [ CUSTOMER-A-IMPORT ]
  BGP      /-101
    Next hop type: Router
    Address: 0x9458234
    Next-hop reference count: 3
    Source: 10.1.1.2
    Next hop: 10.1.1.2 via ge-0/0/0.100, selected
    State: <Hidden Ext>    ## Hidden! Being filtered
    AS path: 65001 I

```

#### Root Cause Analysis:

```

## Step 2: Check import policy
user@PE1> show configuration policy-options policy-statement CUSTOMER-A-IMPORT
term ALLOW-CUSTOMER {
  from {
    protocol bgp;
    as-path CUSTOMER-AS;
    prefix-list CUSTOMER-PREFIXES; ## Potential issue here
  }
}

```

```

        then accept;
}
term DENY-ALL {
    then reject;
}

## Step 3: Check prefix-list
user@PE1> show configuration policy-options prefix-list CUSTOMER-PREFIXES
172.16.0.0/24;
172.16.2.0/24;
## Missing: 172.16.1.0/24

```

#### Solution:

```

[edit policy-options prefix-list CUSTOMER-PREFIXES]
set 172.16.1.0/24

[edit]
commit and-quit

```

#### Verification:

```

user@PE1> show route table CUSTOMER-A.inet.0 172.16.1.0/24

CUSTOMER-A.inet.0: 26 destinations, 31 routes (26 active, 0 holddown, 0 hidden)

172.16.1.0/24      *[BGP/170] 00:00:15, localpref 100
                      AS path: 65001 I, validation-state: unverified
                      > to 10.1.1.2 via ge-0/0/0.100

```

## Troubleshooting Scenario 2: OSPF Routes Not Installing in VRF

**Symptom:** OSPF adjacency is up but routes aren't in the VRF table

#### Initial Check:

```

user@PE1> show ospf neighbor instance CUSTOMER-A
Address          Interface          State     ID          Pri  Dead
10.1.1.2         ge-0/0/0.100      Full      192.168.1.2  128   35

user@PE1> show route protocol ospf table CUSTOMER-A.inet.0

CUSTOMER-A.inet.0: 3 destinations, 3 routes (3 active, 0 holddown, 0 hidden)
## No OSPF routes!

```

#### Diagnostic Process:

```

## Step 1: Check OSPF database
user@PE1> show ospf database instance CUSTOMER-A

        OSPF database, Area 0.0.0.0
        Type      ID          Adv Rtr          Seq      Age  Opt  Cksum Len
Router *192.168.1.1    192.168.1.1  0x80000003  234  0x22 0x6721  48
Router 192.168.1.2    192.168.1.2  0x80000005  235  0x22 0x5c29  60
Summary 172.16.1.0    192.168.1.2  0x80000001  235  0x22 0x9baa  28
## Routes are in database!

## Step 2: Check OSPF route calculation
user@PE1> show ospf route instance CUSTOMER-A
Topology default Route Table:
Prefix          Path  Route      NH          Metric NextHop      Nexthop
                Type  Type      Type          Interface Address/LSP
192.168.1.1    Intra Router  IP          0  lo0.1
192.168.1.2    Intra Router  IP          1  ge-0/0/0.100  10.1.1.2
172.16.1.0/24  Inter Network IP          2  ge-0/0/0.100  10.1.1.2
## Routes calculated but not installed!

## Step 3: Check VRF configuration
user@PE1> show configuration routing-instances CUSTOMER-A
instance-type vrf;
interface ge-0/0/0.100;
route-distinguisher 65000:1;
vrf-target target:65000:1;

```

```

protocols {
    ospf {
        area 0.0.0.0 {
            interface ge-0/0/0.100;
        }
        ## Missing: export policy to redistribute into VRF!
    }
}

```

**Root Cause:** OSPF routes need explicit redistribution into the VRF RIB

**Solution:**

```

[edit routing-instances CUSTOMER-A protocols ospf]
set rib-group OSPF-T0-VRF

[edit routing-options]
set rib-groups OSPF-T0-VRF import-rib [ CUSTOMER-A.inet.0 inet.0 ]

[edit]
commit

```

## Troubleshooting Scenario 3: Intermittent Connectivity Issues

**Symptom:** Customer reports intermittent connectivity between sites

**Initial Investigation:**

```

user@PE1> ping 172.16.2.1 routing-instance CUSTOMER-A count 10
PING 172.16.2.1: 56 data bytes
 64 bytes from 172.16.2.1: icmp_seq=0 ttl=62 time=12.234 ms
 64 bytes from 172.16.2.1: icmp_seq=1 ttl=62 time=11.998 ms
Request timeout for icmp_seq 2
 64 bytes from 172.16.2.1: icmp_seq=3 ttl=62 time=13.445 ms
Request timeout for icmp_seq 4
Request timeout for icmp_seq 5
 64 bytes from 172.16.2.1: icmp_seq=6 ttl=62 time=12.667 ms

--- 172.16.2.1 ping statistics ---
10 packets transmitted, 6 packets received, 40% packet loss

```

**Diagnostic Process:**

```

## Step 1: Check for route flapping
user@PE1> show route 172.16.2.0/24 table CUSTOMER-A.inet.0 detail | match "age|preference"
    State: <Active Int Ext>
    Age: 3      ## Very young! Route is flapping
    Preference: 170

user@PE1> monitor route table CUSTOMER-A.inet.0 172.16.2.0/24
Dec 10 14:32:45 CUSTOMER-A.inet.0 172.16.2.0/24 Delete
Dec 10 14:32:47 CUSTOMER-A.inet.0 172.16.2.0/24 Add
Dec 10 14:32:52 CUSTOMER-A.inet.0 172.16.2.0/24 Delete
Dec 10 14:32:54 CUSTOMER-A.inet.0 172.16.2.0/24 Add

## Step 2: Check BGP flap statistics
user@PE1> show bgp neighbor 10.1.1.2 instance CUSTOMER-A | match flap
  Last flap event: RecvNotification
  Last flap error: Hold Timer Expired Error
  Flap count: 47    ## High flap count!

```

**Root Cause:** BGP Hold Timer expiring due to congestion or MTU issues

**Solution:**

```

## Increase BGP timers temporarily
[edit routing-instances CUSTOMER-A protocols bgp group CE-PEERS]
set hold-time 180
set keep-alive 60

## Check for MTU issues
[edit interfaces ge-0/0/0 unit 100]
set mtu 1500    ## Ensure consistent MTU

```


```
[edit]
commit
```

## Module 18: Additional Layer 3 VPN Troubleshooting

### Part 1: The Conceptual Lecture (The Why)

#### Advanced L3VPN Failure Modes

While Module 17 covered basic troubleshooting, real-world L3VPN failures often involve complex interactions between multiple protocols. Think of L3VPN as a multi-story building:



Problems can occur at ANY layer or BETWEEN layers!

#### MPLS-Related Problems in L3VPN


MPLS issues are particularly insidious because the control plane might look perfect while the data plane fails silently.

##### Common MPLS Failure Patterns:

1. **Label Distribution Failures**
  - LDP session down but BGP still up
  - Label block exhaustion
  - Label filtering or policies
2. **LSP Failures**
  - RSVP-TE path computation failures
  - Bandwidth reservation issues
  - FRR (Fast Reroute) not working
3. **Label Stack Corruption**
  - Wrong labels being imposed
  - Label TTL issues
  - PHP (Penultimate Hop Popping) problems

#### BGP-Related Problems in L3VPN

BGP issues in L3VPN are complex because we're dealing with multiple address families and route reflection:



#### Forwarding Plane Problems

The forwarding plane is where packets actually move. Problems here mean the control plane shows everything working, but packets get dropped:

Packet Flow Through L3VPN:

CE Router → PE Router → P Router → P Router → PE Router → CE Router  
↓      ↓      ↓      ↓      ↓      ↓  
IP      VRF Lookup      Label Swap      Label Swap      Label Pop      IP


+ Label Push

+ VRF Lookup

#### Common Failure Points:

- VRF not in FIB
- Label not in LFIB
- MTU too small for label stack
- Hardware programming errors

## The Troubleshooting Methodology Pyramid



## Part 2: The Junos CLI Masterclass (The How)

### MPLS Troubleshooting in L3VPN Context

#### Step 1: Verify MPLS Label Distribution

```
## Check LDP status
user@PE1> show ldp overview
Instance: master
  Router ID: 192.168.1.1
  LDP inet enabled
  LDP inet6 disabled
  Neighbors: 4 (4 operational)    ## Should match expected P/PE count
  Sessions: 4 (4 operational)

## Check LDP neighbors
user@PE1> show ldp neighbor
Address          Interface      Label space ID      Hold time
192.168.1.2      ge-0/0/1.0   192.168.1.2:0      14
192.168.1.3      ge-0/0/2.0   192.168.1.3:0      13
192.168.1.4      ge-0/0/3.0   192.168.1.4:0      14
192.168.1.100    lo0.0        192.168.1.100:0    42

## Check label bindings
user@PE1> show ldp database
Input label database, 192.168.1.1:0--192.168.1.2:0
Labels received: 245
  Label      Prefix
  299776    192.168.1.1/32
  299777    192.168.1.2/32
  299778    192.168.1.3/32
  3         192.168.1.100/32  ## Implicit null for PHP

Output label database, 192.168.1.1:0--192.168.1.2:0
Labels advertised: 243
  Label      Prefix
  300100    192.168.1.1/32
  300101    192.168.1.4/32
```

#### Step 2: Verify MPLS LSP Status (for RSVP-TE)

```
user@PE1> show mpls lsp
Ingress LSP: 2 sessions
To          From          State Rt P    ActivePath      LSPname
192.168.1.100 192.168.1.1  Up   0 *           T0-PE2-PRIMARY
192.168.1.101 192.168.1.1  Up   0 *           T0-PE3-PRIMARY

Egress LSP: 2 sessions
```

```

To          From        State   Rt Style Labelin Labelout LSPname
192.168.1.1  192.168.1.100 Up      0  1 FF      3      - FROM-PE2
192.168.1.1  192.168.1.101 Up      0  1 FF      3      - FROM-PE3

## Detailed LSP verification
user@PE1> show mpls lsp name T0-PE2-PRIMARY detail
Ingress LSP: 1 sessions

192.168.1.100
  From: 192.168.1.1, State: Up, ActiveRoute: 0, LSPname: T0-PE2-PRIMARY
  ActivePath: (primary)
  LSPtype: Static Configured, Penultimate hop popping
  LoadBalance: Random
  Encoding type: Packet, Switching type: Packet, GPID: IPv4
  Revert timer: 0
  *Primary           State: Up
    Priorities: 7 0
    SmartOptimizeTimer: 180
    Received RRO (ProtectionFlag 1=Available 2=InUse 4=B/W 8=Node 10=SoftPreempt 20=Node-ID):
      192.168.1.1(Label=299856) 10.0.0.1(Label=299857) 10.0.0.5(Label=3) 192.168.1.100
    Computed ERO (S [L] denotes strict [loose] hops): (CSPF metric: 30)
      10.0.0.1 S 10.0.0.5 S 10.0.0.9 S
    Received Record Route:
      192.168.1.1(flag=0x20) 10.0.0.1(flag=0) 10.0.0.5 192.168.1.100

```

### Step 3: Verify VPN Label Allocation

```

user@PE1> show route table CUSTOMER-A.inet.0 detail | match "label|push"
  Label operation: Push 299776
  Label TTL action: prop-ttl
  Load balance label: Label 299776: None;

user@PE1> show route table mpls.0 protocol vpn
mpls.0: 1547 destinations, 1547 routes (1547 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

299776      *[VPN/0] 1w2d 14:35:22
              to table CUSTOMER-A.inet.0, Pop
299777      *[VPN/0] 1w2d 14:35:22
              to table CUSTOMER-B.inet.0, Pop

```

## BGP Troubleshooting in L3VPN Context

### Step 1: Verify MP-BGP Capabilities

```

user@PE1> show bgp neighbor 192.168.1.100 | match "NLRI|afi|safi|Address"
  Address families configured: inet-vpn-unicast inet6-vpn-unicast
  NLRI that we support:
    NLRI:inet-unicast
    NLRI:inet-vpn-unicast      ## Must be present for L3VPN!
    NLRI:inet6-vpn-unicast
  NLRI that peer supports:
    NLRI:inet-unicast
    NLRI:inet-vpn-unicast
    NLRI:inet6-vpn-unicast
  NLRI for this session:
    NLRI:inet-vpn-unicast      ## Negotiated successfully

```

### Step 2: Verify Route Target Processing

```

## Check route target configuration
user@PE1> show configuration routing-instances CUSTOMER-A | display inheritance
vrf-target {
  ## Inherited from group 'VRF-DEFAULTS'
  import target:65000:1;
  export target:65000:1;
}

## Check if routes have correct communities
user@PE1> show route detail table bgp.l3vpn.0 rd-prefix 65000:1:172.16.1.0/24

bgp.l3vpn.0: 2456 destinations, 4912 routes (2456 active, 0 holddown, 0 hidden)
65000:1:172.16.1.0/24 (2 entries, 1 announced)

```

```

*BGP    Preference: 170/-101
      Route Distinguisher: 65000:1
      Next hop type: Indirect, Index: 0
      Address: 0x9d3f7f4
      Next-hop reference count: 2456
      Source: 192.168.1.100
      Protocol next hop: 192.168.1.100
      VPN Label: 299776
      Indirect next hop: 0xe248200 1048576 INH Session ID: 0x15c
      State: <Secondary Active Int Ext>
      Local AS: 65000 Peer AS: 65000
      Age: 1w2d 14:45:32
      Communities: target:65000:1 target:65000:999 ## Multiple targets!

```

### Step 3: Verify BGP Next-Hop Resolution

```

user@PE1> show route resolution unresolved table bgp.l3vpn.0
Tree Index: 1
Tree Index: 2
      Protocol Nexthop: 192.168.1.200
      Indirect nexthop: 0x0 - INH Session ID: 0x0
      RIB Route: No
      Originating RIB: inet.0
      Metric: 0 Node path count: 0
      Forwarding nexthops: 0
      Nexthop: Discarded ## Problem: Next-hop unreachable!

```

## Forwarding Plane Troubleshooting

### Step 1: Verify FIB Programming

```

user@PE1> show route forwarding-table vpn CUSTOMER-A
Routing table: CUSTOMER-A.inet
Internet:
Destination      Type RtRef Next hop          Type Index NhRef Netif
default          perm   0                  dscd    713     1
10.1.1.0/30      intf   0 10.1.1.1        locl    714     1 ge-0/0/0.100
172.16.1.0/24    user   0                  Push 299776, Push 300200(top)
                                         indr  1048577   2
                                         10.0.0.1       ucst    824     2 ge-0/0/1.0

## Verify hardware programming (on MX platforms)
user@PE1> show pfe route vpn CUSTOMER-A ip prefix 172.16.1.0/24
Slot 0
=====
IPv4 Route Table 0, CUSTOMER-A.inet, 0x80000:
Destination          NH IP Addr      Type      NH ID Interface
-----              -----
172.16.1.0/24        10.0.0.1      Unicast   824    ge-0/0/1.0
MPLS Label Stack: 300200 299776 ## Two labels correctly programmed

```

### Step 2: Test with Labeled Ping

```

## MPLS ping to verify label switching
user@PE1> ping mpls rsvp T0-PE2-PRIMARY
!!!!!
--- lsp ping statistics ---
5 packets transmitted, 5 packets received, 0% packet loss

## Detailed MPLS traceroute
user@PE1> traceroute mpls rsvp T0-PE2-PRIMARY
  ttl  Label Protocol  Address      Previous Hop  Probe Status
  1   299856 RSVP-TE  10.0.0.1    (null)        Success
FEC-Stack-Sent: RSVP
  2   299857 RSVP-TE  10.0.0.5    10.0.0.1      Success
FEC-Stack-Sent: RSVP
  3       3 RSVP-TE   10.0.0.9    10.0.0.5      Egress
FEC-Stack-Sent: RSVP

Path 1 via ge-0/0/1.0 destination 192.168.1.100

```

## Complete Troubleshooting Workflow

```

#!/bin/bash
## L3VPN Comprehensive Troubleshooting Script
## =====

## Function: check_mpls_plane
check_mpls_plane() {
    echo "==== MPLS Plane Verification ==="

    ## Check LDP
    show ldp neighbor | match "operational"
    show ldp session | match "Operational"

    ## Check RSVP if used
    show rsvp neighbor
    show mpls lsp summary

    ## Check label database
    show route table mpls.0 summary
    show mpls label usage
}

## Function: check_bgp_plane
check_bgp_plane() {
    echo "==== BGP Control Plane Verification ==="

    ## Check BGP sessions
    show bgp summary | match "inet-vpn"

    ## Check VPNv4 routes
    show route table bgp.l3vpn.0 summary

    ## Check route targets
    show route community-name target:* table bgp.l3vpn.0 | count
}

## Function: check_forwarding_plane
check_forwarding_plane() {
    echo "==== Forwarding Plane Verification ==="

    ## Check VRF forwarding tables
    show route forwarding-table summary

    ## Check LFIB
    show route forwarding-table family mpls

    ## Check PFE programming (MX only)
    request pfe execute target fpc0 command "show route summary"
}

## Main execution
check_mpls_plane
check_bgp_plane
check_forwarding_plane

```

## Part 3: Verification & Troubleshooting (The What-If)

### Scenario 1: MPLS Transport LSP Failure

**Symptom:** All customer sites connected to PE2 are unreachable

**Initial Discovery:**

```

user@PE1> ping 172.16.2.1 routing-instance CUSTOMER-A
PING 172.16.2.1: 56 data bytes
ping: sendto: No route to host
ping: sendto: No route to host
^C
--- 172.16.2.1 ping statistics ---
2 packets transmitted, 0 packets received, 100% packet loss

user@PE1> show route 172.16.2.0/24 table CUSTOMER-A.inet.0

CUSTOMER-A.inet.0: 25 destinations, 30 routes (25 active, 0 holddown, 0 hidden)

172.16.2.0/24      *[BGP/170] 00:45:12, localpref 100, from 192.168.1.100

```

```
AS path: 65000 I, validation-state: unverified
> to 10.0.0.1 via ge-0/0/1.0, Push 299776, Push 300200(top)
```

#### Diagnostic Process:

```
## Step 1: Route is present, check MPLS transport
user@PE1> show mpls lsp to 192.168.1.100
Ingress LSP: 1 sessions
To          From        State Rt P    ActivePath      LSPname
192.168.1.100  192.168.1.1   Dn     0 *              T0-PE2-PRIMARY

## Step 2: Check why LSP is down
user@PE1> show mpls lsp name T0-PE2-PRIMARY extensive | match "error|fail|down"
State: Dn
LastError: CSPF: no route toward 192.168.1.100[4 times]
CSPF failed: no route toward 192.168.1.100

## Step 3: Check IGP reachability
user@PE1> show route 192.168.1.100 protocol ospf

inet.0: 250 destinations, 250 routes (247 active, 0 holddown, 3 hidden)
## No OSPF route to PE2!

## Step 4: Check OSPF status
user@PE1> show ospf neighbor 10.0.0.1
Address          Interface      State     ID           Pri  Dead
## No output - OSPF neighbor is down!

## Step 5: Check interface
user@PE1> show interfaces ge-0/0/1.0 terse
Interface          Admin Link Proto  Local           Remote
ge-0/0/1.0         up     down
```

**Root Cause:** Physical link failure causing OSPF adjacency loss and LSP failure

#### Solution:

```
## Immediate workaround - use LDP instead of RSVP
[edit routing-instances CUSTOMER-A]
set protocols bgp group IBGP family inet-vpn unicast

[edit protocols mpls]
set traffic-engineering mpls-forwarding

## Long-term fix - repair physical link or configure backup path
[edit protocols mpls label-switched-path T0-PE2-PRIMARY]
set primary MAIN-PATH
set secondary BACKUP-PATH standby
```

## Scenario 2: Route Target Configuration Mismatch

**Symptom:** New customer site can't reach existing sites

#### Initial Check:

```
user@PE3> show route table CUSTOMER-A.inet.0

CUSTOMER-A.inet.0: 3 destinations, 3 routes (3 active, 0 holddown, 0 hidden)
+ = Active Route, - = Last Active, * = Both

10.3.3.0/30      *[Direct/0] 00:45:22
                  > via ge-0/0/0.300
10.3.3.1/32      *[Local/0] 00:45:22
                  Local via ge-0/0/0.300
172.16.3.0/24    *[BGP/170] 00:44:18, localpref 100
                  AS path: 65003 I, validation-state: unverified
                  > to 10.3.3.2 via ge-0/0/0.300
## Only local routes - no routes from other PEs!
```

#### Diagnostic Process:

```
## Step 1: Check what we're advertising
user@PE3> show route advertising-protocol bgp 192.168.1.100 table CUSTOMER-A detail
```

```

CUSTOMER-A.inet.0: 3 destinations, 3 routes (3 active, 0 holddown, 0 hidden)
* 172.16.3.0/24 (1 entry, 1 announced)
  BGP group IBGP type Internal
    Route Distinguisher: 65000:3
    VPN Label: 300001
    Nexthop: Self
    Communities: target:65000:3    ## Wrong RT!

## Step 2: Check VRF configuration
user@PE3> show configuration routing-instances CUSTOMER-A
instance-type vrf;
interface ge-0/0/0.300;
route-distinguisher 65000:3;
vrf-target target:65000:3;    ## Should be 65000:1

```

**Root Cause:** Wrong route target configured on new PE

**Solution:**

```

[edit routing-instances CUSTOMER-A]
delete vrf-target
set vrf-target target:65000:1

[edit]
commit

```

**Verification:**

```

user@PE3> show route table CUSTOMER-A.inet.0 | match "172.16"
172.16.1.0/24      *[BGP/170] 00:00:15, localpref 100, from 192.168.1.1
172.16.2.0/24      *[BGP/170] 00:00:15, localpref 100, from 192.168.1.100
172.16.3.0/24      *[BGP/170] 00:44:18, localpref 100

```

### Scenario 3: MTU Issues with Label Stack

**Symptom:** Large packets fail but small packets work

**Testing:**

```

user@PE1> ping 172.16.2.1 routing-instance CUSTOMER-A size 100
PING 172.16.2.1: 100 data bytes
108 bytes from 172.16.2.1: icmp_seq=0 ttl=62 time=12.234 ms
108 bytes from 172.16.2.1: icmp_seq=1 ttl=62 time=11.876 ms

user@PE1> ping 172.16.2.1 routing-instance CUSTOMER-A size 1472 do-not-fragment
PING 172.16.2.1: 1472 data bytes
ping: sendto: Message too long
ping: sendto: Message too long

```

**Diagnostic Process:**

```

## Step 1: Check interface MTU
user@PE1> show interfaces ge-0/0/1.0 | match mtu
  Protocol inet, MTU: 1500
  Protocol mpls, MTU: 1488, Maximum labels: 3

## Step 2: Calculate overhead
## Base Ethernet: 1500 bytes
## MPLS labels: 2 labels x 4 bytes = 8 bytes
## VPN service: 1500 - 8 = 1492 bytes available
## Customer trying to send 1500 byte packets!

## Step 3: Check core interface MTU
user@P1> show interfaces ge-0/0/0 | match mtu
  Link-level type: Ethernet, MTU: 1514    ## Too small!

```

**Root Cause:** Core network MTU too small for label stack

**Solution:**

```

## On all P and PE routers:
[edit interfaces ge-0/0/0]

```

```

set mtu 9192    ## Jumbo frames

## Or on PE facing customer:
[edit interfaces ge-0/0/0 unit 100]
set family inet mtu 1492    ## Reduce customer-facing MTU

```

## Scenario 4: Hardware Programming Failure

**Symptom:** Control plane perfect but no data plane forwarding

**Diagnostic Process:**

```

## Step 1: Everything looks good in control plane
user@PE1> show route 172.16.2.0/24 table CUSTOMER-A.inet.0 detail
## Routes present with correct labels

user@PE1> show route forwarding-table vpn CUSTOMER-A destination 172.16.2.0/24
Routing table: CUSTOMER-A.inet
Destination      Type RtRef Next hop          Type Index NhRef Netif
172.16.2.0/24    user     0                  Push 299776, Push 300200
                           10.0.0.1        ucst      824      2 ge-0/0/1.0

## Step 2: Check PFE programming (MX specific)
user@PE1> request pfe execute target fpc0 command "show route ip vpn CUSTOMER-A prefix 172.16.2.0/24"
SENT: U kern command: show route ip vpn CUSTOMER-A prefix 172.16.2.0/24
FPC0:
## No output - route not in hardware!

## Step 3: Check for hardware errors
user@PE1> show pfe statistics error
FPC0:
Packet Forwarding Engine Hardware Errors:
  DRAM ECC errors: 15    ## Memory errors!
  Route programming failures: 234

```

**Root Cause:** Hardware memory errors preventing FIB programming

**Solution:**

```

## Immediate workaround - reprogram FIB
user@PE1> request pfe execute target fpc0 command "clear route ip vpn CUSTOMER-A"
user@PE1> restart routing

## Long-term solution - replace faulty linecard
user@PE1> request chassis fpc slot 0 offline
## Physical replacement of FPC
user@PE1> request chassis fpc slot 0 online

```

## Master Troubleshooting Checklist

```

L3VPN_Troubleshooting_Checklist:
Layer_1_Physical:
- Check interface status (show interfaces terse)
- Verify light levels (show interfaces diagnostics optics)
- Check error counters (show interfaces extensive | match error)

Layer_2_DataLink:
- Verify VLAN configuration
- Check ARP/ND resolution
- Verify Ethernet OAM if configured

Layer_3_IP:
- Ping PE-CE interfaces
- Check routing table entries
- Verify MTU consistency

MPLS_Transport:
- Verify LDP/RSVP sessions
- Check label bindings
- Test with MPLS ping/traceroute
- Verify label stack operations

BGP_Control_Plane:
- Check BGP session status

```

- Verify address family negotiation
- Check route target configuration
- Verify route distinguisher uniqueness
- Check next-hop resolution

#### VRF Operations:

- Verify VRF configuration
- Check import/export policies
- Verify route installation
- Check VRF-specific forwarding

#### Forwarding Verification:

- Check FIB programming
- Verify hardware tables (PFE)
- Test with traffic generation
- Monitor packet counters

These two modules provide you with comprehensive knowledge of L3VPN troubleshooting, from basic connectivity issues to complex multi-layer failures. The systematic approach, combined with deep understanding of each protocol layer, will prepare you for any L3VPN troubleshooting scenario in your JNCIE-SP exam and real-world deployments.

## JNCIE-SP Multicast Learning Modules

### Module 19: Multicast Overview

#### Part 1: The Conceptual Lecture (The Why)

##### The Problem Multicast Solves

Imagine you're a teacher in a classroom with 30 students. You need to deliver the same lecture to all of them. You have three options:

- 1. Unicast approach:** Whisper the lecture individually to each student (30 separate conversations)
- 2. Broadcast approach:** Use a megaphone that reaches everyone in the entire school (even those not in your class)
- 3. Multicast approach:** Speak normally to only the students who enrolled in your class

This is exactly the problem IP multicast solves in networking. When one source needs to send identical data to multiple receivers, multicast provides an efficient middle ground between wasteful unicast replication and indiscriminate broadcast flooding.

### Core Concepts and Mechanics

#### What is IP Multicast?


IP multicast is a method of sending IP packets to a group of interested receivers in a single transmission. The source sends one copy of the packet, and the network intelligently replicates it only where paths diverge toward different receivers.

Traditional Unicast (One-to-One):  
 Source → Receiver1 (Copy 1)  
 Source → Receiver2 (Copy 2)  
 Source → Receiver3 (Copy 3)  
 [Source sends 3 copies]

Multicast (One-to-Many):  
 Source → Network → Receiver1  
     └── Receiver2  
     └── Receiver3  
 [Source sends 1 copy, network replicates]

#### Multicast Traffic Flow

The multicast traffic flow involves several key components working together:



## Key Components:

1. **Multicast Source (S)**: Any device sending data to a multicast group
2. **Multicast Group (G)**: A logical identifier (IP address) representing the set of receivers
3. **Multicast Receivers**: Devices interested in receiving traffic for a specific group
4. **First-Hop Router (FHR)**: Router directly connected to the source
5. **Last-Hop Router (LHR)**: Router directly connected to receivers
6. **Distribution Tree**: The path multicast traffic takes through the network

## Multicast Addressing

Multicast uses a special range of IP addresses to identify groups:

**IPv4 Multicast Address Range: 224.0.0.0 - 239.255.255.255**

This range is subdivided into:

Range	Purpose	Scope
224.0.0.0/24	Local Network Control	Link-local
224.0.1.0 - 238.255.255.255	Globally Scoped	Internet-wide
239.0.0.0/8	Administratively Scoped	Private use

### Special Addresses to Remember:

- 224.0.0.1: All hosts on subnet
- 224.0.0.2: All routers on subnet
- 224.0.0.5: OSPF routers
- 224.0.0.6: OSPF designated routers
- 224.0.0.13: PIM routers

### MAC Address Mapping:

IPv4 multicast addresses map to Ethernet MAC addresses using a specific algorithm:

```
Multicast IP: 239.192.1.1
Binary:      11101111.11000000.00000001.00000001
             ↓ (Last 23 bits)
MAC Prefix:  01:00:5E (fixed)
Final MAC:   01:00:5E:40:01:01
```

Note: Because only 23 bits map,  $2^{23} = 8,388,608$  different multicast IPs can map to the same MAC address!

## Reverse Path Forwarding (RPF) Check

The RPF check is crucial for preventing loops in multicast networks. Unlike unicast routing which cares about the destination, multicast routing cares about the source.


### The RPF Principle:

"Accept multicast packets only if they arrive on the interface that would be used to reach the source via unicast routing."

#### RPF Check Process:

1. Packet arrives with source S on interface ge-0/0/1
2. Router checks: "What interface would I use to reach S?"
3. If answer = ge-0/0/1: PASS (forward packet)
4. If answer ≠ ge-0/0/1: FAIL (drop packet)

#### Visual Example:



If Router D receives multicast from S via Router B: RPF PASS ✓

If Router D receives same packet via Router C: RPF FAIL ✗ (loop prevention)

## Multicast Routing Tables

Multicast routers maintain separate tables from unicast routing:

1. **Multicast Routing Information Base (MRIB):**
  - Contains topology information for RPF checks
  - Derived from unicast routing table
2. **Multicast Forwarding Information Base (MFIB):**
  - Contains (S,G) and (\*,G) state entries
  - Determines packet forwarding behavior

### Entry Types:

- **(S,G):** Source-specific entry for source S sending to group G
- **(\*G):** Any-source entry for any source sending to group G

```
Example MFIB Entry:  
(192.168.1.1, 239.1.1.1)  
  Incoming Interface: ge-0/0/0 (RPF interface)  
  Outgoing Interface List: ge-0/0/1, ge-0/0/2, ge-0/0/3  
  
Packet flow: ge-0/0/0 → Router → Replicate to all OIFs
```

## Part 2: The Junos CLI Masterclass (The How)

### Junos Multicast Configuration Hierarchy

The multicast configuration in Junos is organized under several hierarchies:

```
[edit]  
└── protocols {  
    ├── igmp {          # IGMP for IPv4 host management  
    │   └── interface <name>  
    ├── pim {           # PIM for multicast routing  
    │   ├── rp {         # Rendezvous Point configuration  
    │   │   └── interface <name>  
    │   └── msdp { }     # Multicast Source Discovery Protocol  
    └── routing-options {  
        └── multicast { # Global multicast options  
            └── scope <name>  
    }  
    └── interfaces {  
        └── <name> {  
            └── unit <number> {  
                └── family inet {  
                    └── address <address> {  
                        └── multicast; # Enable multicast
```

### Basic Multicast Configuration Pattern

#### Step 1: Enable multicast on interfaces

```
[edit interfaces]  
set ge-0/0/0 unit 0 family inet address 10.1.1.1/24  
set ge-0/0/1 unit 0 family inet address 10.2.2.1/24  
  
## While not always required, explicit multicast enablement ensures proper operation  
[edit protocols]  
set pim interface ge-0/0/0.0  
set pim interface ge-0/0/1.0
```

#### Step 2: Configure PIM (Protocol Independent Multicast)

```
[edit protocols pim]  
## Enable PIM on all participating interfaces  
set interface ge-0/0/0.0 mode sparse      # Sparse mode (most common)  
set interface ge-0/0/1.0 mode sparse  
set interface lo0.0 mode sparse           # Include loopback for RP  
  
## Configure static Rendezvous Point  
set rp static address 192.168.1.1       # RP address  
set rp static group-ranges 239.0.0.0/8  # Groups this RP serves
```

### Step 3: Configure IGMP for receiver interfaces

```
[edit protocols igmp]
## Enable IGMP where receivers connect
set interface ge-0/0/1.0 version 2      # IGMPv2 is most common
set interface ge-0/0/1.0 static group 239.1.1.1 # Static group (optional)
```

### Complete Reference Configuration:

```
## Complete multicast configuration for a typical router
protocols {
    pim {
        ## Define Rendezvous Point
        rp {
            static {
                address 10.0.0.1;          ## RP loopback address
                group-ranges {
                    239.0.0.0/8;         ## Administratively scoped
                }
            }
            local {                  ## This router is the RP
                address 10.0.0.1;
                group-ranges {
                    224.0.0.0/4;       ## Serve all multicast
                }
            }
        }
        ## Enable PIM on interfaces
        interface ge-0/0/0.0 {
            mode sparse;
            version 2;           ## PIMv2 (default)
        }
        interface ge-0/0/1.0 {
            mode sparse;
        }
        interface lo0.0 {
            mode sparse;
        }
    }
    igmp {
        ## Configure IGMP for receiver-facing interfaces
        interface ge-0/0/1.0 {
            version 2;
            ## Optional: static group membership
            static {
                group 239.1.1.1 {
                    source 192.168.1.100; ## Source-specific
                }
            }
        }
    }
}

## Multicast-specific routing options
routing-options {
    multicast {
        ## Define administrative scoping
        scope 239.1.0.0/16 {
            interface ge-0/0/0.0;      ## Boundary interface
        }
    }
}
```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential Verification Commands

#### 1. Verify PIM Neighbors:

```
user@router> show pim neighbors
Instance: PIM.master
B = Bidirectional Capable, G = Generation Identifier
H = Hello Option Holdtime, L = Hello Option LAN Prune Delay,
P = Hello Option DR Priority
```

Interface	IP	V	Mode	Option	Uptime	Neighbor addr
ge-0/0/0.0	4	2	Sparse	HPLG	00:45:23	10.1.1.2
ge-0/0/1.0	4	2	Sparse	HPLG	00:45:20	10.2.2.2

## 2. Verify Multicast Routes:

```
user@router> show multicast route extensive
Family: INET
Group: 239.1.1.1
  Source: 192.168.1.100/32
  Upstream interface: ge-0/0/0.0
  Downstream interface list:
    ge-0/0/1.0
  Session description: Administratively Scoped
  Statistics: 1000 kBps, 1500 pps
  Next-hop ID: 1048576
  Upstream protocol: PIM
  Route state: Active
```

## 3. Verify IGMP Groups:

```
user@router> show igmp group
Interface: ge-0/0/1.0
  Group: 239.1.1.1
  Group mode: Include
  Source: 192.168.1.100
  Last reported by: 10.2.2.100
  Timeout: 174 sec
```

## Common Troubleshooting Scenarios

### Scenario 1: No Multicast Traffic Received

*Symptom:* Receivers not getting multicast traffic despite IGMP join

*Diagnostic Commands:*

```
user@router> show pim join extensive
user@router> show multicast route inactive-paths
user@router> show route forwarding-table family inet multicast
```

*Sample Problem Output:*

```
user@router> show pim join
Instance: PIM.master Family: INET
R = Rendezvous Point Tree, S = Sparse, W = Wildcard

Group: 239.1.1.1
  Source: *
  Flags: sparse
  Upstream interface: Local      ## Problem: No upstream!
```

*Cause:* No route to Rendezvous Point or RPF failure

*Solution:*

```
## Ensure RP is reachable
[edit protocols pim]
set rp static address 10.0.0.1

## Verify unicast routing to RP
[edit protocols ospf]
set area 0.0.0.0 interface lo0.0 passive
```

### Scenario 2: RPF Check Failures

*Symptom:* Multicast packets dropped at router

*Diagnostic Commands:*

```
user@router> show multicast rpf 192.168.1.100
user@router> show pim statistics
```

*Sample Problem Output:*

```
user@router> show multicast rpf 192.168.1.100
Multicast RPF table: inet.0, 12 entries
192.168.1.100/32
  Protocol: OSPF
  Interface: ge-0/0/1.0      ## Wrong interface!
  Neighbor: 10.2.2.2
```

*Cause:* Multicast arriving on different interface than unicast path

*Solution:*

```
## Option 1: Fix unicast routing
[edit protocols ospf]
set area 0.0.0.0 interface ge-0/0/0.0 metric 1

## Option 2: Static multicast RPF route
[edit routing-options]
set multicast rpf-check-policy MULTICAST_RPF

[edit policy-options]
set policy-statement MULTICAST_RPF term 1 from source-address-filter 192.168.1.0/24 exact
set policy-statement MULTICAST_RPF term 1 then accept
```

### Scenario 3: IGMP Version Mismatch

*Symptom:* Hosts can't join multicast groups

*Diagnostic Commands:*

```
user@router> show igmp interface
user@router> monitor traffic interface ge-0/0/1.0 matching igmp
```

*Sample Problem Output:*

```
user@router> show igmp interface
Interface: ge-0/0/1.0
  Querier: 10.2.2.1 (Self)
  Version: 3          ## Router using v3
  Groups: 0          ## No groups joined!
```

*Cause:* Hosts using IGMPv2, router configured for IGMPv3

*Solution:*

```
[edit protocols igmp]
set interface ge-0/0/1.0 version 2
```

### Scenario 4: Multicast Boundary Blocking

*Symptom:* Certain multicast groups not forwarding

*Diagnostic Commands:*

```
user@router> show multicast scope
user@router> show configuration routing-options multicast
```

*Sample Problem Output:*

```
user@router> show multicast scope
Scope           Interfaces
239.0.0.0/8     ge-0/0/0.0    ## Blocks admin-scoped
```

*Cause:* Administrative scoping blocking desired groups

*Solution:*

```

## Adjust scope to allow specific groups
[edit routing-options multicast]
delete scope 239.0.0.0/8
set scope 239.255.0.0/16 interface ge-0/0/0.0

```

## Module 20: Introduction to IGMP

### Part 1: The Conceptual Lecture (The Why)

#### The Problem IGMP Solves

Think of multicast like a newspaper subscription service. The newspaper company (multicast source) wants to deliver papers only to houses that have subscribed. But there's a challenge: How does the delivery truck (router) know which houses on each street want the newspaper?

This is where IGMP (Internet Group Management Protocol) comes in. IGMP is the protocol that allows hosts to tell their local router: "I want to receive this multicast traffic" or "I'm no longer interested."

Without IGMP, routers would have two bad choices:

1. Forward all multicast to all segments (wasteful)
2. Forward no multicast at all (useless)


#### Core Concepts and Mechanics

##### What is IGMP?

IGMP is a communication protocol used between hosts and their immediately neighboring routers. It operates at Layer 3 (Network Layer) and enables hosts to:

- Join multicast groups (subscribe)
- Leave multicast groups (unsubscribe)
- Report their group membership status

IGMP Communication Model:



##### Key IGMP Components:

1. **IGMP Querier**: The router that sends queries (typically the one with lowest IP)
2. **IGMP Report**: Message from host saying "I want group X"
3. **IGMP Query**: Message from router asking "Who wants what groups?"
4. **IGMP Leave**: Message from host saying "I no longer want group X" (v2+)

#### IGMP Versions and Evolution

##### IGMPv1 (RFC 1112) - The Beginning:

The original IGMP was simple but limited:

IGMPv1 Operation:

1. Router: "Anyone want any multicast?" (General Query)
2. Host: "I want 239.1.1.1!" (Report)
3. Router waits... waits... waits... (No explicit leave)
4. After timeout (260 seconds): Assumes host left

##### Problems with IGMPv1:

- No explicit leave mechanism (slow convergence)
- No querier election (multiple routers cause issues)
- Simple but inefficient

## IGMPv2 (RFC 2236) - The Improvement:

IGMPv2 added crucial features:

IGMPv2 Enhancements:

1. Explicit Leave messages
2. Group-specific queries
3. Querier election process
4. Faster leave processing

IGMPv2 Operation:

1. Router: "Anyone want any multicast?" (General Query)
2. Host A: "I want 239.1.1.1!" (Report)
3. Host B: "I'm done with 239.1.1.1" (Leave)
4. Router: "Anyone else want 239.1.1.1?" (Group-Specific Query)
5. If no response: Stop forwarding immediately

## IGMPv3 (RFC 3376) - Source Filtering:

IGMPv3 introduced source-specific multicast (SSM):

IGMPv3 Source Filtering:

- Include Mode: "I want group G only from sources S1, S2"
- Exclude Mode: "I want group G from all sources except S3, S4"

Example:

Host → Router: "I want 239.1.1.1 only from 192.168.1.100"

Result: More efficient, secure multicast delivery

## IGMP Message Types and Formats

### Message Types by Version:

Version	Message Type	Type Code	Purpose
v1	Query	0x11	Router asks for members
v1	Report	0x12	Host joins group
v2	Query	0x11	Router asks for members
v2	v2 Report	0x16	Host joins group
v2	Leave	0x17	Host leaves group
v3	Query	0x11	Router asks (with options)
v3	v3 Report	0x22	Host report (with sources)

### IGMP Message Format (v2):





## IGMP Timers and State Machine

### Critical Timers:

1. **Query Interval:** How often router sends general queries (default: 125 seconds)
2. **Query Response Interval:** Max time host waits before responding (default: 10 seconds)
3. **Group Membership Interval:** Time before group times out (default: 260 seconds)
4. **Last Member Query Interval:** Time between group-specific queries after leave (default: 1 second)

IGMP State Machine:





## Part 2: The Junos CLI Masterclass (The How)

### IGMP Configuration Hierarchy in Junos

```

[edit protocols igmp]
└── interface <interface-name> {
    ├── version <1|2|3>;
    ├── static {
    │   └── group <group-address> {
    │       └── source <source-address>; # v3 only
    │   }
    ├── group-limit <number>;
    ├── query-interval <seconds>;
    ├── query-response-interval <seconds>;
    └── robust-count <number>;
}
└── query-last-member-interval <seconds>;
└── traceoptions {
    └── flag <flag-name>;
}
}

```

### Comprehensive IGMP Configuration Examples

#### Basic IGMPv2 Configuration:

```

[edit protocols igmp]
## Standard IGMPv2 setup for LAN interfaces
set interface ge-0/0/0.0 version 2

## Version 2 is default, but explicit is better
set interface ge-0/0/1.0 {
    version 2;
    ## Adjust timers for faster convergence
    query-interval 30;          ## Query every 30 seconds
    query-response-interval 5;   ## Hosts respond within 5 seconds
    robust-count 2;             ## Robustness variable (packet loss tolerance)
}

```

#### IGMPv3 with Source-Specific Multicast:

```

[edit protocols igmp]
set interface ge-0/0/2.0 {
    version 3;
    ## IGMPv3 allows source filtering
    static {
        group 232.1.1.1 {           ## SSM range (232.0.0.0/8)
            source 10.1.1.100;      ## Specific source
            source 10.1.1.101;      ## Another allowed source
        }
    }
    ## Immediate leave for faster convergence
    immediate-leave;
}

```

#### Advanced IGMP Configuration with Access Control:

```

[edit protocols igmp]
set interface ge-0/0/3.0 {
    version 2;
    ## Limit groups per interface (DoS protection)
    group-limit 10;
    ## Access list for group filtering
    group-policy ALLOWED_GROUPS;
    ## Passive mode (no queries sent)
    passive;
    ## Static groups (always forwarded)
    static {
        group 239.1.1.1;
    }
}

```

```

        group 239.2.2.2;
    }
}

[edit policy-options]
set policy-statement ALLOWED_GROUPS {
    term accept-admin-scoped {
        from {
            route-filter 239.0.0.0/8 orlonger;
        }
        then accept;
    }
    term deny-all {
        then reject;
    }
}

```

#### Complete Production IGMP Configuration:

```

## Full IGMP configuration with all best practices
[edit protocols]
igmp {
    ## Global IGMP parameters
    query-last-member-interval 1;      ## 1 second for fast leave

    ## Subscriber-facing interface
    interface ge-0/0/0.0 {
        version 2;                      ## Most compatible
        query-interval 125;             ## Standard interval
        query-response-interval 10;     ## Standard response
        robust-count 2;                ## Handle packet loss

        ## Immediate leave for single-host segments
        immediate-leave;

        ## Group limiting for security
        group-limit 20;

        ## Access control
        group-policy MULTICAST_POLICY;
    }

    ## Server-facing interface (static groups)
    interface ge-0/0/1.0 {
        version 3;                      ## Support SSM
        passive;                        ## Don't send queries

        static {
            ## Always forward these groups
            group 239.1.1.1 {
                source 192.168.1.10;
            }
            group 239.1.1.2 {
                source 192.168.1.11;
            }
        }
    }

    ## Enable IGMP snooping for switches
    interface ae0.0 {
        version 2;
        ## Proxy reporting for efficiency
        proxy;
    }

    ## Troubleshooting
    traceoptions {
        file igmp-log size 10m files 5;
        flag packets detail;
        flag membership detail;
        flag state detail;
    }
}

## Supporting policy
policy-options {

```

```

policy-statement MULTICAST_POLICY {
    term allow-admin {
        from {
            route-filter 239.0.0.0/8 orlonger;
        }
        then accept;
    }
    term allow-ssm {
        from {
            route-filter 232.0.0.0/8 orlonger;
        }
        then accept;
    }
    term default-deny {
        then reject;
    }
}
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential IGMP Verification Commands

#### 1. Verify IGMP Interface Status:

```

user@router> show igmp interface detail
Physical Interface: ge-0/0/0.0
    Querier: 10.1.1.1 (Self)
    State: Querier
    Version: 2
    Groups: 3
    Query Interval: 125.0
    Query Response Interval: 10.0
    Last Member Query Interval: 1.0
    Robust Count: 2
    Immediate Leave: On

```

#### 2. Verify IGMP Groups:

```

user@router> show igmp group detail
Interface: ge-0/0/0.0
    Group: 239.1.1.1
        Group mode: Exclude (IGMPv2)
        Source list: (empty)
        Last reported by: 10.1.1.100
        Group timeout: 240 seconds
        Type: Dynamic

    Group: 239.2.2.2
        Group mode: Include (IGMPv3)
        Source list:
            192.168.1.10 Timeout: 255
            192.168.1.11 Timeout: 255
        Last reported by: 10.1.1.101

```

#### 3. Verify IGMP Statistics:

```

user@router> show igmp statistics
IGMP packet statistics for all interfaces
IGMP Message Type      Received      Sent
Membership Query       100          5234
V1 Membership Report   0             0
V2 Membership Report   3421         0
V3 Membership Report   234          0
Leave Group            89           0

```

## Common IGMP Troubleshooting Scenarios

### Scenario 1: Hosts Not Joining Groups

*Symptom:* No IGMP groups shown despite hosts attempting to join

*Diagnostic Commands:*

```
user@router> show igmp interface
user@router> monitor traffic interface ge-0/0/0.0 matching "igmp"
user@router> show log messages | match IGMP
```

*Sample Problem Output:*

```
user@router> show igmp interface
## No output - IGMP not enabled!
```

*Cause:* IGMP not configured on interface

*Solution:*

```
[edit protocols]
set igmp interface ge-0/0/0.0
```

### Scenario 2: IGMP Version Incompatibility

*Symptom:* Some hosts can join, others cannot

*Diagnostic Commands:*

```
user@router> show igmp group
user@router> show igmp interface detail
```

*Sample Problem Output:*

```
user@router> monitor traffic interface ge-0/0/0.0 matching igmp
15:32:01.123456 In IP 10.1.1.100 > 224.0.0.22: IGMPv3 Report ## v3 host
15:32:02.234567 In IP 10.1.1.101 > 224.0.0.2: IGMPv2 Report ## v2 host

user@router> show igmp interface
Interface: ge-0/0/0.0
Version: 3 ## Router in v3-only mode!
```

*Cause:* Router using IGMPv3, but some hosts only support IGMPv2

*Solution:*

```
## Use version 2 for compatibility
[edit protocols igmp]
set interface ge-0/0/0.0 version 2

## Or use version 3 with compatibility mode
[edit protocols igmp]
set interface ge-0/0/0.0 version 3
set interface ge-0/0/0.0 version-1-2-compatibility
```

### Scenario 3: Groups Not Timing Out

*Symptom:* Groups remain active even after all hosts leave

*Diagnostic Commands:*

```
user@router> show igmp group detail
user@router> clear igmp membership ## Manual clear
```

*Sample Problem Output:*

```
user@router> show igmp group detail
Interface: ge-0/0/0.0
Group: 239.1.1.1
Type: Static ## Problem: Configured as static!
Timeout: Never
```

*Cause:* Groups configured as static instead of dynamic

*Solution:*

```

## Remove static configuration
[edit protocols igmp interface ge-0/0/0.0]
delete static group 239.1.1.1

## Groups will now be learned dynamically

```

#### Scenario 4: Duplicate Multicast Traffic

*Symptom:* Hosts receiving duplicate multicast packets

*Diagnostic Commands:*

```

user@router> show igmp interface
user@router> show pim neighbors

```

*Sample Problem Output:*

```

user@router> show igmp interface
Interface: ge-0/0/0.0
Querier: 10.1.1.2    ## Not us!

user@router-2> show igmp interface
Interface: ge-0/0/0.0
Querier: 10.1.1.2    ## Also querier!

```

*Cause:* Multiple IGMP queriers on same segment

*Solution:*

```

## Ensure only one querier per segment
## Router with lowest IP becomes querier

## Option 1: Make this router passive
[edit protocols igmp]
set interface ge-0/0/0.0 passive

## Option 2: Adjust IP to be lowest (become querier)
[edit interfaces]
set ge-0/0/0 unit 0 family inet address 10.1.1.1/24

```

## Module 21: Multicast Routing Protocols (PIM)

### Part 1: The Conceptual Lecture (The Why)

#### The Problem PIM Solves

IGMP tells routers which multicast groups are needed on each network segment, but it doesn't tell them HOW to build the delivery tree across multiple routers. Think of it this way:


- **IGMP** = The subscription list (who wants what)
- **PIM** = The delivery route planning (how to get it there)

Protocol Independent Multicast (PIM) creates and maintains the distribution trees that carry multicast traffic from sources to receivers across a routed network.

#### Core PIM Concepts

##### Why "Protocol Independent"?

PIM doesn't care about how you learned your unicast routes (OSPF, IS-IS, BGP, static). It uses whatever unicast routing table exists for its RPF checks. This makes PIM flexible and deployable in any environment.




#### PIM Operating Modes

## PIM Dense Mode (PIM-DM) - The Flood-and-Prune Approach:

Dense mode assumes everyone wants the traffic until they say otherwise:

PIM-DM Operation:

1. Source starts sending to 239.1.1.1
2. Flood to ALL PIM neighbors
3. Routers with no receivers send Prune
4. Prunes time out after 3 minutes
5. Flooding resumes (repeat cycle)




## PIM Sparse Mode (PIM-SM) - The Explicit-Join Approach:

Sparse mode assumes nobody wants traffic until they explicitly request it:

PIM-SM Operation:

1. Receivers join group via IGMP
2. Last-hop router sends PIM Join toward RP
3. Source sends to RP (register process)
4. RP forwards down the shared tree
5. Optional: Switch to shortest path tree

Shared Tree (\*,G):



## The Rendezvous Point (RP)

The RP is the meeting point for sources and receivers in PIM-SM:

Before RP knows about source:  
Receivers: "We want 239.1.1.1" → Join travels to RP  
RP: "Noted, waiting for sources"

When source starts:  
Source → First-Hop Router: "I have data for 239.1.1.1"  
First-Hop → RP: Register message (encapsulated data)  
RP: "Great! I have receivers waiting" → Forwards data  
RP → First-Hop: "Send me native multicast" (Join)


## PIM Message Types

### Core PIM Messages:

Message Type	Purpose	Direction
Hello	Neighbor discovery	All PIM routers
Join/Prune	Build/tear down trees	Toward source/RP
Register	Notify RP of new source	First-hop → RP
Register-Stop	Stop registration	RP → First-hop
Assert	Resolve duplicate forwarders	Between neighbors
Bootstrap	RP discovery	Throughout domain
Candidate-RP	Announce RP candidacy	To BSR

### PIM Message Format Example (Join/Prune):

### PIM Join/Prune Message:



## PIM Tree Types and Transitions

### 1. Rendezvous Point Tree (RPT) / Shared Tree (\*,G):


- All sources for a group share the same tree
- Rooted at the RP
- Suboptimal paths possible

### 2. Shortest Path Tree (SPT) / Source Tree (S,G):

- Unique tree per source
- Rooted at the source
- Optimal paths

#### SPT Switchover Process:

1. Last-hop router receives traffic via RPT
2. Threshold exceeded (first packet or rate-based)
3. Last-hop sends (S,G) Join toward source
4. Receives traffic via SPT
5. Sends (S,G,rpt) Prune toward RP
6. Traffic flows on optimal path



## Part 2: The Junos CLI Masterclass (The How)

### PIM Configuration Hierarchy

```
[edit protocols pim]
  rp {
    local {                      # Rendezvous Point configuration
    static {                      # This router as RP
    auto-rp {                     # Static RP mapping
    bootstrap {                  # Cisco Auto-RP compatibility
    interface <name> {
      mode <sparse|dense>;
      version <1|2>;
      priority <0-65535>;       # DR election priority
      hello-interval <seconds>;
    spt-threshold {              # SPT switchover control
    dense-groups {               # Groups to run in dense mode
    traceoptions {
```

## Comprehensive PIM Configuration Patterns

### Basic PIM Sparse Mode Setup:

```
[edit protocols pim]
## Define static RP for simple deployments
rp {
```

```

        static {
            address 10.0.0.1;          ## RP address (usually loopback)
            group-ranges {
                224.0.0.0/4;         ## All multicast groups
            }
        }
    }

## Enable PIM on all interfaces
interface ge-0/0/0.0 {
    mode sparse;                 ## Sparse mode (recommended)
}
interface ge-0/0/1.0 {
    mode sparse;
}
interface lo0.0 {
    mode sparse;                 ## Include loopback for RP
}

```

#### **Advanced PIM with Bootstrap Router:**

```

[edit protocols pim]
## Bootstrap Router configuration (dynamic RP)
rp {
    bootstrap {
        family inet {
            ## This router as BSR candidate
            bsr-candidate lo0.0 {
                priority 100;      ## Higher = preferred
                hash-mask-length 30;
            }

            ## This router as RP candidate
            rp-candidate lo0.0 {
                group-ranges {
                    239.0.0.0/8;
                }
                priority 100;
                hold-time 150;
            }
        }
    }

    ## Backup static RP
    static {
        address 10.0.0.2 {
            group-ranges {
                224.0.0.0/4;
            }
            override;           ## Use static even if BSR exists
        }
    }
}

## Interface configuration
interface ge-0/0/0.0 {
    mode sparse;
    priority 100;                  ## DR priority (higher wins)
    hello-interval 30;             ## Hello timer
}

```

#### **PIM Source-Specific Multicast (SSM) Configuration:**

```

[edit protocols pim]
## SSM doesn't need an RP
ssm-groups {
    232.0.0.0/8;                 ## IANA SSM range
    239.255.0.0/16;              ## Custom SSM range
}

## No RP needed for SSM groups
rp {
    static {
        address 10.0.0.1 {
            group-ranges {

```

```

        ## Exclude SSM ranges
        224.0.0.0/4 {
            exclude {
                232.0.0.0/8;
                239.255.0.0/16;
            }
        }
    }
}

interface all {                      ## Shortcut for all interfaces
    mode sparse;
}

```

#### Complete Production PIM Configuration:

```

## Comprehensive PIM configuration with all features
[edit protocols pim]
## Rendezvous Point configuration
rp {
    ## Local RP (if this router is RP)
    local {
        address 10.0.0.1;
        group-ranges {
            239.0.0.0/8;          ## Admin-scoped groups
        }
    }

    ## Bootstrap router for dynamic RP selection
    bootstrap {
        family inet {
            bsr-candidate lo0.0 {
                priority 200;
                hash-mask-length 30;
            }
            rp-candidate lo0.0 {
                group-ranges {
                    239.0.0.0/8;
                }
                priority 200;
                hold-time 150;
            }
        }
    }

    ## Static RP as fallback
    static {
        address 10.0.0.2 {
            group-ranges {
                224.0.0.0/4;
            }
        }
    }
}

## SPT switchover control
spt-threshold {
    infinity {                      ## Never switch to SPT
        group-ranges {
            239.1.0.0/16;      ## For these groups
        }
    }
}

## Dense mode for specific groups (rare)
dense-groups {
    224.0.1.0/24;                ## IANA dense groups
}

## Interface configurations
interface ge-0/0/0.0 {
    mode sparse;
    version 2;                  ## PIMv2 (default)
    priority 150;               ## High DR priority
}

```

```

hello-interval 30;
join-prune-interval 60;

## Limit multicast rate
multicast-rate-limit {
    maximum-bandwidth 100m;
}
}

interface ge-0/0/1.0 {
    mode sparse;
    priority 50;          ## Lower DR priority

    ## BFD for fast failure detection
    bfd-liveness-detection {
        minimum-interval 300;
        multiplier 3;
    }
}

interface ge-0/0/2.0 {
    mode sparse;
    ## Passive interface (no PIM messages sent)
    passive;
}

interface lo0.0 {
    mode sparse;
}

## SSM configuration
ssm-groups {
    232.0.0.0/8;          ## Standard SSM
    239.255.0.0/16;       ## Custom SSM
}

## PIM neighbor policy
neighbor-policy ALLOWED_NEIGHBORS;

## Troubleshooting
traceoptions {
    file pim-log size 10m files 5;
    flag hello detail;
    flag join detail;
    flag register detail;
    flag rp detail;
    flag spt detail;
}

## Supporting routing options
[edit routing-options]
multicast {
    ## RPF check exceptions
    rpf-check-policy RPF_EXCEPTIONS;

    ## Multicast forwarding cache
    forwarding-cache {
        timeout 300;
        threshold {
            suppress 5000;
            reuse 4000;
        }
    }
}

## Supporting policies
[edit policy-options]
policy-statement ALLOWED_NEIGHBORS {
    term allow-local {
        from {
            protocol pim;
            neighbor 10.0.0.0/8;
        }
        then accept;
    }
    term deny-all {

```

```

        then reject;
    }
}

policy-statement RPF_EXCEPTIONS {
    term multicast-sources {
        from {
            source-address-filter {
                192.168.1.0/24 exact;
            }
        }
        then {
            rpf-check loose;
        }
    }
}

```

## Part 3: Verification & Troubleshooting (The What-If)

### Essential PIM Verification Commands

#### 1. Verify PIM Neighbors:

```

user@router> show pim neighbors detail
Instance: PIM.master

Interface: ge-0/0/0.0
Address: 10.1.1.2, IPv4, PIM version: 2
Hello Option - Holdtime: 105 seconds, 97 remaining
Hello Option - DR Priority: 100
Hello Option - Generation ID: 1234567890
Hello Option - LAN Prune Delay: delay 500 ms, override 2500 ms
BFD: Enabled, Status: Up
Designated Router
Uptime: 2d 14:23:45

```

#### 2. Verify PIM Join State:

```

user@router> show pim join extensive
Instance: PIM.master Family: INET
R = Rendezvous Point Tree, S = Sparse, W = Wildcard

Group: 239.1.1.1
Source: *
RP: 10.0.0.1
Flags: sparse,rp-tree,wildcard
Upstream protocol: Static
Upstream interface: ge-0/0/0.0
Upstream neighbor: 10.1.1.2
Upstream state: Joined
Downstream neighbors:
    Interface: ge-0/0/1.0
        10.2.2.2 State: Join Flags: SRW Timeout: 180

```

#### 3. Verify RP Information:

```

user@router> show pim rps extensive
Instance: PIM.master
Address family: INET

RP: 10.0.0.1
Learned via: static
Group-Ranges:
    224.0.0.0/4
Active groups using RP:
    239.1.1.1
    239.2.2.2
Register state:
    Limit: 0 kbps
    Current: 145 kbps

```

### Common PIM Troubleshooting Scenarios

### **Scenario 1: PIM Neighbors Not Forming**

*Symptom:* No PIM neighbors visible

*Diagnostic Commands:*

```
user@router> show pim interfaces  
user@router> show pim statistics  
user@router> monitor traffic interface ge-0/0/0.0 matching "pim"
```

*Sample Problem Output:*

```
user@router> show pim interfaces  
## No output!  
  
user@router> show configuration protocols pim  
## PIM not configured
```

*Cause:* PIM not enabled on interfaces

*Solution:*

```
[edit protocols pim]  
set interface ge-0/0/0.0 mode sparse  
set interface lo0.0 mode sparse
```

### **Scenario 2: Multicast Not Flowing Through RP**

*Symptom:* Sources registering but receivers not getting traffic

*Diagnostic Commands:*

```
user@router> show pim rps  
user@router> show pim register  
user@router> show multicast route
```

*Sample Problem Output:*

```
user@router> show pim register  
PIM Register Messages:  
Group          Source        First Hop       State  
239.1.1.1     192.168.1.100  10.1.1.2      Register-Stop  
  
user@router> show pim join  
Group: 239.1.1.1  
  Source: *  
  Upstream interface: None    ## Problem: No upstream!
```

*Cause:* RP not reachable or RPF failure to RP

*Solution:*

```
## Ensure RP is reachable via unicast  
[edit protocols ospf]  
set area 0.0.0.0 interface lo0.0 passive  
  
## Or add static route to RP  
[edit routing-options static]  
set route 10.0.0.1/32 next-hop 10.1.1.2
```

### **Scenario 3: SPT Switchover Not Happening**

*Symptom:* Traffic continues flowing through RP (suboptimal path)

*Diagnostic Commands:*

```
user@router> show pim join extensive  
user@router> show multicast route source-prefix 192.168.1.100
```

*Sample Problem Output:*

```

user@router> show pim join extensive
Group: 239.1.1.1
  Source: *
  Flags: sparse,rp-tree,wildcard    ## Still on RPT

  Source: 192.168.1.100
  Flags: none                      ## No (S,G) state!

```

**Cause:** SPT threshold set to infinity

**Solution:**

```

## Remove infinity threshold
[edit protocols pim]
delete spt-threshold infinity

## Or set specific threshold
[edit protocols pim]
set spt-threshold {
    rate 10;      ## Switch at 10 kbps
}

```

#### Scenario 4: Register Messages Dropped

**Symptom:** Sources can't register with RP

**Diagnostic Commands:**

```

user@router> show pim statistics | match register
user@router> show pim register detail

```

**Sample Problem Output:**

```

user@router> show pim statistics
PIM Message Type      Received      Sent
Register              0            1523
Register-Stop          1523          0      ## All stopped!

user@router> show log messages | match PIM
PIM: Register message rate limit exceeded

```

**Cause:** Register rate limiting or filtering

**Solution:**

```

## Adjust register limits
[edit protocols pim rp]
set register-limit {
    maximum 1000;          ## Increase limit
    rate 500;              ## kbps
}

## Check for register filtering
[edit protocols pim rp]
delete register-policy

```

#### Advanced Troubleshooting with Traceoptions:

```

## Enable detailed PIM debugging
[edit protocols pim]
set traceoptions {
    file pim-debug size 50m files 5;
    flag all detail;
}

## Monitor the trace
user@router> monitor start pim-debug
user@router> show log pim-debug | match "JOIN|REGISTER|ASSERT"

## Common issues to look for in traces:
## - "RPF check failed" - Unicast routing issue
## - "Assert lost" - Multiple routers forwarding

```

```
## - "Neighbor timeout" - Connectivity issue
## - "Register-Stop" - RP rejecting source
```

## Summary and Key Takeaways

You've now completed a comprehensive journey through IP multicast, IGMP, and PIM. Here are the essential concepts to master:

### Multicast Overview

- **Purpose:** Efficient one-to-many communication
- **Addressing:** 224.0.0.0/4 for IPv4 multicast
- **RPF Check:** Critical loop prevention mechanism
- **State Types:** (S,G) for source-specific, (\*,G) for any-source

### IGMP Mastery

- **Versions:** v2 most common, v3 adds source filtering
- **Timers:** Query interval, response time, membership timeout
- **Troubleshooting:** Version mismatches, multiple queriers, static vs dynamic

### PIM Excellence

- **Modes:** Sparse mode (explicit join) preferred over dense (flood-and-prune)
- **RP:** Central registration point in sparse mode
- **Trees:** RPT (\*,G) vs SPT (S,G) and switchover process
- **Messages:** Hello, Join/Prune, Register, Assert

### Configuration Best Practices

1. Always enable PIM on all participating interfaces including loopback
2. Use sparse mode unless specifically required otherwise
3. Configure redundant RPs using BSR or anycast
4. Implement proper IGMP timers for convergence
5. Use policy to control multicast group access
6. Enable BFD for faster failure detection

### Troubleshooting Methodology

1. Verify layer 2 connectivity first
2. Check IGMP state on last-hop routers
3. Verify PIM neighbors and RP reachability
4. Confirm RPF checks are passing
5. Trace the join path from receiver to source
6. Use traceoptions when other methods fail

Remember: Multicast is stateful and requires all routers in the path to participate. A single misconfigured router can break the entire flow. Always think about the complete path from source to receivers when troubleshooting.

## JNCIE-SP Master Course: Modules 22–23 & Lab 6

### BGP Multicast VPNs (MVPNs) – Complete Learning Module

---

## Module 22: BGP MVPN Overview

### Part 1: The Conceptual Lecture (The Why)

#### 1.1 The Fundamental Problem: Multicast Without MVPNs

Before we discuss MVPNs, let's understand what problem we're solving.

##### What is Multicast?

In a traditional network, when you send data, you use **unicast**: one sender, one receiver. The sender puts a destination address on the packet, and it travels to that one place.

But imagine you're a video streaming company and you want to broadcast a live sports event to **thousands** of customers simultaneously. If you used unicast, you'd send the same video stream independently to each customer—one stream to Customer A, another identical stream to Customer B, and so on. This consumes enormous amounts of bandwidth and server resources.

**Multicast solves this problem.** In multicast, one sender transmits a *single* stream, and multiple receivers can receive it simultaneously. The network makes copies of packets *only when necessary* to reach different destinations. This is far more efficient.

Multicast addresses are special. Instead of representing a single host, a multicast address represents a **group**. Any device can join this group to receive traffic destined to that group address. In IPv4, multicast addresses are in the range 224.0.0.0/4 (that is, 224.0.0.0 to 239.255.255.255).

#### The Problem in VPN Networks:

Now, imagine you're a Service Provider running Layer 3 VPNs (L3VPNs) for multiple customers. Each customer has their own isolated virtual network.

Within a single customer's VPN, you might want multicast to work—perhaps Customer A wants to broadcast a corporate video to all their branch offices. But here's the challenge:

1. **Isolation:** Customer A's multicast traffic must not leak into Customer B's VPN.
2. **Scalability:** Multicast groups might span across the provider's backbone. The backbone must carry the traffic efficiently.
3. **Control:** The Service Provider needs a way to manage which customers can use multicast and where that traffic flows.

**Without a proper solution**, multicast in VPN environments becomes chaotic. Multicast packets might:

- Leak between customers (security violation)
- Create broadcast storms
- Use inefficient forwarding paths

**BGP MVPNs are the solution.** They extend the L3VPN model to support multicast while maintaining customer isolation and efficient forwarding.

---

## 1.2 Core Concepts: What is a BGP MVPN?

A **BGP MVPN** (Multicast VPN) is a framework that allows Service Providers to offer multicast services to VPN customers while maintaining complete isolation between customers.

Think of it this way:

- **L3VPN (what we already know):** Uses BGP to distribute unicast routes and labels across the provider backbone. Each customer has their own routing instance (VRF) and routing tables.
- **MVPN (what we're learning now):** Uses BGP to distribute *multicast* routes and trees across the provider backbone, while keeping each customer's multicast traffic isolated in their own "virtual multicast domain."

#### Key Differences from Regular Multicast:

In regular multicast (without VPNs), a backbone router learns about multicast groups directly from customers and builds multicast trees across the backbone. But this doesn't work for VPNs because:

- Multiple customers might use the *same* multicast group address (e.g., 224.1.1.1)—you can't distinguish whose traffic is whose.
- There's no mechanism to enforce customer isolation.

BGP MVPNs solve this by:

1. **Encoding customer identity** into multicast information using **VPN-IPv4 multicast addresses** (similar to how L3VPNs use VPN-IPv4 unicast addresses).
  2. **Using BGP** to advertise multicast group memberships and sources across the backbone.
  3. **Building customer-specific multicast trees** inside the backbone, separate from unicast trees.
- 

## 1.3 MVPN Architecture: The Players

Let's identify the key components:

#### 1. Customer Edge (CE) Device:

- The customer's router at their site.
- Runs multicast routing protocols (e.g., PIM) within the customer's network.
- May send or receive multicast traffic.
- Has no special MVPN awareness—it just does multicast as normal.

#### 2. Provider Edge (PE) Router:

- The Service Provider's router connected to customers.
- Runs multicast in the VRF (virtual routing/forwarding instance) for each customer.
- Learns about customer multicast groups via PIM or other protocols.
- Advertises this information to other PEs in the backbone via BGP.
- Makes decisions about whether to forward multicast traffic into the backbone.


### 3. Provider (P) Router:

- Interior backbone routers.
- Participates in the backbone multicast tree infrastructure.
- May carry customer multicast traffic (but doesn't need to know which VPN it belongs to—that's encapsulated).

### 4. Rendezvous Point (RP):

- A router within a customer's network or across multiple sites that acts as a meeting point for multicast sources and receivers.
- PIM uses RPs to build shared multicast trees.

### Example Architecture:



Notice: Both customers use the *same* multicast address (224.1.1.1), but they're completely isolated. PE1 keeps them separate using the VRF mechanism, and the backbone carries the traffic from PE1 to PE2 on a backbone multicast tree.

## 1.4 MVPN Operation: Control Plane

### Step-by-Step Data Flow:

**Scenario:** Customer A has a multicast source at Site 1 and receivers at Site 2. Both sites connect to the provider network via PEs.

### Phase 1: Group Membership Discovery

#### Timeline:

T0: Receiver B at Site 2 wants to join group 224.1.1.1 |  
   – CE-B sends IGMP (Internet Group Management Protocol) |  
     join message locally. |  
   – CE-B's routing protocol learns about the group. |  
   – CE-B advertises the group to PE2 via PIM. |

T1: PE2 (Provider Edge at Site 2) learns:  
   "In VRF-B, there's a receiver for 224.1.1.1" |  
   PE2 does NOT immediately forward upstream. |  
   Instead, PE2 advertises this information to the |  
   other PEs via BGP MVPN routes. |

T2: PE1 receives BGP MVPN route from PE2 saying:  
   "VRF-B has a receiver for 224.1.1.1" |  
   PE1 stores this information and waits. |

### Phase 2: Source Registration

T3: Source A at Site 1 sends first packet to |

224.1.1.1. This packet reaches CE-A, which performs PIM Source Registration (or just participates in PIM Join).

T4: CE-A signals to PE1 (via PIM) about the source.  
PE1 learns: "VRF-A has source S sending to group 224.1.1.1"  
PE1 advertises this via BGP MVPN.

T5: PE2 receives BGP MVPN route from PE1 saying:  
"VRF-A has source S for group 224.1.1.1"  
PE2 recognizes: "My VRF-B has a receiver for 224.1.1.1, and there's a source for it in VRF-A. I should build a path to that source."  
PE2 initiates joining the source's tree.

### Phase 3: Tree Building

T6: The backbone PIM routers (including PE1 and P-routers) build a multicast tree from the source in VRF-A to the receiver in VRF-B. This tree is shared between the PEs and uses special techniques to keep traffic in a backbone multicast domain.

T7: Multicast traffic from Source A flows:  
- From Source A to CE-A (normal routing)  
- From CE-A to PE1 (normal routing)  
- From PE1, encapsulated, into the backbone multicast tree (PE1 → P-routers → PE2)  
- From PE2 to CE-B (normal routing)  
- From CE-B to Receiver B (normal routing)

#### BGP Routes Used (Control Plane Signaling):

BGP carries MVPN routes to communicate:

1. **Group membership routes** (e.g., "PE2 has receivers for group 224.1.1.1 in VRF-A")
2. **Source information routes** (e.g., "PE1 has source 10.1.1.1 sending to 224.1.1.1 in VRF-A")
3. **Tree information routes** (e.g., "Use PE1 as the multicast tree root")

These routes use special BGP extensions (MVPN AFI/SAFI - Address Family Identifier/Sub-Address Family Identifier).

## 1.5 MVPN Data Plane: How Traffic Flows

The **data plane** is the actual forwarding of multicast packets.

#### Encapsulation:

When a PE router sends multicast traffic into the backbone, it must ensure the traffic is only delivered to PEs that need it (those with receivers in that VPN). To do this, MVPNs use **encapsulation**:


1. The multicast packet arrives at PE1 from the customer.
2. PE1 adds a **Multiprotocol Label Switching (MPLS) label** to the packet (similar to L3VPN unicast).
3. The label tells backbone routers: "This packet belongs to the MVPN backbone tree and should follow backbone multicast routing."
4. The packet travels across the backbone on a **backbone multicast tree**.
5. When the packet reaches PE2, PE2 removes the label and forwards the original multicast packet to its local customer VRF.

#### Backbone Multicast Tree vs. Customer Multicast Tree:

Customer A's Network  
(VRF-A Multicast Tree)

Backbone  
(MVPN Backbone Tree)

Customer B's Network  
(VRF-B Multicast Tree)



## 1.6 Key Concepts Summary

Let me summarize the essential ideas:

Concept	Explanation
<b>Multicast Address</b>	Special address (224.0.0.0/4) representing a group. Any sender can transmit to it; any receiver can join it.
<b>MVPN</b>	Extends L3VPN to support multicast. Each VPN has its own multicast domain; different customers can use the same multicast addresses.
<b>VPN-IPv4 Multicast</b>	Like VPN-IPv4 unicast, but for multicast routes. Encodes customer identity (RD) + multicast group + source.
<b>BGP MVPN Routes</b>	BGP carries information about multicast group memberships, sources, and tree-building directives.
<b>PE Role</b>	Learns about customer multicast activity (via PIM), advertises it to other PEs via BGP, and builds backbone multicast paths.
<b>Backbone Multicast Tree</b>	The path multicast traffic takes across the provider backbone. Built using MPLS labels to ensure traffic stays within the MVPN infrastructure.
<b>Encapsulation</b>	Multicast packets are labeled when entering the backbone and delabeled when exiting. This ensures isolation and correct forwarding.

## Part 2: The Junos CLI Masterclass (The How)

### 2.1 MVPN Configuration Hierarchy in Junos

In Junos OS, MVPN configuration is hierarchical and builds on L3VPN concepts. Here's the structure:

```

[edit routing-instances <instance-name>]
  instance-type vrf
  interface <interface>
  routing-options {
    multicast {
      ...multicast-specific settings...
    }
  }
  protocols {
    pim {
      ...PIM configuration for the VRF...
    }
    bgp {
      ...BGP for the VRF (learned routes)...
    }
  }
}

```

Additionally, at the top-level (backbone) configuration:

```

[edit]
protocols {
  pim {
    ...backbone PIM...
  }
  bgp {
    group <peer-group> {
      family mvpn { # Special BGP address family for MVPN!
        ...
      }
    }
  }
}

```

```
}
```

**Key Insight:** MVPN configuration is split:

- **VRF-level:** Configure multicast within each customer's routing instance.
- **Backbone-level:** Configure how PEs advertise and receive multicast information to/from other PEs.

## 2.2 Pre-Configuration Assumptions

Before we configure MVPN, the following must already be in place:

1. **L3VPN Foundation:** Routing instances (VRFs) are created, customer interfaces are assigned, and unicast routing is working.
2. **Backbone Connectivity:** PE-to-PE connectivity exists (e.g., via BGP or static routes in the backbone).
3. **Backbone IGP:** An interior gateway protocol (e.g., OSPF, IS-IS) runs on the backbone to ensure PE-to-PE reachability.
4. **Backbone MPLS:** LDP (Label Distribution Protocol) or RSVP is running in the backbone to distribute labels.

If any of these are missing, MVPN will not work. This is because MVPN relies on:

- VRFs to separate customers.
- BGP to advertise MVPN routes.
- MPLS to encapsulate and label-switch multicast traffic.

## 2.3 Step-by-Step Configuration Walkthrough

Let's build a complete, working MVPN configuration. We'll configure a scenario with two sites:

**Scenario:**

- **PE1 connects to Site A** (Customer VPN-A). Site A has a multicast source.
- **PE2 connects to Site B** (Customer VPN-B). Site B has multicast receivers.
- **Backbone connects PE1 and PE2** with a multicast-capable core.

**Assumption:** L3VPN is already configured (VRF created, interfaces assigned, BGP unicast working).

### Step 1: Enable Multicast Routing in the Backbone

The backbone must run PIM to support multicast trees.

```
[edit protocols pim]
set rp local address 10.255.255.1    # This PE (PE1) is the backbone RP
set interface ge-0/0/0 mode sparse-dense
set interface ge-0/0/1 mode sparse-dense
set interface ge-0/0/2 mode sparse-dense
# Add all backbone interfaces
```

**Explanation:**

- **rp local address :** Designates this router as the Rendezvous Point for the backbone multicast domain. Receivers and sources use this RP for signaling.
- **mode sparse-dense :** Allows PIM to operate in both sparse mode (needs an RP) and dense mode (floods). Sparse-dense is flexible for mixed topologies.
- List all backbone interfaces (those carrying core traffic).

### Step 2: Configure MVPN Address Family in BGP

Enable BGP to carry MVPN routes (in addition to regular L3VPN routes).

```
[edit protocols bgp group internal]
set type internal
set local-address 10.255.255.1      # PE1's loopback
set family inet-vpn unicast          # Regular L3VPN (already done)
set family inet-vpn multicast         # NEW: MVPN multicast routes
```

```
# Add neighbors (other PEs)
set neighbor 10.255.255.2 description "PE2" # PE2's loopback
```

#### Explanation:

- `family inet-vpn multicast`: Enables BGP to advertise multicast information using the BGP MVPN AFI/SAFI.
- This tells BGP: "In addition to unicast VPN routes, also exchange multicast VPN routes with this neighbor."

---

### Step 3: Configure MVPN Within the Customer VRF

Now configure multicast behavior within the VRF for Customer A (on PE1).

```
[edit routing-instances VPN-A]
set instance-type vrf
set interface ge-1/0/0          # Customer-facing interface

# Configure multicast options
set routing-options multicast {
    rpf-check disable          # Trust multicast packets
}

# Enable PIM in this VRF
set protocols pim {
    rp static address 192.168.1.1      # Customer's internal RP
    interface ge-1/0/0 mode sparse-dense
}

# Ensure BGP is running (already configured for L3VPN)
# Routes in this VRF are learned from the CE via PIM
```

#### Explanation:

- `instance-type vrf`: This is a customer virtual routing instance.
- `rpf-check disable`: Relaxes strict reverse path forwarding in multicast, allowing traffic from CEs.
- `rp static address 192.168.1.1`: The customer's RP is at 192.168.1.1 (internal to the customer network).
- `interface ge-1/0/0 mode sparse-dense`: The customer interface participates in PIM. Multicast packets from the CE are accepted here.

---

### Step 4: Configure MVPN Signaling (BGP MVPN Routes)

Tell BGP how to advertise MVPN information for this VRF.

```
[edit routing-instances VPN-A]
set vrf-import [target:65000:100]      # Import unicast routes (L3VPN)
set vrf-export [target:65000:100]      # Export unicast routes
set vrf-route-import-target [
    target:65000:100 import-multicast # Import MVPN multicast
]
set vrf-route-export-target [
    target:65000:100 export-multicast # Export MVPN multicast
]
```

#### Explanation:

- `vrf-import/vrf-export target`: These control which routes are imported/exported for *unicast* (already known from L3VPN).
- `vrf-route-import-target [... import-multicast]`: Tells the PE: "Import MVPN multicast routes from other PEs that use this target."
- `vrf-route-export-target [... export-multicast]`: Tells the PE: "When you learn about multicast groups in this VRF, advertise them to other PEs using this target."

---

### Step 5: Repeat Configuration for PE2 and VPN-B

Apply similar configuration to PE2 for Customer B:

```
[edit protocols pim]
set rp local address 10.255.255.2    # PE2's loopback
set interface ge-0/0/0 mode sparse-dense
set interface ge-0/0/1 mode sparse-dense
```

```

[edit protocols bgp group internal]
# (Same as PE1, but with PE2's local address)
set local-address 10.255.255.2
set neighbor 10.255.255.1 description "PE1"

[edit routing-instances VPN-B]
set instance-type vrf
set interface ge-1/0/0          # Different customer interface

set routing-options multicast {
    rpf-check disable
}

set protocols pim {
    rp static address 192.168.2.1      # Customer B's RP
    interface ge-1/0/0 mode sparse-dense
}

set vrf-route-import-target [
    target:65000:101 import-multicast
]
set vrf-route-export-target [
    target:65000:101 export-multicast
]

```

## Step 6: Complete Configuration Reference

Here's a consolidated configuration for PE1:

```

# PE1 Configuration (Junos)

# =====#
# BACKBONE CONFIGURATION
# =====#

[edit protocols pim]
set rp local address 10.255.255.1
set interface ge-0/0/0.0 mode sparse-dense
set interface ge-0/0/1.0 mode sparse-dense
set interface lo0.0 mode sparse-dense

[edit protocols bgp group internal]
set type internal
set local-address 10.255.255.1
set family inet-vpn unicast
set family inet-vpn multicast      # Enable MVPN address family
set neighbor 10.255.255.2 description "PE2"

# =====#
# CUSTOMER VRF CONFIGURATION (VPN-A)
# =====#

[edit routing-instances VPN-A]
set instance-type vrf
set interface ge-1/0/0.0

# Multicast options
set routing-options multicast {
    rpf-check disable
}

# PIM in the VRF
set protocols pim {
    rp static address 192.168.1.1
    interface ge-1/0/0.0 mode sparse-dense
}

# MVPN import/export
set vrf-import [target:65000:100]
set vrf-export [target:65000:100]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]

```

```

# =====
# CUSTOMER VRF CONFIGURATION (VPN-B) - If PE1 also serves VPN-B
# =====

[edit routing-instances VPN-B]
set instance-type vrf
set interface ge-2/0/0.0

set routing-options multicast {
    rpf-check disable
}

set protocols pim {
    rp static address 192.168.2.1
    interface ge-2/0/0.0 mode sparse-dense
}

set vrf-import [target:65000:101]
set vrf-export [target:65000:101]
set vrf-route-import-target [target:65000:101 import-multicast]
set vrf-route-export-target [target:65000:101 export-multicast]

```

## 2.4 Configuration Logic and Patterns

### Pattern 1: MVPN Import/Export Targets

Just as L3VPN uses route targets to control unicast route distribution, MVPN uses route targets (with the `import-multicast` / `export-multicast` keywords) to control multicast information flow.

- If PE1 and PE2 share the same export target, they'll exchange MVPN routes and build backbone multicast trees for that customer.
- If they use different targets, they won't exchange MVPN routes (useful for separating customer VPNs).

### Pattern 2: MVPN Encapsulation

When a PE originates a multicast packet into the backbone, the Junos OS automatically:

1. Adds an MPLS label (similar to L3VPN unicast).
2. Marks the packet for backbone multicast forwarding.
3. Inserts a special route distinguisher (RD) and multicast group information.

This is all transparent to the configuration; Junos handles it automatically.

### Pattern 3: BGP MVPN Routes

BGP carries several types of MVPN routes:

Route Type	Meaning
Intra-AS I-PMSI A-D	Announcement: "This PE is part of the MVPN backbone for this VRF"
S-PMSI A-D	Announcement: "I have a source (S) for group (G); use this tree"
Leaf A-D	Announcement: "I (PE) have receivers for this group in my VRF"

The PE automatically generates and sends these routes based on:

- PIM Join messages from customers.
- Multicast group memberships learned from CEs.

## Part 3: Verification & Troubleshooting (The What-If)

### 3.1 Essential Verification Commands

When an MVPN is configured, use these commands to verify operation:

#### Command 1: Verify BGP MVPN Routes Are Advertised

```
show route advertising-protocol bgp 10.255.255.2 family inet-vpn multicast
```

**Sample Output:**

```

inet-vpn-multicast.0: 15 destinations, 20 routes (15 active)
Prefix                      Nexthop          MED   Lclprf   AS path
1:65000:100:192.168.1.0,255/32,0/32
*                           Self                  100
1:65000:100:192.168.1.1,255,192.168.1.2/32,0/32
*                           Self                  100
2:65000:100:192.168.1.1,255,10.1.1.0,255/32,0/32
*                           Self                  100

```

#### What to Look For:

- Routes should be marked with asterisk (\*) = active.
- Routes should have Lclprf (Local Preference) and be reachable.
- **Route Types:**
  - Type 1: Intra-AS I-PMSI (leaf advertisements)
  - Type 2: S-PMSI (source-to-group routes)
  - Type 3: Leaf A-D (customer has members for a group)

**Interpretation:** Routes are being advertised to PE2. This means PE1 has learned multicast information in its VRFs and is signaling it to the backbone.

---

#### Command 2: Verify BGP MVPN Routes Received

```
show route receiving-protocol bgp 10.255.255.2 family inet-vpn multicast
```

#### Sample Output:

```

inet-vpn-multicast.0: 15 destinations, 20 routes (15 active)
Prefix                      Nexthop          MED   Lclprf   AS path
1:65000:101:192.168.2.0,255/32,0/32
*                           10.255.255.2        100     I
1:65000:101:192.168.2.1,255,192.168.2.2/32,0/32
*                           10.255.255.2        100     I

```

#### What to Look For:

- Routes should be active (\*).
- Nexthop should be the remote PE (10.255.255.2).
- This confirms that PE1 has received MVPN routes from PE2.

**Interpretation:** PE1 knows about multicast groups in PE2's VRFs. This is necessary for the backbone to build multicast trees.

---

#### Command 3: Verify Backbone Multicast State

```
show multicast route family inet
```

#### Sample Output (Backbone Perspective):

```

Family: INET

Group: 224.0.0.0, Source: 10.255.255.1/32
  Upstream interface: lsi.0
  Session id: 0x64
  State: Active Joined
  Flags: None
  Upstream state: Join to 10.255.255.1
  Session Statistics:
    Packets: 1500, Bytes: 1500000, Packets/Sec: 10

Group: 224.0.0.0, Source: 10.1.1.1/32 (Customer Multicast)
  Upstream interface: ge-0/0/0.0
  Session id: 0x65
  State: Active Joined
  Flags: RPF
  Downstream interfaces: ge-0/0/1.0

```

#### **What to Look For:**

- State: Active Joined = The multicast tree is operational.
- Upstream interface = The direction towards the source.
- Downstream interfaces = The direction towards receivers.
- Packets counter should be non-zero = Traffic is flowing.

**Interpretation:** The backbone has built active multicast trees for customer sources and groups.

---

#### **Command 4: Verify PIM State in a VRF**

```
show pim join instance VPN-A family inet
```

#### **Sample Output:**

```
Family: INET

Group: 224.1.1.1
Source: 10.1.1.1
Upstream interface: ge-1/0/0.0
Upstream state: Join to 10.1.1.1
Join timer: 58
Prune timer: inactive
RPF neighbor: 192.168.1.2
Flags: sparse
Downstream interfaces:
    ge-1/0/0.0
```

#### **What to Look For:**

- Source and Group should be listed.
- Upstream state: Join = This PE has successfully joined the multicast tree.
- Downstream interfaces should list customer interfaces with receivers.

**Interpretation:** Within the VRF, PIM is operating correctly, tracking sources and group memberships.

---

#### **Command 5: Verify VRF Route Targets**

```
show route table VPN-A.inet-vpn multicast
```

#### **Sample Output:**

```
VPN-A.inet-vpn-multicast.0: 8 destinations, 10 routes (8 active)
A Destination      P Prf   Metric 1   Metric 2   Next hop          AS path
* 1:65000:100:192.168.1.1,255,10.1.1.1/32,0/32
    I 100                  10.255.255.2
* 2:65000:100:192.168.1.1,255,10.1.1.0,255/32,0/32
    I 100                  10.255.255.2
```

#### **What to Look For:**

- Rows marked with \* = routes are installed.
- Prf: 100 = BGP internal preference (correct).
- Nexthop shows remote PE (10.255.255.2) for routes learned from backbone.

**Interpretation:** MVPN routes are correctly installed in the VRF's multicast routing table.

---

## **3.2 Troubleshooting Scenarios**

---

### **Scenario 1: Multicast Packets Drop at PE Ingress**

#### **Symptom:**

- A customer multicast source at Site A sends traffic to 224.1.1.1.
- The traffic arrives at CE-A, but never exits to the backbone.
- Remote PE doesn't see the traffic.

#### Diagnostic Commands:

```
# Check if PIM is active in the VRF
show pim join instance VPN-A family inet

# Check if the customer interface is receiving multicast
show interfaces ge-1/0/0.0 statistics
(Look for input packets)

# Check MVPN route advertisements
show route advertising-protocol bgp 10.255.255.2 family inet-vpn multicast
(Look for routes matching the customer's source and group)

# Check RPF check status
show pim statistics instance VPN-A
(Look for "RPF failures" counter)
```

#### Sample Output Showing Problem:

```
show pim join instance VPN-A family inet

Family: INET
(Empty or no output - no multicast state!)

show pim statistics instance VPN-A
...
RPF Failures: 42
```

#### Root Causes:

1. **PIM not enabled on customer interface:** Verify set protocols pim interface ge-1/0/0.0 exists.
2. **rpf-check enabled:** If rpf-check is not disabled, strict RPF may reject customer packets. Disable it: set routing-options multicast rpf-check disable .
3. **No BGP MVPN routes advertised:** BGP must carry multicast routes. Verify set family inet-vpn multicast is set in BGP config.
4. **Customer RP misconfigured:** Verify the customer's RP address is reachable from the CE.

#### Solution:

```
# Ensure PIM is enabled on the customer interface
[edit routing-instances VPN-A]
set protocols pim interface ge-1/0/0.0 mode sparse-dense

# Disable strict RPF
[edit routing-instances VPN-A]
set routing-options multicast rpf-check disable

# Verify RP is reachable
ping 192.168.1.1 routing-instance VPN-A
# Should succeed

# Force PIM to rejoin
clear pim join instance VPN-A
```

#### Verification After Fix:

```
show pim join instance VPN-A family inet
(Should now show active (S,G) states)

show route advertising-protocol bgp 10.255.255.2 family inet-vpn multicast
(Should show MVPN routes marked with *)
```

#### Scenario 2: BGP MVPN Routes Not Exchanged Between PEs

##### Symptom:

- PE1 and PE2 are configured with MVPN.

- PE1 sends multicast from a customer, but PE2 doesn't learn about it.
- No MVPN routes appear in `show route advertising-protocol bgp`.

#### Diagnostic Commands:

```
# Check BGP session status
show bgp summary family inet-vpn multicast

# Check if BGP MVPN address family is active
show bgp group internal
(Look for "Family: inet-vpn multicast")

# Check MVPN route targets
show route table VPN-A.inet-vpn multicast advertising
(Look for routes with export-multicast target)

# Check BGP route rejection
show route advertising-protocol bgp 10.255.255.2 family inet-vpn multicast suppress
```

#### Sample Output Showing Problem:

```
show bgp summary family inet-vpn multicast
(No output or empty table - MVPN AFI not active!)

show bgp group internal
Type: Internal
Local Address: 10.255.255.1
...
Families: inet-vpn unicast           <-- Missing "inet-vpn multicast"!
```

#### Root Causes:

1. **MVPN address family not enabled in BGP:** family `inet-vpn multicast` missing from BGP config.
2. **Incorrect BGP MVPN targets:** VRF export targets don't match other PE's import targets.
3. **BGP session flapping:** BGP session to the remote PE is unstable.
4. **Label issues:** Backbone labels not distributed correctly by LDP/RSPV.

#### Solution:

```
# Enable MVPN address family in BGP
[edit protocols bgp group internal]
set family inet-vpn multicast

# Verify targets match
[edit routing-instances VPN-A]
show vrf-route-import-target
show vrf-route-export-target
# Should show "target:65000:100 import-multicast" and "export-multicast"

# Check BGP session
show bgp neighbor 10.255.255.2
(Should show "State: Established" and "Families: inet-vpn unicast, inet-vpn multicast")

# Soft-reset BGP to refresh MVPN routes
clear bgp neighbor 10.255.255.2 soft-inbound family inet-vpn multicast
```

#### Verification After Fix:

```
show bgp summary family inet-vpn multicast
(PE2 neighbor should show "1" in the "RcvdUpdates" column)

show route advertising-protocol bgp 10.255.255.2 family inet-vpn multicast
(Routes should appear)
```

### Scenario 3: Multicast Receiver Doesn't Receive Traffic

#### Symptom:

- PE1 learns about a multicast source at Site A.
- PE2 has a receiver for that group at Site B.

- PE1 advertises MVPN routes to PE2.
- But the receiver at Site B doesn't get traffic.

#### Diagnostic Commands:

```
# Check if PE2 has received MVPN routes
show route receiving-protocol bgp 10.255.255.1 family inet-vpn multicast

# Check PIM join state in PE2's VRF
show pim join instance VPN-B family inet

# Verify the multicast tree from source to receiver
show multicast rpf 224.1.1.1 10.1.1.1
(Shows the reverse path from receiver to source)

# Check backbone tree status
show multicast route family inet detail
(Look for source 10.1.1.1 and verify downstream interfaces)

# Check if receiver is actually requesting the group
show igmp group instance VPN-B
```

#### Sample Output Showing Problem:

```
show pim join instance VPN-B family inet
(No entry for group 224.1.1.1 - receiver didn't join!)

show igmp group instance VPN-B
(Empty - no IGMP membership reports)
```

#### Root Causes:

1. **Receiver not joined:** The CE at Site B isn't sending IGMP joins. Issue on customer network.
2. **PE2 not reporting receiver to backbone:** PIM on PE2 isn't triggering upstream joins. Verify PIM is configured.
3. **Source tree not built:** The backbone tree from PE1 to PE2 didn't establish. Check backbone PIM.
4. **Multicast RPF fails:** Backward path from receiver to source is asymmetric.

#### Solution:

```
# Ensure PIM is configured in VPN-B on PE2
[edit routing-instances VPN-B protocols pim]
set interface ge-1/0/0.0 mode sparse-dense

# Force receiver to rejoin
# (On the customer CE:)
# Restart IGMP or send explicit joins

# Verify backbone PIM state
show pim neighbors

# Soft-reset MVPN
clear pim join instance VPN-B
```

#### Verification After Fix:

```
show pim join instance VPN-B family inet
(Should show active (S,G) entry for 224.1.1.1)

show igmp group instance VPN-B
(Should show receiver joined to 224.1.1.1)

# Wait for traffic to flow
show multicast route family inet
(Packets counter should increment)
```

---

#### Scenario 4: MVPN Route Targets Misconfigured (Customer Isolation Broken)

##### Symptom:

- Customer A's multicast is leaking into Customer B's VRF.

- Both customers use the same group address (e.g., 224.1.1.1).
- Traffic appears to be duplicated or garbled.

#### Diagnostic Commands:

```
# Check route targets for both VRFs
[edit routing-instances VPN-A] show vrf-route-export-target
[edit routing-instances VPN-B] show vrf-route-export-target

# Check what routes each VRF imports
show route table VPN-A.inet-vpn multicast
show route table VPN-B.inet-vpn multicast

# Check if VPN-A's routes appear in VPN-B
show route table VPN-B.inet-vpn multicast | match "VPN-A"
```

#### Sample Output Showing Problem:

```
[edit routing-instances VPN-A] show vrf-route-export-target
target:65000:100 export-multicast

[edit routing-instances VPN-B] show vrf-route-export-target
target:65000:100 export-multicast           <-- SAME TARGET AS VPN-A!

# This means VPN-B is importing VPN-A's multicast routes!
show route table VPN-B.inet-vpn multicast | match "VPN-A"
1:65000:100:... (VPN-A's routes showing up in VPN-B!)
```

#### Root Cause:

Both VRFs use the same export/import targets, causing route leakage. Target 65000:100 is meant for VPN-A only.

#### Solution:

Assign unique targets to each VRF:

```
# VPN-A uses target 65000:100
[edit routing-instances VPN-A]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]

# VPN-B uses target 65000:101
[edit routing-instances VPN-B]
set vrf-route-import-target [target:65000:101 import-multicast]
set vrf-route-export-target [target:65000:101 export-multicast]
```

#### Verification After Fix:

```
show route table VPN-B.inet-vpn multicast
(Should NOT show any VPN-A routes; only VPN-B routes)

show route table VPN-A.inet-vpn multicast
(Should NOT show any VPN-B routes; only VPN-A routes)
```

### 3.3 Monitoring Command Reference Table

Command	Purpose	Key Output
show route family inet-vpn multicast	List all MVPN routes learned	Routes, nexthops, preferences
show pim join instance <VRF>	Multicast tree state within VRF	(S,G) entries, join/prune state
show multicast route family inet	Backbone multicast tree state	Upstream/downstream interfaces, packet counts
show igmp group instance <VRF>	IGMP membership in a VRF	Groups receivers joined, interface counts
show route table <VRF>.inet-vpn multicast	Routes in a specific VRF multicast table	VRF-specific MVPN routes
show bgp summary family inet-vpn multicast	BGP MVPN session summary	Session state, route counts
show interfaces <if>.0 statistics	Interface counters	Input/output packets, errors
show pim statistics instance <VRF>	PIM statistics within VRF	Joins sent, prunes, RPF failures

# Module 23: Configuring BGP MVPNs

## Part 1: The Conceptual Lecture (Deep Dive into MVPN Configuration)

### 1.1 MVPN Configuration Building Blocks

In Module 22, we learned the concepts and high-level configuration structure. Now, let's dive deeper into the logic and nuances of MVPN configuration.

#### Five Core Configuration Areas:

##### 1. Backbone Multicast Foundation

- Enable PIM in the backbone.
- Designate RPs for the backbone.
- Ensure MPLS labels are distributed.

##### 2. BGP MVPN Address Family

- Enable the `inet-vpn multicast` address family in BGP.
- Advertise and receive MVPN routes between PEs.

##### 3. Customer VRF Multicast Setup

- Enable PIM within each customer's VRF.
- Configure route targets for multicast import/export.
- Set multicast options (e.g., RPF).

##### 4. Multicast Tree Architecture Selection

- **I-PMSI (Inclusive PMSI):** Shared backbone tree for all groups in a VRF.
- **S-PMSI (Selective PMSI):** Dedicated trees per (source, group) pair.
- **None:** No backbone tree (for simple scenarios).

##### 5. Advanced Features

- Multicast loop prevention.
- Bandwidth management.
- MVPN failover and protection.


### 1.2 Inclusive PMSI vs. Selective PMSI

#### I-PMSI (Inclusive Provider-Managed Multicast Service Interface):

Think of I-PMSI as a **common highway** for all multicast traffic in a VRF.

- **Single tree per VRF:** When a PE joins multicast, it builds one backbone tree for that entire VRF.
- **All groups share one tree:** Every customer multicast group uses the same backbone path.
- **Simpler:** Easier to configure and manage.
- **Inefficient:** Traffic for many groups crosses the backbone on a single tree, even if only some groups have remote receivers.

Backbone Tree (I-PMSI) for VPN-A:




All groups (224.1.1.1, 224.1.1.2, 224.1.1.3, ...) flow across this one tree.

#### S-PMSI (Selective Provider-Managed Multicast Service Interface):

Think of S-PMSI as **dedicated lanes** for high-bandwidth flows.

- **Multiple trees per VRF:** For each (source, group) pair with significant traffic, a dedicated tree is built.
- **Optimized paths:** Trees are built only for groups that actually have remote receivers.
- **Higher bandwidth groups get their own tree:** Prevents congestion for popular groups.
- **More complex:** Requires signaling to determine which groups warrant dedicated trees.

Backbone Trees (S-PMSI) for VPN-A:



Group1 has dedicated tree: PE1 → PE2

Group2 has dedicated tree: PE1 → PE3

Other groups might still use I-PMSI

#### When to Use Which:

Scenario	Recommendation	Reason
Small VPN, few multicast groups	I-PMSI	Simpler, sufficient capacity
Large VPN, many low-bandwidth groups	I-PMSI	Overhead of S-PMSI not justified
VPN with a few high-bandwidth groups	S-PMSI	Dedicated trees reduce congestion
Service provider with tight SLA	S-PMSI	Precise control over bandwidth

### 1.3 MVPN Route Types in Detail

BGP carries several types of MVPN-specific routes. Understanding them is crucial for troubleshooting.

#### Type 1: Intra-AS I-PMSI A-D (Auto-Discovery)

**Purpose:** "I (this PE) am part of the MVPN for this VRF."

**Format:** RD:0:PE-address:0

**Example:**

```
1:65000:100:10.255.255.1:0
```

**Explanation:**

- 1 = Type 1 route.
- 65000:100 = Route Distinguisher (VPN-A).
- 10.255.255.1 = This PE's loopback address.
- :0 = Always zero for Type 1.

**When Used:**

- PE1 sends this to announce it has VRF-A (via iBGP to PE2).
- PE2 recognizes this and knows PE1 is part of the same VPN.
- On receiving Type 1, PEs can build I-PMSI backbone trees.

#### Type 2: S-PMSI A-D

**Purpose:** "I have a source (S) sending to group (G) in this VRF. Use this specific tree."

**Format:** RD:source:group:0:0

**Example:**

```
2:65000:100:10.1.1.1:224.1.1.1:0:0
```

**Explanation:**

- 2 = Type 2 route.
- 65000:100 = Route Distinguisher (VRF-A).
- 10.1.1.1 = Source IP.
- 224.1.1.1 = Multicast group.
- :0:0 = Additional fields (usually zero).

**When Used:**

- When S-PMSI is enabled and a source is detected, PE1 sends Type 2.
  - PE2 (with receivers for 224.1.1.1) receives this and joins the S-PMSI tree.
  - Traffic for this (S,G) pair is carried on the dedicated tree.
- 

### Type 3: Leaf A-D

**Purpose:** "I (this PE) have a receiver for group G (from source S) in this VRF."

**Format:** RD:source:group:PE-address:0

**Example:**

```
3:65000:100:10.1.1.1:224.1.1.1:10.255.255.2:0
```

**Explanation:**

- 3 = Type 3 route.
- 65000:100 = Route Distinguisher.
- 10.1.1.1 = Source.
- 224.1.1.1 = Group.
- 10.255.255.2 = This PE's address (indicates it has receivers).
- :0 = Zero field.

**When Used:**

- PE2 (which has receivers for the group) sends Type 3 to PE1.
  - PE1 recognizes: "PE2 has receivers for this group."
  - PE1 uses this to decide whether to forward traffic to the backbone.
- 

### Type 4: Source Active A-D

**Purpose:** "A source is active in this VRF, sending to this group. Please note for RPF checks."

**Format:** Similar to Type 2, used primarily for RPF validation.

---

### Type 5: Leaf A-D (Aggregated)

**Purpose:** Alternative format for Type 3, used in some deployments for optimization.

---

## 1.4 MVPN Configuration Patterns

### Pattern 1: Single VRF with I-PMSI Only

Simple scenario: One VPN, no S-PMSI needed.

Configuration focus:

- Enable PIM in backbone and VRF.
- Enable BGP MVPN address family.
- Set import/export multicast targets.
- Junos automatically builds I-PMSI trees.

### Pattern 2: Multi-VRF with S-PMSI for Large Groups

Complex scenario: Multiple VPNs, some groups are high-bandwidth.

Configuration focus:

- Identical baseline to Pattern 1.
- Additionally, define S-PMSI thresholds.
- Groups exceeding bandwidth thresholds get dedicated trees.

### Pattern 3: MVPN with Internet Access

Scenario: Customer multicast sources are external (over Internet).

#### Configuration focus:

- Separate I-PMSI tree for internal multicast.
- Import external multicast via BGP.
- Use policy to control which external groups are allowed.

## Part 2: The Junos CLI Masterclass (The How) – Advanced Configuration

In Module 22, we covered the basics. Now, let's explore advanced configurations and options.

### 2.1 MVPN with I-PMSI (Default Behavior)

#### Scenario:

- PE1 and PE2 serve customers with multicast.
- Start with I-PMSI (simplest approach).
- All multicast in a VRF shares one backbone tree.

#### Configuration (PE1):

```
# =====
# BACKBONE CONFIGURATION
# =====

[edit protocols pim]
set rp local address 10.255.255.1    # PE1 is backbone RP
set interface ge-0/0/0.0 mode sparse-dense
set interface ge-0/0/1.0 mode sparse-dense
set interface lo0.0 mode sparse-dense

[edit protocols bgp group internal]
set type internal
set local-address 10.255.255.1
set family inet-vpn unicast
set family inet-vpn multicast      # Enable MVPN

set neighbor 10.255.255.2 description "PE2"

# =====
# VRF CONFIGURATION (I-PMSI MODE)
# =====

[edit routing-instances VPN-A]
set instance-type vrf
set interface ge-1/0/0.0          # Customer interface

# Multicast options
set routing-options multicast {
    rpf-check disable
    interface-seed {
        ge-1/0/0.0           # Customer interface is source of multicast
    }
}

# PIM in VRF
set protocols pim {
    rp static address 192.168.1.1      # Customer's RP
    interface ge-1/0/0.0 mode sparse-dense
}

# I-PMSI configuration
set protocols pim mvpn {
    i-pmsi                      # Enable I-PMSI for all groups
    rp static address 192.168.1.1
}

# BGP MVPN: Import and export multicast routes
set vrf-import [target:65000:100]
set vrf-export [target:65000:100]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]
```

#### Explanation of Key Commands:

- set protocols pim mvpn i-pmsi : Enables I-PMSI mode. Junos will build a single backbone tree for all multicast in this VRF.
  - set protocols pim mvpn rp static address : Designates an RP for the backbone multicast tree.
  - interface-seed : Identifies which interfaces are sources of multicast (customer interfaces).
- 

## 2.2 MVPN with S-PMSI (Advanced)

### Scenario:

- Same as above, but now add S-PMSI for bandwidth-intensive groups.
- Groups exceeding 10 Mbps get dedicated trees.

### Configuration (PE1):

```
[edit routing-instances VPN-A]

# Multicast options (same as I-PMSI)
set routing-options multicast {
    rpf-check disable
    interface-seed {
        ge-1/0/0.0
    }
}

# PIM in VRF (same as I-PMSI)
set protocols pim {
    rp static address 192.168.1.1
    interface ge-1/0/0.0 mode sparse-dense
}

# S-PMSI configuration (new)
set protocols pim mvpn {
    i-pmsi                      # I-PMSI as default
    s-pmsi {
        group-threshold 10000      # Groups exceeding 10 Mbps (10000 Kbps)
                                    # will get S-PMSI trees
        maximum-p2mp-spmsi 100     # Max 100 S-PMSI trees per VRF
    }
    rp static address 192.168.1.1
}

# BGP MVPN (same as I-PMSI)
set vrf-import [target:65000:100]
set vrf-export [target:65000:100]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]
```

### Explanation of S-PMSI Options:

Option	Purpose	Example
group-threshold	Bandwidth threshold for triggering S-PMSI	10000 = 10 Mbps
maximum-p2mp-spmsi	Max S-PMSI trees per VRF	100 prevents tree explosion
leaf-ad-join	Enable Leaf A-D routes for optimization	Boolean

### Behavior:

- Groups with traffic rate below 10 Mbps use I-PMSI.
  - When a group exceeds 10 Mbps, Junos automatically creates an S-PMSI tree.
  - The source PE (PE1) sends Type 2 (S-PMSI) routes to advertise the new tree.
  - Receiver PEs (like PE2) receive these routes and join the S-PMSI tree.
  - Traffic switches to the dedicated tree, offloading from I-PMSI.
- 

## 2.3 MVPN with Multicast Loop Prevention

### Scenario:

- Customers have redundant connections to the backbone (for resilience).

- This can create multicast loops if not handled carefully.
- Use RPF (Reverse Path Forwarding) and multicast policies to prevent loops.

#### Configuration (PE1):

```
[edit routing-instances VPN-A]

set routing-options multicast {
    rpf-check disable                      # Already disabled; customers trust
    interface-seed {
        ge-1/0/0.0
        ge-2/0/0.0                         # Redundant customer interface
    }
}

set protocols pim {
    rp static address 192.168.1.1
    interface ge-1/0/0.0 mode sparse-dense
    interface ge-2/0/0.0 mode sparse-dense # Both interfaces in PIM
}

# Multicast Policy: Prevent loops by restricting which interfaces
# multicast traffic can egress
set policy-options policy-statement multicast-egress {
    from interface ge-1/0/0.0
    then accept
}

set protocols pim mvpn {
    i-pmsi
    rp static address 192.168.1.1
    multicast-policy multicast-egress # Apply policy
}
```

#### How This Works:

- Multicast packets from redundant interface ge-2/0/0.0 are accepted (ingress).
- But egress policies ensure that multicast sent out through the backbone doesn't loop back through both customer interfaces.
- PIM's Join/Prune mechanisms further prevent duplicate traffic.

## 2.4 MVPN Hub-and-Spoke Configuration

#### Scenario:

- Customer has a central hub site (HQ) and multiple spoke sites (branches).
- Multicast traffic should flow through the hub (e.g., corporate video).
- Spokes should not communicate multicast directly with each other.

#### Configuration (PE1 serving the hub):

```
[edit routing-instances VPN-A-HUB]
set instance-type vrf
set interface ge-1/0/0.0          # Hub site interface

set routing-options multicast {
    rpf-check disable
    interface-seed {
        ge-1/0/0.0
    }
}

set protocols pim {
    rp static address 192.168.1.1
    interface ge-1/0/0.0 mode sparse-dense
}

set protocols pim mvpn {
    i-pmsi
    rp static address 192.168.1.1
    root-monitor                 # PE1 monitors hub for issues
}
```

```

set vrf-import [target:65000:100]
set vrf-export [target:65000:100]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]

```

#### Configuration (PE2 serving a spoke):

```

[edit routing-instances VPN-A-SPOKE]
set instance-type vrf
set interface ge-1/0/0.0          # Spoke site interface

set routing-options multicast {
    rpf-check disable
    interface-seed {
        ge-1/0/0.0
    }
}

set protocols pim {
    rp static address 192.168.1.1      # Same RP as hub
    interface ge-1/0/0.0 mode sparse-dense
}

set protocols pim mvpn {
    i-pmsi
    rp static address 192.168.1.1
    root-preference 100                # Lower preference to prefer hub tree
}

set vrf-import [target:65000:100]
set vrf-export [target:65000:100]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]

```

#### Behavior:

- Hub PE1 becomes the "root" of the backbone multicast tree for this VPN.
- All spoke PEs (like PE2) preferentially route multicast through PE1.
- Spoke-to-spoke multicast is avoided (traffic paths through hub).

## 2.5 Complete Configuration Reference

#### PE1 (Full MVPN Configuration with I-PMSI + S-PMSI + Hub-and-Spoke):

```

# =====
# BACKBONE
# =====

[edit protocols pim]
set rp local address 10.255.255.1
set interface ge-0/0/0.0 mode sparse-dense
set interface ge-0/0/1.0 mode sparse-dense
set interface lo0.0 mode sparse-dense

[edit protocols bgp group internal]
set type internal
set local-address 10.255.255.1
set family inet-vpn unicast
set family inet-vpn multicast
set neighbor 10.255.255.2 description "PE2"
set neighbor 10.255.255.3 description "PE3"

# =====
# VRF: VPN-A (Standard I-PMSI + S-PMSI)
# =====

[edit routing-instances VPN-A]
set instance-type vrf
set interface ge-1/0/0.0

set routing-options {
    multicast {

```

```

        rpf-check disable
        interface-seed ge-1/0/0.0
    }

}

set protocols pim {
    rp static address 192.168.1.1
    interface ge-1/0/0.0 mode sparse-dense
}

set protocols pim mvpn {
    i-pmsi
    s-pmsi {
        group-threshold 10000
        maximum-p2mp-spmsi 100
    }
    rp static address 192.168.1.1
}

set vrf-import [target:65000:100]
set vrf-export [target:65000:100]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]

# =====
# VRF: VPN-B (Hub-and-Spoke with I-PMSI)
# =====

[edit routing-instances VPN-B-HUB]
set instance-type vrf
set interface ge-2/0/0.0

set routing-options {
    multicast {
        rpf-check disable
        interface-seed ge-2/0/0.0
    }
}

set protocols pim {
    rp static address 192.168.2.1
    interface ge-2/0/0.0 mode sparse-dense
}

set protocols pim mvpn {
    i-pmsi
    rp static address 192.168.2.1
    root-monitor           # PE1 is the hub root
}

set vrf-import [target:65000:101]
set vrf-export [target:65000:101]
set vrf-route-import-target [target:65000:101 import-multicast]
set vrf-route-export-target [target:65000:101 export-multicast]

```

## Part 3: Verification & Troubleshooting (The What-If) – Advanced

### 3.1 Verifying I-PMSI Operation

**Command:** Check I-PMSI Tree Status

```
show route table VPN-A.inet-vpn multicast advertising
```

**Sample Output:**

```

VPN-A.inet-vpn-multicast.0: 10 destinations, 12 routes (10 active)
A Destination                  P Prf Metric Next hop
* 1:65000:100:10.255.255.1:0   B 100    10.255.255.1
* 1:65000:100:10.255.255.2:0   B 100    10.255.255.2
* 1:65000:100:10.255.255.3:0   B 100    10.255.255.3

```

**Interpretation:**

- Type 1 routes (marked with 1:...) represent I-PMSI Auto-Discovery advertisements.
  - Three PEs (10.255.255.1, 2, 3) are advertising their participation in this MVPN.
  - These routes are the foundation for building I-PMSI backbone trees.
- 

#### Command: Verify I-PMSI Backbone Tree

```
show multicast route family inet instance <default>
```

#### Sample Output (Backbone Perspective):

```
Family: INET

Group: 224.0.0.0, Source: 10.255.255.1/32
Upstream interface: lsi.0
Session id: 0x100
State: Active Joined
Flags: None
Upstream state: Join to 10.255.255.1
Session Statistics:
  Packets: 50000, Bytes: 50000000, Packets/Sec: 100

Downstream interfaces:
  ge-0/0/1.0          # PE2
  ge-0/0/2.0          # PE3
```

#### Interpretation:

- The I-PMSI tree is built: Root is 10.255.255.1 (PE1), and all VPN-A multicast flows through this tree.
  - Downstream interfaces ge-0/0/1.0 and ge-0/0/2.0 are PE2 and PE3.
  - Packet counters are non-zero, indicating traffic is actively flowing.
- 

## 3.2 Verifying S-PMSI Operation

#### Command: Check S-PMSI Routes

```
show route table VPN-A.inet-vpn multicast | match "^\d:"
```

#### Sample Output:

```
2:65000:100:10.1.1.1:224.1.1.1:0:0
*
B 100      10.255.255.2
2:65000:100:10.1.1.2:224.1.1.2:0:0
*
B 100      10.255.255.2
```

#### Interpretation:

- Type 2 routes represent S-PMSI advertisements for specific (source, group) pairs.
  - Two groups (224.1.1.1 and 224.1.1.2) have exceeded the bandwidth threshold (10 Mbps) and have dedicated S-PMSI trees.
  - These trees originate from PE2 (nexthop), meaning PE2 is the tree root for these groups.
- 

#### Command: Verify S-PMSI Tree Details

```
show pim join instance VPN-A family inet | match "Source: 10.1.1.1"
```

#### Sample Output:

```
Group: 224.1.1.1
Source: 10.1.1.1
Upstream interface: lsi.0          # S-PMSI tunnel interface
Upstream state: Join to 10.1.1.1
Join timer: 59
Flags: sparse
```

```
Downstream interfaces:  
  ge-1/0/0.0                                # Customer interface
```

#### Interpretation:

- The source 10.1.1.1 sending to 224.1.1.1 has joined an S-PMSI tree (note `lsi.0` = Label Switched Interface).
- This dedicated tree reduces latency and provides QoS guarantees compared to shared I-PMSI.

## 3.3 Advanced Troubleshooting Scenarios

### Scenario 5: S-PMSI Tree Not Building Despite High Bandwidth

#### Symptom:

- A multicast group (224.1.1.1) consistently sends >10 Mbps traffic.
- Configuration specifies `group-threshold 10000` (10 Mbps).
- But S-PMSI route is not created; traffic still uses I-PMSI.

#### Diagnostic Commands:

```
# Check S-PMSI configuration  
[edit routing-instances VPN-A] show protocols pim mvpn  
  
# Monitor bandwidth of specific groups  
show pim join instance VPN-A family inet detail  
  
# Check BGP Type 2 routes  
show route table VPN-A.inet-vpn multicast | match "^2:"  
  
# Check for policy rejections  
show policy policy-statement <policy> | match "reject"  
  
# Verify multicast statistics  
show pim statistics instance VPN-A  
  
# Check if S-PMSI is actually enabled  
show route table inet.2                               # Internal multicast table
```

#### Sample Output Showing Problem:

```
[edit routing-instances VPN-A] show protocols pim mvpn  
(No output or S-PMSI section missing!)  
  
show route table VPN-A.inet-vpn multicast | match "^2:"  
(No Type 2 routes, even with high-bandwidth groups)
```

#### Root Causes:

1. **S-PMSI not enabled:** `s-pmsi { ... }` section missing from configuration.
2. **Threshold too high:** Group bandwidth not actually exceeding the threshold.
3. **Policy rejecting S-PMSI:** An import/export policy silently drops Type 2 routes.
4. **Backbone connectivity issue:** No path to build S-PMSI tree.

#### Solution:

```
# Verify S-PMSI is configured  
[edit routing-instances VPN-A protocols pim mvpn]  
set s-pmsi {  
    group-threshold 10000                      # Ensure this matches actual bandwidth  
    maximum-p2mp-spmsi 100  
}  
  
# Verify the group is actually exceeding threshold  
show pim join instance VPN-A family inet group 224.1.1.1 detail  
# Check the "Packets/Sec" to confirm bandwidth  
  
# Check policies  
show route table VPN-A.inet-vpn multicast suppress  
# If output shows suppressed Type 2 routes, a policy is blocking them
```

```

# If policy is the issue, verify import/export targets
show route table VPN-A.inet-vpn multicast advertising | match "^2:"
# Type 2 should be advertised by this PE

# Force BGP to refresh
clear bgp neighbor 10.255.255.2 soft-inbound
clear pim statistics instance VPN-A

```

#### Verification After Fix:

```

show route table VPN-A.inet-vpn multicast | match "^2:"
# Type 2 routes should now appear

show pim join instance VPN-A family inet detail | match "lsi"
# S-PMSI trees (lsi.0 interface) should be active

```

### Scenario 6: Multicast Reaching Wrong VRF (Route Target Misconfiguration)

#### Symptom:

- Multicast from Customer A (VPN-A) is appearing in Customer B's VRF (VPN-B).
- Both VRFs use the same multicast group 224.1.1.1.
- Traffic should be completely isolated.

#### Diagnostic Commands:

```

# Check import/export targets for both VRFs
show route table VPN-A.inet-vpn multicast
show route table VPN-B.inet-vpn multicast

# Check if VPN-A's routes appear in VPN-B
show route table VPN-B.inet-vpn multicast | match "65000:100"

# Verify BGP MVPN routes
show route family inet-vpn multicast | match "VPN-A"
show route family inet-vpn multicast | match "VPN-B"

# Check VRF configuration
[edit routing-instances VPN-A] show vrf-route-export-target
[edit routing-instances VPN-B] show vrf-route-import-target

```

#### Sample Output Showing Problem:

```

[edit routing-instances VPN-A] show vrf-route-export-target
target:65000:100 export-multicast

[edit routing-instances VPN-B] show vrf-route-import-target
target:65000:100 import-multicast      <-- Importing VPN-A's target!
target:65000:101 import-multicast

show route table VPN-B.inet-vpn multicast
1:65000:100:10.255.255.1:0          <-- VPN-A's Type 1 route in VPN-B!
2:65000:100:10.1.1.1:224.1.1.1:0:0    <-- VPN-A's S-PMSI route in VPN-B!

```

#### Root Cause:

VPN-B's import targets include 65000:100 , which matches VPN-A's export target. This causes VPN-A's MVPN routes to be imported into VPN-B's multicast table, creating cross-VPN leakage.

#### Solution:

Ensure each VRF has unique import/export targets:

```

[edit routing-instances VPN-A]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]

[edit routing-instances VPN-B]
set vrf-route-import-target [target:65000:101 import-multicast]
set vrf-route-export-target [target:65000:101 export-multicast]

```

```
# For shared VPNs (e.g., partner sites), use the same target
[edit routing-instances Partner-VPN]
set vrf-route-import-target [target:65000:100 import-multicast]
set vrf-route-export-target [target:65000:100 export-multicast]
```

#### Verification After Fix:

```
show route table VPN-B.inet-vpn multicast | match "65000:100"
# Should return no results (VPN-A routes excluded)

show route table VPN-A.inet-vpn multicast | match "65000:100"
# Should show VPN-A routes only
```

### Scenario 7: Backbone PIM Neighbors Not Joining Multicast Trees

#### Symptom:

- PE1 has active multicast in VPN-A.
- PE1 sends Type 1 (I-PMSI) routes to PE2.
- But backbone PIM shows no tree state.
- PE2 has no I-PMSI tree listed in show multicast route .

#### Diagnostic Commands:

```
# Check backbone PIM neighbors
show pim neighbors

# Check multicast routing table (backbone)
show multicast route family inet

# Check if Type 1 routes are received by PE2
show route receiving-protocol bgp 10.255.255.1 family inet-vpn multicast

# Check backbone PIM configuration
show protocols pim

# Verify MPLS labels in backbone
show route table mpls.0

# Check for PIM errors
show pim statistics
```

#### Sample Output Showing Problem:

```
show pim neighbors
(Empty – no PIM neighbors; PE1 is isolated!)

show multicast route family inet
(No routes at all)

show route receiving-protocol bgp 10.255.255.1 family inet-vpn multicast
(No Type 1 routes received)

show pim statistics
Joins Sent: 0
Joins Failed: 0
```

#### Root Causes:

- PIM not configured on backbone interfaces:** PIM must run between PEs.
- No backbone RP:** Type 1 routes use the backbone RP for tree building; if the RP is misconfigured, trees don't build.
- BGP MVPN address family not enabled:** Junos won't send Type 1 routes without family inet-vpn multicast in BGP.
- Interface not added to PIM:** Backbone interfaces must be explicitly added to PIM configuration.

#### Solution:

```
# Verify backbone PIM is running
[edit protocols pim]
set rp local address 10.255.255.1      # This PE is the RP
```

```

set interface ge-0/0/0.0 mode sparse-dense
set interface ge-0/0/1.0 mode sparse-dense
set interface lo0.0 mode sparse-dense
# Add all backbone interfaces!

# Verify BGP MVPN is enabled
[edit protocols bgp group internal]
set family inet-vpn multicast

# Verify BGP MVPN routes are sent
show route advertising-protocol bgp 10.255.255.2 family inet-vpn multicast
# Should list Type 1 routes

```

#### Verification After Fix:

```

show pim neighbors
# Neighbors should be listed (e.g., 10.255.255.2)

show multicast route family inet
# I-PMSI backbone trees should appear

show route receiving-protocol bgp 10.255.255.2 family inet-vpn multicast
# Type 1 routes should be received

```

### Scenario 8: RPF Check Failures Blocking Multicast

#### Symptom:

- Multicast traffic from the customer arrives at PE1.
- PE1 receives it but immediately drops the packets.
- No traffic reaches the backbone.
- Counters show high RPF failures.

#### Diagnostic Commands:

```

# Check RPF failures
show pim statistics instance VPN-A

# Check routing path
show route routing-table instance VPN-A table inet.0 10.1.1.1/32

# Verify RPF configuration
[edit routing-instances VPN-A] show routing-options multicast

# Check interface RPF
show pim interfaces detail instance VPN-A

# Examine packet drops
show interfaces ge-1/0/0.0 statistics

```

#### Sample Output Showing Problem:

```

show pim statistics instance VPN-A
RPF Failures: 1000      <-- Very high!

[edit routing-instances VPN-A] show routing-options multicast
rpf-check          (enabled, causing the problem!)

show route routing-table instance VPN-A table inet.0 10.1.1.1/32
(Asymmetric routing; source reachable via different interface than reverse path)

```

#### Root Causes:

1. **RPF check enabled:** rpf-check should be disabled for multicast from customers.
2. **Asymmetric routing:** Forward path to source via interface A, but reverse path via interface B. RPF rejects this.
3. **Route update lag:** Routing tables not yet converged.

#### Solution:

```

# Disable RPF check (trust customers)
[edit routing-instances VPN-A routing-options multicast]
set rpf-check disable

# If asymmetric routing is necessary, use a more lenient policy
[edit routing-instances VPN-A routing-options multicast]
set rpf-check loose           # Allow asymmetric paths

```

#### Verification After Fix:

```

show pim statistics instance VPN-A | match "RPF Failures"
# Should show zero or very low failures

show multicast route family inet
# Traffic should now flow through backbone

```

### 3.4 Comprehensive Verification Checklist

Use this checklist when deploying or troubleshooting MVPNs:

#### MVPN Verification Checklist:


- [ ] Backbone Connectivity
  - [ ] Ping all PE loopback addresses
  - [ ] Verify IGP (OSPF/IS-IS) routing
  - [ ] Confirm MPLS labels distributed via LDP/RSPV
- [ ] BGP Configuration
  - [ ] BGP sessions established (show bgp summary)
  - [ ] inet-vpn unicast routes exchanged
  - [ ] inet-vpn multicast address family enabled
  - [ ] Type 1 routes advertised and received
- [ ] Backbone PIM
  - [ ] PIM neighbors active (show pim neighbors)
  - [ ] Backbone RP configured
  - [ ] PIM interfaces in mode sparse-dense or sparse
  - [ ] Multicast routes in backbone (show multicast route family inet)
- [ ] VRF Configuration
  - [ ] VRF created and interfaces assigned
  - [ ] PIM enabled in VRF
  - [ ] RP configured for VRF
  - [ ] RPF check disabled (unless asymmetric routing required)
  - [ ] Multicast import/export targets configured
- [ ] Customer Multicast
  - [ ] CE multicast interfaces active (ping within customer)
  - [ ] PIM joins from CE received at PE (show pim join)
  - [ ] Multicast routes in VRF table (show route table VRF.inet-vpn multicast)
  - [ ] Type 1/2/3 routes generated by PE
- [ ] End-to-End Flow
  - [ ] Multicast source at Site A sends traffic
  - [ ] PE1 receives and forwards to backbone
  - [ ] Backbone PIM tree carries traffic
  - [ ] PE2 receives traffic from backbone
  - [ ] PE2 forwards to receiver at Site B
  - [ ] Receiver successfully receives multicast

## Lab 6: MVPNs – Hands-On Exercises

### Exercise 1: Basic MVPN Deployment (I-PMSI)

**Objective:** Deploy a basic MVPN with I-PMSI across two PEs serving two customer VPNs.

**Topology:**



### Configuration Steps:

#### Step 1: Configure PE1

```

# Configure backbone PIM and BGP MVPN

delete interfaces
delete routing-instances
delete protocols pim
delete protocols bgp

set interfaces lo0 unit 0 family inet address 10.255.255.1/32

set interfaces ge-0/0/0 unit 0 family inet address 10.3.3.1/30
set interfaces ge-1/0/0 unit 0 family inet address 10.1.1.1/30
set interfaces ge-2/0/0 unit 0 family inet address 10.4.4.1/30

set protocols ospf area 0.0.0.0 interface lo0.0 passive
set protocols ospf area 0.0.0.0 interface ge-0/0/0.0
set protocols ospf area 0.0.0.0 interface ge-3/0/0.0

set protocols ldp interface lo0.0
set protocols ldp interface ge-0/0/0.0
set protocols ldp interface ge-3/0/0.0

set protocols pim rp local address 10.255.255.1
set protocols pim interface ge-0/0/0.0 mode sparse-dense
set protocols pim interface lo0.0 mode sparse-dense

set protocols bgp group internal type internal
set protocols bgp group internal local-address 10.255.255.1
set protocols bgp group internal family inet-vpn unicast
set protocols bgp group internal family inet-vpn multicast
set protocols bgp group internal neighbor 10.255.255.2

# Configure Customer VRF-A

set routing-instances VPN-A instance-type vrf
set routing-instances VPN-A interface ge-1/0/0.0
set routing-instances VPN-A routing-options multicast rpf-check disable
set routing-instances VPN-A routing-options multicast interface-seed ge-1/0/0.0

set routing-instances VPN-A protocols pim rp static address 192.168.1.1
set routing-instances VPN-A protocols pim interface ge-1/0/0.0 mode sparse-dense

set routing-instances VPN-A protocols pim mvpn i-pmsi
set routing-instances VPN-A protocols pim mvpn rp static address 192.168.1.1

set routing-instances VPN-A vrf-import target:65000:100
set routing-instances VPN-A vrf-export target:65000:100
set routing-instances VPN-A vrf-route-import-target target:65000:100 import-multicast
set routing-instances VPN-A vrf-route-export-target target:65000:100 export-multicast

# Configure unicast routing to reach other PE (for MVPN)
set routing-options static route 0.0.0.0/0 next-hop 10.3.3.2
set routing-instances VPN-A routing-options static route 10.4.4.0/30 next-hop 10.1.1.2

```

#### Step 2: Configure PE2

```

# Similar to PE1, but with:

set interfaces lo0 unit 0 family inet address 10.255.255.2/32

```

```

set interfaces ge-0/0/0 unit 0 family inet address 10.3.3.2/30
set interfaces ge-1/0/0 unit 0 family inet address 10.2.2.1/30

set protocols bgp group internal local-address 10.255.255.2
set protocols bgp group internal neighbor 10.255.255.1

# Configure Customer VRF-B (same RT for this exercise, different customer)

set routing-instances VPN-B instance-type vrf
set routing-instances VPN-B interface ge-1/0/0.0
set routing-instances VPN-B routing-options multicast rpf-check disable

set routing-instances VPN-B protocols pim rp static address 192.168.2.1
set routing-instances VPN-B protocols pim interface ge-1/0/0.0 mode sparse-dense

set routing-instances VPN-B protocols pim mvpn i-pmsi
set routing-instances VPN-B protocols pim mvpn rp static address 192.168.2.1

set routing-instances VPN-B vrf-import target:65000:101
set routing-instances VPN-B vrf-export target:65000:101
set routing-instances VPN-B vrf-route-import-target target:65000:101 import-multicast
set routing-instances VPN-B vrf-route-export-target target:65000:101 export-multicast

```

### Step 3: Verification

```

# On PE1, verify BGP MVPN routes:
show route family inet-vpn multicast

# Expected output:
# Type 1 routes for both VPN-A and VPN-B

# Verify PIM neighbors:
show pim neighbors

# Verify multicast trees in VRFs:
show pim join instance VPN-A family inet
show pim join instance VPN-B family inet

# End-to-end multicast test:
# (On CE-A1) Start multicast source:
# mcast_send -a 224.1.1.1 -p 5555

# (On CE-B2) Join multicast group:
# mcast_receive -a 224.1.1.1 -p 5555

# Packets should flow from CE-A1 → PE1 → PE2 → CE-B2

```

## Exercise 2: S-PMSI Configuration and Bandwidth Monitoring

**Objective:** Enable S-PMSI and verify that high-bandwidth groups get dedicated trees.

### Configuration:

```

# On PE1 and PE2, modify MVPN configuration to include S-PMSI:

[edit routing-instances VPN-A protocols pim mvpn]
set i-pmsi
set s-pmsi group-threshold 5000          # 5 Mbps threshold
set s-pmsi maximum-p2mp-spmsi 50         # Max 50 S-PMSI trees
set rp static address 192.168.1.1

```

### Monitoring:

```

# Check Type 2 (S-PMSI) routes:
show route table VPN-A.inet-vpn multicast | match "^2:"

# Generate high-bandwidth multicast traffic:
# (On CE-A1) mcast_send -a 224.1.1.1 -p 5555 -b 10M (10 Mbps)

# Wait 30 seconds, then check for S-PMSI tree:
show route table VPN-A.inet-vpn multicast | match "^2:"

```

```
# Verify S-PMSI tree in PIM state:  
show pim join instance VPN-A family inet detailed | match "lsi"  
# lsi.0 interface indicates S-PMSI tunnel
```

## Exercise 3: Troubleshooting Multicast Isolation

**Objective:** Intentionally break customer isolation and fix it.

### Step 1: Create the Problem

```
# Misconfigure VPN-B to import VPN-A's multicast routes:  
  
[edit routing-instances VPN-B]  
set vrf-route-import-target target:65000:100 import-multicast # Wrong!
```

### Step 2: Verify the Problem

```
# On PE2, check if VPN-A's routes appear in VPN-B:  
show route table VPN-B.inet-vpn multicast | match "65000:100"  
  
# Expected: Routes from VPN-A (violation!)
```

### Step 3: Fix the Problem

```
# Correct VPN-B's import target:  
  
[edit routing-instances VPN-B]  
delete vrf-route-import-target target:65000:100 import-multicast  
set vrf-route-import-target target:65000:101 import-multicast
```

### Step 4: Verify the Fix

```
show route table VPN-B.inet-vpn multicast | match "65000:100"  
# Should return no results (isolation restored)
```

## Exercise 4: Hub-and-Spoke Multicast

**Objective:** Configure a customer with hub (PE1) and spokes (PE2, PE3) where multicast must flow through the hub.

### Topology:

```
Hub (PE1 - VPN-C-HUB)  
|  
└ Backbone └─ Spoke1 (PE2 - VPN-C-SPOKE1)  
    └ Spoke2 (PE3 - VPN-C-SPOKE2)
```

Multicast must flow: Spoke1 → Hub → Spoke2 (never Spoke1 → Spoke2 directly)

### Configuration:

```
# On PE1 (Hub):  
[edit routing-instances VPN-C-HUB]  
set instance-type vrf  
set interface ge-3/0/0.0  
  
set routing-options multicast rpf-check disable  
set protocols pim rp static address 192.168.3.1  
set protocols pim interface ge-3/0/0.0 mode sparse-dense  
  
set protocols pim mvpn i-pmsi  
set protocols pim mvpn rp static address 192.168.3.1  
set protocols pim mvpn root-monitor # Hub is the root  
  
set vrf-import target:65000:102  
set vrf-export target:65000:102  
set vrf-route-import-target target:65000:102 import-multicast  
set vrf-route-export-target target:65000:102 export-multicast
```

```

# On PE2/PE3 (Spokes):
[edit routing-instances VPN-C-SPOKE]
set instance-type vrf
set interface ge-1/0/0.0

set routing-options multicast rpf-check disable
set protocols pim rp static address 192.168.3.1
set protocols pim interface ge-1/0/0.0 mode sparse-dense

set protocols pim mvpn i-pmsi
set protocols pim mvpn rp static address 192.168.3.1
set protocols pim mvpn root-preference 100 # Prefer hub as root

set vrf-import target:65000:102
set vrf-export target:65000:102
set vrf-route-import-target target:65000:102 import-multicast
set vrf-route-export-target target:65000:102 export-multicast

```

#### Verification:

```

# On PE2 (Spoke), check multicast forwarding:
show pim join instance VPN-C-SPOKE family inet

# Upstream interface should point towards Hub (PE1), not directly to Spoke2 (PE3)

# Generate multicast traffic from Spoke1 to Spoke2:
# (At Spoke1) mcast_send -a 224.3.3.3 -p 5555
# (At Spoke2) mcast_receive -a 224.3.3.3 -p 5555

# Verify traffic flows through Hub by monitoring:
show multicast route family inet
# Root should be PE1, with both Spoke1 and Spoke2 as downstream

```

## Exercise 5: Debugging MVPN with Real-World Scenario

**Scenario:** A customer reports that multicast from their HQ (attached to PE1) is not reaching a branch office (attached to PE2). Diagnose the issue.

#### Given Information:

- Customer multicast source: 10.1.1.100 (HQ)
- Multicast group: 224.50.50.50
- Receiver: 10.2.2.200 (Branch)
- VRF: VPN-CORP (target: 65000:200)

#### Troubleshooting Steps:

```

# Step 1: Verify multicast is reaching PE1 from HQ
ping 224.50.50.50 routing-instance VPN-CORP
# or
traceroute multicast 10.1.1.100 224.50.50.50 routing-instance VPN-CORP

# Expected: Trace shows path from HQ to PE1

# Step 2: Check if PE1 has learned the multicast source
show pim join instance VPN-CORP family inet group 224.50.50.50

# Expected output:
# Group: 224.50.50.50
#   Source: 10.1.1.100
#     Upstream interface: ge-1/0/0.0
#     Upstream state: Join

# If no output: Multicast not reaching PE1 from HQ
#   → Check CE-PE connectivity
#   → Verify PIM enabled on CE interface

# Step 3: Check if PE1 is advertising MVPN routes
show route advertising-protocol bgp 10.255.255.2 family vpn multicast | match "224.50.50.50"

# Expected: Type 2 or Type 3 routes for this group

```

```

# If no routes: PE1 hasn't informed PE2 about the group
#   → Check BGP MVPN session (show bgp neighbor 10.255.255.2)
#   → Verify export multicast target configured

# Step 4: Check if PE2 has received the routes from PE1
show route receiving-protocol bgp 10.255.255.1 family inet-vpn multicast | match "224.50.50.50"

# Expected: Type 2/3 routes received

# If no routes: BGP session or route filtering issue
#   → Soft-reset BGP: clear bgp neighbor 10.255.255.1

# Step 5: Check if PE2 has joined the multicast tree
show pim join instance VPN-CORP family inet group 224.50.50.50

# Expected:
# Group: 224.50.50.50
# Source: 10.1.1.100
#   Upstream interface: lsi.0 (or ge-0/0/X – backbone interface)
#   Upstream state: Join

# If no join: PE2 hasn't found receivers yet
#   → Check if branch sent IGMP joins (show igmp group instance VPN-CORP)
#   → Verify PIM enabled on branch interface

# Step 6: Check multicast forwarding in backbone
show multicast route family inet source 10.1.1.100 group 224.50.50.50

# Expected:
# Group: 224.50.50.50, Source: 10.1.1.100
#   Upstream interface: (towards PE1)
#   Downstream interfaces: (towards PE2)
#   Packets: (non-zero count)

# If "Packets: 0": Traffic isn't flowing
#   → PE2 might not have receivers (check step 5)
#   → Or backbone tree not fully built

# Step 7: End-to-end check
# On branch receiver, verify reception:
show igmp group instance VPN-CORP
# Should show group 224.50.50.50 with receiver count > 0

# Monitor multicast traffic:
show multicast statistics instance VPN-CORP
# "Packets received" should be non-zero

```

#### Common Issues and Fixes:

Issue	Diagnostic Command	Fix
Receiver not joined	show pim join instance VPN-CORP (no group)	Enable IGMP or PIM on branch interface
PE2 not advertising downstream	show route advertising-protocol bgp ... (no routes)	Verify export target matches
Backbone tree not active	show multicast route family inet (no tree)	Check backbone PIM neighbors, RP config
Traffic dropping at PE1 ingress	show pim statistics (high RPF failures)	Disable or loosen RPF check
Branch receiving wrong source	show pim join (wrong upstream)	Check branch CE RP configuration

## Exercise 6: MVPN Scalability Testing

**Objective:** Deploy multiple customer VRFs with MVPN and measure scalability.

**Test Scenario:** Create 10 customer VRFs, each with 5 multicast groups. Monitor PE resource usage.

#### Configuration Template:

```

# Generate 10 VRF configurations (pseudo-code)

for i in 1 to 10 do:
    VRF_NAME = "VPN-CUSTOMER-" + i
    RT = "65000:" + (100 + i)
    INTERFACE = "ge-1/" + i + "/0.0"

```

```

set routing-instances ${VRF_NAME} instance-type vrf
set routing-instances ${VRF_NAME} interface ${INTERFACE}

set routing-instances ${VRF_NAME} routing-options multicast rpf-check disable
set routing-instances ${VRF_NAME} protocols pim rp static address 192.168.(100+i).1
set routing-instances ${VRF_NAME} protocols pim interface ${INTERFACE} mode sparse-dense

set routing-instances ${VRF_NAME} protocols pim mvpn i-pmsi
set routing-instances ${VRF_NAME} protocols pim mvpn rp static address 192.168.(100+i).1

set routing-instances ${VRF_NAME} vrf-import target:${RT}
set routing-instances ${VRF_NAME} vrf-export target:${RT}
set routing-instances ${VRF_NAME} vrf-route-import-target target:${RT} import-multicast
set routing-instances ${VRF_NAME} vrf-route-export-target target:${RT} export-multicast
end for

```

#### **Monitoring Commands:**

```

# Monitor PE resource usage:
show system processes extensive | match "memory"

# Monitor BGP MVPN routes:
show route count family inet-vpn multicast

# Monitor multicast state on PE:
show pim routes           # Total PIM state

# Monitor interfaces:
show interfaces statistics

```

#### **Expected Results:**

- BGP MVPN routes: ~50 (10 VRFs × Type 1 + multiple Type 2/3 per VRF)
- PIM join state: ~50 (5 multicast groups per VRF)
- Memory usage: Should remain stable if PE is sized correctly.

## **Exercise 7: MVPN Convergence Testing**

**Objective:** Simulate a backbone link failure and measure MVPN convergence time.

#### **Setup:**

- Two PEs with MVPN deployed.
- Multicast flowing from PE1 to PE2.
- Introduce a 2nd backbone path (redundancy).

#### **Test Steps:**

```

# 1. Start multicast traffic from PE1 to PE2
#   (On CE1) mcast_send -a 224.1.1.1 -p 5555 -b 10M

# 2. Measure baseline latency
#   (On CE2) mcast_receive -a 224.1.1.1 -p 5555
#   Note: RTT, jitter, packet loss

# 3. Fail the primary backbone link
#   (On P-router) down ge-0/0/0.0
#   or delete interface route

# 4. Timestamp the failure

# 5. Monitor PIM convergence on PE1 and PE2
#   (Every 100ms) show pim join instance VPN-A family inet group 224.1.1.1
#   Watch for upstream state change

# 6. Monitor multicast traffic recovery
#   (On CE2) mcast_receive output
#   Measure: Time until receiver resumes receiving packets

# 7. Restore the link
#   (On P-router) up ge-0/0/0.0

```

```
# 8. Measure convergence time  
#     Convergence Time = Time to resume receiving packets
```

### Expected Results:

- Convergence time: <500ms (depends on PIM timers and MPLS FRR)
  - No packet loss (if FRR is configured)
  - Jitter: Should spike during failure, normalize after recovery
- 
- 

## Summary

You now have a comprehensive understanding of BGP MVPNs:

**Module 22** established the foundational concepts:

- What multicast is and why VPNs need it.
- The architecture (CE, PE, P, RP) and roles.
- Control plane (BGP) and data plane (MPLS encapsulation).
- MVPN route types and their purposes.

**Module 23** covered practical configuration:

- I-PMSI for shared trees.
- S-PMSI for dedicated high-bandwidth trees.
- Hub-and-spoke topologies.
- Advanced MVPN scenarios.

**Lab 6** provided hands-on practice:

- From basic deployment to troubleshooting.
- Scalability testing and convergence measurement.

### Key Takeaways:

1. **MVPNs extend L3VPN** to support multicast with customer isolation.
2. **BGP carries MVPN routes** (Types 1–5) to signal group membership and tree information.
3. **Junos handles encapsulation transparently** – configuration enables it, and the OS manages labels and forwarding.
4. **Verification is methodical** – check BGP routes, PIM state, multicast trees, and end-to-end flow.
5. **Troubleshooting follows patterns** – isolate issues to backbone, BGP, VRF, or customer connectivity.

For the **JNCIE-SP exam**, you should now be able to:

- Design MVPN topologies for different customer scenarios.
- Configure I-PMSI and S-PMSI on Junos OS.
- Troubleshoot multicast isolation, route target issues, and tree convergence problems.
- Verify MVPN operations using show commands and interpret output.
- Explain the control and data plane of BGP MVPNs.

Good luck on your JNCIE-SP journey!